

Introduction to Neural Networks

From Brain to Business: How Machines Learn to Predict

Neural Networks for Business Applications

November 23, 2025

By the End of This Lecture, You Will Be Able To:

- **Explain** how biological neurons inspire artificial neural networks
- **Calculate** the output of an artificial neuron given inputs and weights
- **Design** a simple multilayer network architecture for a business problem
- **Trace** information flow through forward propagation
- **Describe** how networks learn by minimizing prediction errors
- **Evaluate** when neural networks are appropriate for business predictions
- **Assess** the ethical implications of automated prediction systems

The Prediction Challenge: Can We Predict Markets?

The Business Question:

- Can we predict if a stock price will rise or fall tomorrow?
- Traditional methods: Statistical analysis, expert intuition, rule-based systems
- Challenge: Markets are **complex, non-linear systems**
- Many interacting factors: price history, volume, sentiment, volatility, interest rates

Why This Matters:

- Better investment decisions (\rightarrow higher returns)
- Risk management (\rightarrow protect capital)
- Portfolio optimization (\rightarrow balanced exposure)
- Automated trading strategies (\rightarrow scalability)

Our journey begins with understanding how nature solved similar prediction problems

What We Need: A Learning System

Requirements for Market Prediction System

A system that can:

1. Process multiple inputs simultaneously
2. Learn patterns from historical data
3. Handle **non-linear** relationships
4. Improve predictions over time
5. Generalize to new market conditions

Inspiration from Nature

The human brain solves complex pattern recognition tasks every day by learning from experience.

Can we mimic this for business predictions?

► Let's explore how brains work...

Next: Understanding biological neurons as the foundation

Biological Neuron Structure

- **Dendrites:** Receive signals from other neurons
- **Soma:** Integrates incoming signals with weights
- **Axon:** Transmits output to next neurons
- **Synapses:** Connection points with varying strengths

Key Principle

Neurons process multiple weighted inputs and fire when threshold is exceeded.


Business AI Insights

1. **Multiple inputs combined**
→ Consider many market factors
2. **Weighted connections**
→ Some factors matter more
3. **Non-linear activation**
→ Threshold effects (tipping points)
4. **Layered processing**
→ Abstract reasoning emerges

We can create mathematical models that learn the same way!

Next: See the visual comparison of biological vs artificial neurons

From Biological Intelligence to Business AI



The artificial neuron mathematically mimics biological signal integration and activation

Step 1: Weighted Sum

Mimics soma integration:

$$z = \sum_{i=1}^n w_i x_i + b$$

- x_i : Input features (market data)
- w_i : Weights (learned from data)
- b : Bias term (baseline adjustment)

Like dendrites weighting signals differently

Step 2: Activation Function

Mimics axon firing:

$$y = f(z) = \frac{1}{1 + e^{-z}}$$

- f : Activation function (non-linearity)
- Output: probability between 0 and 1
- Mimics neuron “firing” at threshold


Complete Formula:

$$y = \sigma \left(\sum_{i=1}^n w_i x_i + b \right)$$

Next: See a concrete example with real market numbers

Single Neuron Computation: Step-by-Step Example

How a Neuron Computes: Step-by-Step



With market inputs (price=100, volume=85, sentiment=120), the neuron predicts 100% probability of price increase

Activation Functions: Why Non-Linearity Matters

The Problem

Without activation functions:

- Neural networks = fancy linear regression
- Real business relationships are **non-linear!**

Three Common Functions

- **Sigmoid:** Smooth (0,1) range
Perfect for probabilities
- **ReLU:** Fast, efficient
Most popular in modern networks
- **Tanh:** Zero-centered (-1,1)
Alternative to sigmoid

Business Non-Linearity Examples




1. **Diminishing returns**
Doubling marketing spend doesn't double sales
2. **Threshold effects**
Sentiment must reach tipping point
3. **Saturation**
Engagement plateaus beyond certain point
4. **Network effects**
Value increases non-linearly with users

Activation functions let networks capture these patterns!

Next: Visual comparison of these three activation functions

Activation Functions: Visual Comparison

Activation Functions: Adding Non-Linearity



Each activation function has different properties: Sigmoid for probabilities, ReLU for speed, Tanh for zero-centered outputs

The Limitation: Why One Neuron Is Not Enough

What One Neuron Can Do

- Draw a single straight line (hyperplane)
- Separate **linearly separable** patterns
- Example: “Buy if price > 100 AND volume > 50”

Business Analogy

One simple rule for decision-making

Limitation

Only linear decision boundaries possible

What One Neuron Cannot Do

- Complex, curved decision boundaries
- XOR-like patterns: “Buy when high price XOR high volume”
- **Real-world market patterns!**



Business Reality

Multiple interacting factors, non-linear relationships, conditional dependencies

Solution: Multiple Layers!

Next: See the XOR problem that proves one neuron's limitation

Why One Neuron Is Not Enough



Solution: Use Multiple Layers (Hidden Layers) to Create Non-Linear Decision Boundaries

Left: Linearly separable (one neuron works). Right: XOR pattern (one neuron fails, need hidden layers)

Multi-Layer Architecture

- **Input Layer:** Raw market features
No computation, like sensory neurons
- **Hidden Layer(s):** Pattern detection
Feature combinations, like association cortex
Learns “high volume + rising price = momentum”
- **Output Layer:** Final prediction
Probability output, like motor neurons

Result: Buy/Sell decision

Hierarchical Learning Principle

- **Layer 1:** Detects simple patterns
“price rising”, “volume high”
- **Layer 2:** Combines into complex patterns
“strong momentum”, “weak support”
- **Layer 3:** Makes strategic decisions
“high probability buy signal”


Key Insight

Each layer builds on previous layer's abstractions

Next: See the full network architecture with all connections

Neural Network Architecture Diagram

Neural Network Architecture: Building Intelligence in Layers



5 inputs → 6 hidden neurons → 1 output. Total: 36 weights to learn from data

The Forward Pass Process

1. **Input:** Feed market features
2. **Hidden Layer:**

$$a^{(1)} = \sigma(W^{(1)}x + b^{(1)})$$

Detects patterns like “rising trend”

3. **Output Layer:**

$$y = \sigma(W^{(2)}a^{(1)} + b^{(2)})$$

Efficiency

All neurons in layer compute in parallel!

Concrete Example

Input: price=105.2, volume=0.75, sentiment=0.62

Hidden Layer: Detects patterns

Output: $y = 0.742$


Interpretation:

- 74.2% confidence price will rise
- If $y > 0.5 \rightarrow \text{BUY}$
- If $y < 0.5 \rightarrow \text{SELL}$

Next: See forward propagation with actual numbers and calculations

Forward Propagation: Detailed Example

Forward Propagation: Making a Prediction



Data flows left to right: inputs (105.2, 0.75, 0.62) → hidden activations → output (0.742) = 74% buy confidence

Learning Process Steps

1. Predict with random weights
2. Measure error (Loss Function):

$$L = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

3. Adjust weights (Gradient Descent):

$$w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$$

4. Repeat thousands of times

η = learning rate (how fast we learn)

Concrete Example

Initial: Network predicts 55% price rise
Actual: Price fell ($y = 0$)
Error: $(0 - 0.55)^2 = 0.30$



Learning Step:

- Calculate gradient (direction of error)
- Move weights opposite direction
- Error decreases each iteration

Like a trader learning from past mistakes

Next: Visualize the loss landscape that we're trying to navigate

Loss Landscape: The Error Surface



Goal: Find the weights that minimize the loss
(The red star shows the optimal solution)

Goal: Find weights (red star) that minimize loss. Different starting points converge to same optimum through gradient descent

Algorithm Steps

1. Calculate gradient:

$$\frac{\partial L}{\partial w} = \text{slope of loss}$$

2. Step opposite direction:

$$w_{new} = w_{old} - \eta \times \text{gradient}$$

3. Repeat until convergence

Learning Rate Trade-offs

- **Too small:** Slow learning
- **Too large:** Unstable, overshoots
- **Just right:** Steady progress

Business Analogy

Like a trader learning from mistakes:

- **Fast learning phase**
Rapid improvement from obvious patterns
- **Steady progress**
Fine-tuning strategy
- **Convergence**
Optimal trading rules learned


Key Insight

Gradient tells us which direction reduces error fastest

Next: See how loss decreases over training iterations

Gradient Descent: Optimization in Action

Gradient Descent: Learning by Stepping Downhill



Left: 1D visualization showing steps toward minimum. Right: Loss decreasing over 100 training iterations

Putting It Together: Market Prediction Case Study

Business Application

- **Goal:** Predict if stock price rises/falls
- **Data:** 60 days historical market data
- **Features:** 4 input variables per day

Input Features

1. Historical Stock Price
2. Trading Volume (normalized)
3. Market Sentiment (0-1)
4. Volatility Index (0-1)

Target Variable

Binary output:

- 1 = price increased
- 0 = price decreased

Network outputs: $p(\text{price rise})$


Training Setup

- Training: First 45 days (learn)
- Test: Last 15 days (evaluate)
- Network: 4 → 6 → 1

Next: See the actual market data used for training

Market Data: Input Features for Neural Network

Market Data: Input Features for Neural Network



Neural Network Input: All 4 features for each day | Output: Probability of price increase tomorrow

60 days of market data: price trend (up/down markers), volume bars, sentiment score, volatility index

Training Results: Before vs After

The Experiment

- **Before training:**
Random weights → coin flip predictions
- **After training:**
Learned weights → intelligent predictions
- **Test set:** 30 days unseen data

Key Results

- **Before:** ≈50% accuracy
- **After:** ≈70% accuracy
- **Improvement:** +20 percentage points

What the Network Learned

- High volume + rising price + positive sentiment = likely rise
- Low volume + high volatility = uncertain
- Sentiment lags price but confirms trends


Discovered from data alone!

Reality Check

- 70% is **good** for markets
- 100% is **impossible**
- 70% over many trades = profit

Next: See detailed before/after comparison with prediction accuracy

Prediction Results: Before vs After Training



Top: Actual prices. Middle: Before training (random, 50% accuracy). After training (learned patterns, 70% accuracy)

Our Journey (Concept → Visualization):

1. Biology: Neurons integrate weighted signals → Chart
2. Mathematics: $y = f(\sum w_i x_i + b)$ → Chart
3. Non-linearity: Activation functions enable complexity → Chart
4. Limitation: One neuron fails on XOR → Chart
5. Architecture: Layers learn hierarchical patterns → Chart
6. Forward prop: Making predictions → Chart
7. Learning: Loss landscape → Chart
8. Optimization: Gradient descent → Chart
9. Application: Market data → Chart
10. Results: 50% → 70% accuracy → Chart

Key Insight:

Neural networks **learn from data**, not explicit programming!

The separation of concept and visualization helped you see both theory and practice

Use Neural Networks When:

- Large dataset (thousands+ examples)
- Complex patterns (non-linear relationships)
- Difficult to specify rules (too many cases)
- Pattern recognition (images, speech, text)
- Acceptable black-box (don't need to explain every decision)
- Computational resources available

Example Applications:

- Customer churn prediction
- Fraud detection
- Recommendation systems
- Image recognition
- Natural language processing

Do NOT Use When:

- Small dataset (few hundred examples)
- Simple linear relationships
- Need interpretability (regulatory requirements)
- Rules are known (use rule-based system instead)
- High stakes, low data (medical diagnosis with small samples)
- Real-time constraints (millisecond decisions)

Better Alternatives:

- Linear/logistic regression (interpretable)
- Decision trees (explainable)
- Expert systems (known rules)
- Statistical models (small data)

Choose the right tool for the problem - neural networks are powerful but not always appropriate

Technical Limitations:

- **Data hungry:** Need thousands of examples
- **Black box:** Hard to explain decisions to regulators
- **Overfitting:** May memorize training data, fail on new data
- **No guarantees:** Markets are inherently unpredictable
- **Computational cost:** Training requires GPUs, time, energy
- **Hyperparameter sensitivity:** Architecture choices matter

Ethical Responsibilities:

- **Fairness:** Biased data → biased predictions
 - Historical hiring data may encode discrimination
 - Loan approval systems may disadvantage minorities
- **Transparency:** Can you explain decisions to stakeholders?
 - GDPR requires “right to explanation”
 - Use interpretability tools (SHAP, LIME)
- **Accountability:** Who is responsible when AI fails?
 - Flash crash: Automated trading gone wrong
 - Wrong medical diagnosis: Who pays?
- **Societal impact:** Unintended consequences
 - Job displacement through automation
 - Algorithmic trading destabilizes markets

With great predictive power comes great responsibility!

Always consider ethical implications before deploying AI systems in real-world business contexts

Appendix: Mathematical Details (Backpropagation)

How Gradient Descent Works: The Chain Rule

For a simple 2-layer network, we compute gradients layer-by-layer working backwards:

Output Layer Gradient:

$$\begin{aligned}\frac{\partial L}{\partial w^{(2)}} &= \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial z^{(2)}} \cdot \frac{\partial z^{(2)}}{\partial w^{(2)}} \\ &= (\hat{y} - y) \cdot \sigma'(z^{(2)}) \cdot a^{(1)}\end{aligned}$$

Hidden Layer Gradient (using chain rule):

$$\begin{aligned}\frac{\partial L}{\partial w^{(1)}} &= \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial z^{(2)}} \cdot \frac{\partial z^{(2)}}{\partial a^{(1)}} \cdot \frac{\partial a^{(1)}}{\partial z^{(1)}} \cdot \frac{\partial z^{(1)}}{\partial w^{(1)}} \\ &= (\hat{y} - y) \cdot \sigma'(z^{(2)}) \cdot w^{(2)} \cdot \sigma'(z^{(1)}) \cdot x\end{aligned}$$

Common Loss Functions:

- **Mean Squared Error (Regression):** $L = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- **Binary Cross-Entropy (Classification):** $L = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$

Backpropagation efficiently computes how each weight contributed to the final error

Practical Tips for Building Business Neural Networks:

- **Start simple:** Try linear/logistic regression first as baseline
- **Feature engineering matters:** Good inputs > complex architecture
- **Avoid overfitting:** Use validation sets, regularization (L1/L2), dropout
- **Hyperparameter tuning:** Learning rate, architecture, batch size, epochs
- **Interpretability tools:** SHAP values, attention weights, feature importance
- **Monitor training:** Plot loss curves, check for convergence
- **Cross-validation:** Don't rely on single train/test split

Recommended Resources:

Books:

- Deep Learning (Goodfellow et al.)
- Neural Networks and Deep Learning (Nielsen)
- Hands-On Machine Learning (Geron)

Courses:

- Andrew Ng - Machine Learning (Coursera)
- Fast.ai - Deep Learning for Coders
- MIT 6.S191 - Intro to Deep Learning

Best way to learn: Build real projects with real data!

Tools & Frameworks:

- PyTorch (flexible, research-friendly)
- TensorFlow/Keras (production-ready)
- scikit-learn (simpler models)

Practice Datasets:

- Kaggle competitions
- Yahoo Finance API
- UCI ML Repository

Appendix: Practice Problem for Business Students

Design Challenge: You are a data scientist at a retail company.

Problem: Predict customer churn (will customer leave next month?)

Available Data:

- Customer demographics (age, location, income)
- Purchase history (frequency, recency, monetary value)
- Customer service interactions (calls, complaints, resolutions)
- Website engagement (visits, time spent, pages viewed)

Your Tasks (Work in Groups):

1. Design a neural network architecture: How many layers? How many neurons per layer?
2. What would be your input features? Raw data or engineered features?
3. What activation functions would you use and where?
4. What loss function is appropriate for this problem?
5. How would you evaluate model performance? (accuracy, precision, recall, F1?)
6. What are potential ethical concerns with automated churn prediction?
7. How would you explain predictions to business stakeholders?
8. When would you NOT use a neural network for this problem?

Discuss in groups and present your design - there's no single right answer!