

Introduction to Neural Networks

From Brain to Business: How Machines Learn to Predict

Neural Networks for Business Applications

November 24, 2025

Learning Objectives

- **Explain** how biological neurons inspire artificial neural networks
- **Calculate** the output of an artificial neuron given inputs and weights
- **Design** a simple multilayer network architecture
- **Trace** information flow through forward propagation
- **Describe** how networks learn by minimizing prediction errors
- **Evaluate** when neural networks are appropriate for business

The Prediction Challenge: Can We Predict Markets?

The Business Question

Can we predict if a stock price will rise or fall tomorrow?

- Traditional: Statistical analysis, expert intuition
- Challenge: Markets are **complex, non-linear**
- Many factors: price, volume, sentiment, volatility

The Limitation

Rule-based systems cannot capture all interactions

Why This Matters

- Better investment decisions
- Risk management
- Portfolio optimization
- Automated trading strategies

What We Need

A system that learns patterns from data, not explicit rules

Our journey begins with understanding how nature solved similar prediction problems

System Requirements

1. Process multiple inputs simultaneously
2. Learn patterns from historical data
3. Handle **non-linear** relationships
4. Improve predictions over time
5. Generalize to new conditions

Key Insight

We need a system that learns, not one we program

Inspiration from Nature

The human brain solves complex pattern recognition every day

Brain Capabilities

- Processes millions of inputs
- Learns from experience
- Handles ambiguity
- Generalizes to new situations

Can we mimic this for business predictions?

Next: Understanding biological neurons as the foundation

Part 1: Foundations

From biological neurons to artificial intelligence

Biological Neuron Structure

- **Dendrites:** Receive signals
- **Soma:** Integrates weighted signals
- **Axon:** Transmits output
- **Synapses:** Variable connection strengths

Key Principle

Fire when weighted sum exceeds threshold


Business AI Insights

1. Multiple inputs combined
2. Weighted connections (importance)
3. Non-linear activation (thresholds)
4. Layered processing (abstraction)

Mathematical models can learn the same way!

Next: See the visual comparison of biological vs artificial neurons

From Biological Intelligence to Business AI



The artificial neuron mathematically mimics biological signal integration and activation

Step 1: Weighted Sum

$$z = \sum_{i=1}^n w_i x_i + b$$

- x_i : Inputs (market data)
- w_i : Weights (learned)
- b : Bias (baseline)

Step 2: Activation

$$y = \sigma(z) = \frac{1}{1 + e^{-z}}$$


- Adds non-linearity
- Output: probability (0 to 1)
- Mimics neuron firing

Complete: $y = \sigma(\sum w_i x_i + b)$

Next: See a concrete example with real market numbers

Single Neuron Computation: Step-by-Step Example

How a Neuron Computes: Step-by-Step



With market inputs (price=100, volume=85, sentiment=120), the neuron predicts 100% probability of price increase

Part 2: Building Blocks

Activation functions and their role in learning

Activation Functions: Why Non-Linearity Matters

The Problem

Without activation functions:

- Networks = linear regression
- Cannot learn complex patterns

Three Common Functions

- **Sigmoid:** (0,1) for probabilities
- **ReLU:** Fast, efficient
- **Tanh:** Zero-centered (-1,1)

Business Non-Linearity




1. Diminishing returns
2. Threshold effects
3. Saturation points
4. Network effects

Activation functions capture these patterns!

Next: Visual comparison of these three activation functions

Activation Functions: Visual Comparison

Activation Functions: Adding Non-Linearity



Each activation function has different properties: Sigmoid for probabilities, ReLU for speed, Tanh for zero-centered outputs

The Limitation: Why One Neuron Is Not Enough

What One Neuron Can Do

- Single straight decision boundary
- Separate linearly separable patterns
- Simple rules only

Analogy: One rule for decisions



What One Neuron Cannot Do

- Complex, curved boundaries
- XOR-like patterns
- Real-world market interactions

Solution: Multiple Layers!

Next: See the XOR problem that proves one neuron's limitation

Why One Neuron Is Not Enough



Solution: Use Multiple Layers (Hidden Layers) to Create Non-Linear Decision Boundaries

Left: Linearly separable (one neuron works). Right: XOR pattern (one neuron fails, need hidden layers)

Part 3: Network Architecture

Building layers of intelligence

Multi-Layer Architecture

- **Input:** Raw features (no computation)
- **Hidden:** Pattern detection
- **Output:** Final prediction

Result: Buy/Sell decision

Hierarchical Learning


- **Layer 1:** Simple patterns
- **Layer 2:** Complex patterns
- **Layer 3:** Strategic decisions

Each layer builds on previous abstractions

Next: See the full network architecture with all connections

Neural Network Architecture Diagram

Neural Network Architecture: Building Intelligence in Layers



5 inputs, 6 hidden neurons, 1 output. Total: 36 weights to learn from data

Forward Propagation: How Networks Make Predictions

The Forward Pass

1. **Input:** Feed market features
2. **Hidden:** $a = \sigma(Wx + b)$
3. **Output:** $y = \sigma(Wa + b)$

All neurons compute in parallel!

Example

Input: price=105.2, volume=0.75
Output: $y = 0.742$


Interpretation:

- 74.2% confidence price rises
- $y > 0.5$: **BUY**
- $y < 0.5$: **SELL**

Next: See forward propagation with actual numbers and calculations

Forward Propagation: Detailed Example

Forward Propagation: Making a Prediction



Data flows left to right: inputs, hidden activations, output (0.742) = 74% buy confidence

Part 4: Learning Process

How networks learn from mistakes

Learning Steps

1. Predict with random weights

2. Measure error:

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

3. Adjust weights:

$$w_{new} = w_{old} - \eta \nabla L$$

4. Repeat until convergence

Example

Predicted: 55% rise, Actual: fell

$$\text{Error: } (0 - 0.55)^2 = 0.30$$


Learning:

- Calculate gradient direction
- Move weights to reduce error

Like a trader learning from mistakes

Next: Visualize the loss landscape that we're trying to navigate

Loss Landscape: The Error Surface



Goal: Find weights (red star) that minimize loss. Gradient descent navigates this surface

Algorithm

1. Calculate gradient (slope)
2. Step opposite direction
3. Repeat until convergence

Learning Rate Trade-offs

- **Too small:** Slow
- **Too large:** Unstable
- **Just right:** Steady

Business Analogy

Like a trader learning:


- Fast learning from obvious patterns
- Steady fine-tuning
- Convergence to optimal rules

Gradient shows fastest error reduction

Next: See how loss decreases over training iterations

Gradient Descent: Optimization in Action

Gradient Descent: Learning by Stepping Downhill



Left: 1D visualization showing steps toward minimum. Right: Loss decreasing over 100 training iterations

Part 5: Application

Putting it all together with market prediction

Putting It Together: Market Prediction Case Study

Business Application

- **Goal:** Predict price direction
- **Data:** 60 days market data

Input Features

1. Stock Price
2. Trading Volume
3. Market Sentiment
4. Volatility Index

Target Variable

Binary: 1 = up, 0 = down
Network outputs: $p(\text{rise})$


Setup

- Train: 45 days
- Test: 15 days
- Network: 4-6-1

Next: See the actual market data used for training

Market Data: Input Features for Neural Network

Market Data: Input Features for Neural Network



Neural Network Input: All 4 features for each day | Output: Probability of price increase tomorrow

60 days of market data: price trend, volume bars, sentiment score, volatility index

Training Results: Before vs After

The Experiment

- Before: Random weights (coin flip)
- After: Learned weights
- Test: 30 days unseen data

Results

- **Before:** 50% accuracy
- **After:** 70% accuracy
- **Gain:** +20 points

What Network Learned


- Volume + price + sentiment patterns
- Volatility indicates uncertainty
- Sentiment confirms trends

Discovered from data alone!

70% is good for markets (100% impossible)

Next: See detailed before/after comparison with prediction accuracy

Prediction Results: Before vs After Training



Top: Actual prices. Middle: Before (random). Bottom: After (learned patterns, 70% accuracy)

Our Journey

1. Biological inspiration
2. Mathematical model
3. Activation functions
4. Network architecture
5. Forward propagation

Learning Process

6. Loss functions
7. Gradient descent
8. Training iterations
9. Pattern discovery
10. Improved predictions

Key Takeaway

Neural networks **learn patterns from data** rather than following explicit rules. This enables them to discover complex relationships that would be impossible to program manually.

Each concept was paired with a visualization to reinforce understanding

When to Use Neural Networks

Use Neural Networks When

- Large dataset (thousands+ examples)
- Complex patterns
- Difficult to specify rules
- Pattern recognition tasks
- Black-box acceptable

Applications

Churn, fraud, recommendations, images, NLP

Do NOT Use When

- Small dataset
- Simple relationships
- Need interpretability
- Rules are known
- Real-time constraints

Alternatives

Regression, decision trees, expert systems

Choose the right tool - neural networks are powerful but not always appropriate

Technical Limitations

- Data hungry
- Black box decisions
- Overfitting risk
- No guarantees
- Computational cost

Ethical Concerns

- **Fairness:** Biased data leads to biased predictions
- **Transparency:** GDPR requires explanations
- **Accountability:** Who is responsible?
- **Impact:** Job displacement, market stability

With great predictive power comes great responsibility!

Always consider ethical implications before deploying AI systems

Appendix: Mathematical Details (Backpropagation)

Output Layer Gradient

$$\frac{\partial L}{\partial w^{(2)}} = (\hat{y} - y) \cdot \sigma'(z) \cdot a$$

Hidden Layer Gradient

$$\frac{\partial L}{\partial w^{(1)}} = \delta^{(2)} \cdot w^{(2)} \cdot \sigma'(z^{(1)}) \cdot x$$

Loss Functions

MSE (Regression):

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

Cross-Entropy (Classification):

$$L = - \sum [y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

Backpropagation efficiently computes how each weight contributed to the error

Practical Tips

- Start simple (baseline first)
- Feature engineering matters
- Avoid overfitting (validation, dropout)
- Tune hyperparameters
- Monitor training curves

Books

Goodfellow (Deep Learning), Nielsen, Geron

Courses

- Andrew Ng (Coursera)
- Fast.ai
- MIT 6.S191

Tools

PyTorch, TensorFlow, scikit-learn

Practice

Kaggle, Yahoo Finance, UCI Repository

Best way to learn: Build real projects with real data!

Appendix: Practice Problem for Business Students

Design Challenge

You are a data scientist at a retail company.

Problem: Predict customer churn

Data Available:

- Demographics
- Purchase history
- Service interactions
- Website engagement

Your Tasks

1. Design network architecture
2. Select input features
3. Choose activation functions
4. Select loss function
5. Define evaluation metrics
6. Identify ethical concerns
7. Plan stakeholder explanation

Discuss in groups - there's no single right answer!