

# Introduction to Neural Networks

From Brain to Business: How Machines Learn to Predict

## Neural Networks for Business Applications

November 24, 2025

### Learning Objectives

- **Explain** how biological neurons inspire artificial neural networks
- **Calculate** the output of an artificial neuron given inputs and weights
- **Design** a simple multilayer network architecture
- **Trace** information flow through forward propagation
- **Describe** how networks learn by minimizing prediction errors
- **Evaluate** when neural networks are appropriate for business

# The Prediction Challenge: Can We Predict Markets?

## The Business Question

Can we predict if a stock price will rise or fall tomorrow?

- Traditional: Statistical analysis, expert intuition
- Challenge: Markets are **complex, non-linear**
- Many factors: price, volume, sentiment, volatility

## The Limitation

Rule-based systems cannot capture all interactions

## Why This Matters

- Better investment decisions
- Risk management
- Portfolio optimization
- Automated trading strategies

## What We Need

A system that learns patterns from data, not explicit rules

---

Our journey begins with understanding how nature solved similar prediction problems

## System Requirements

1. Process multiple inputs simultaneously
2. Learn patterns from historical data
3. Handle **non-linear** relationships
4. Improve predictions over time
5. Generalize to new conditions

## Key Insight

We need a system that learns, not one we program

## Inspiration from Nature

The human brain solves complex pattern recognition every day

## Brain Capabilities

- Processes millions of inputs
- Learns from experience
- Handles ambiguity
- Generalizes to new situations

*Can we mimic this for business predictions?*

---

Next: Understanding biological neurons as the foundation

## Part 1: Foundations

From biological neurons to artificial intelligence

[1] – [2] – [3] – [4] – [5]

## Biological Neuron Structure

- **Dendrites:** Receive signals
- **Soma:** Integrates weighted signals
- **Axon:** Transmits output
- **Synapses:** Variable connection strengths

## Key Principle

Fire when weighted sum exceeds threshold

## Business AI Insights

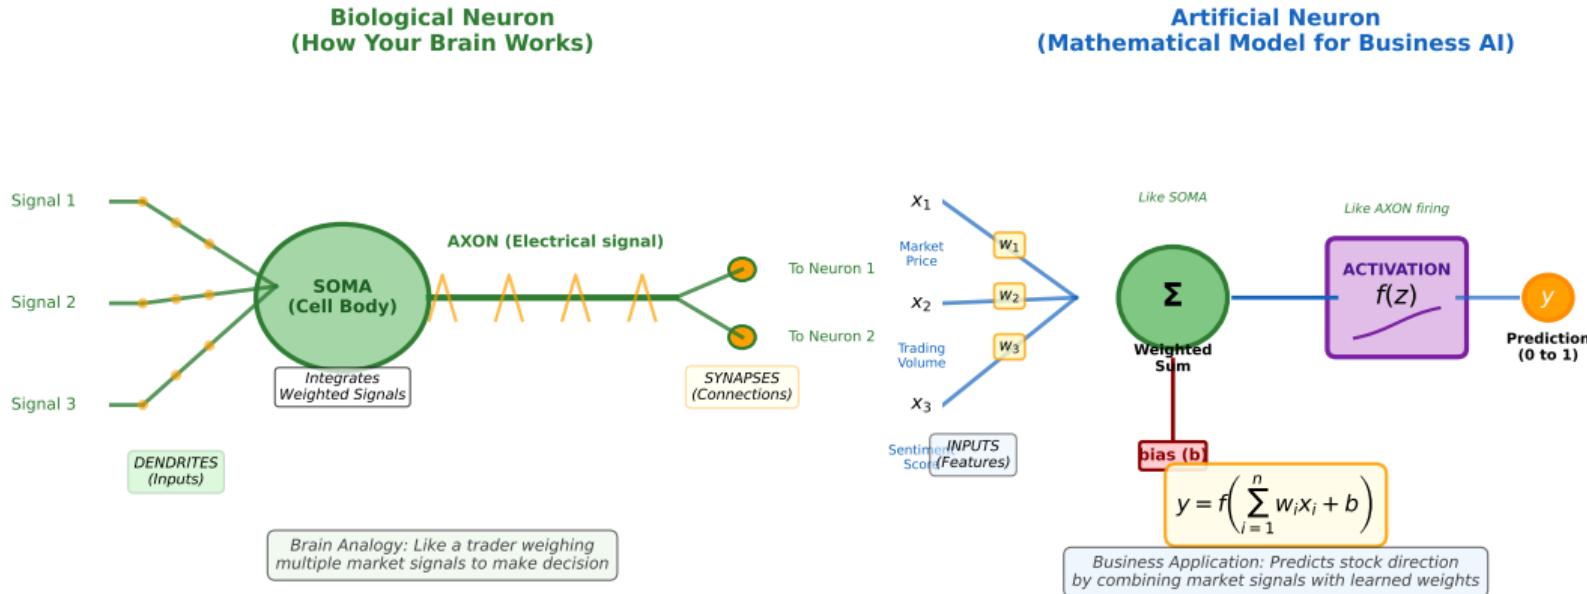
1. Multiple inputs combined
2. Weighted connections (importance)
3. Non-linear activation (thresholds)
4. Layered processing (abstraction)

*Mathematical models can learn the same way!*

---

Next: See the visual comparison of biological vs artificial neurons

## From Biological Intelligence to Business AI



Observe: Which biological components map directly to mathematical operations?

## Step 1: Weighted Sum

$$z = \sum_{i=1}^n w_i x_i + b$$

- $x_i$ : Inputs (market data)
- $w_i$ : Weights (learned)
- $b$ : Bias (baseline)

## Step 2: Activation

$$y = \sigma(z) = \frac{1}{1 + e^{-z}}$$

- Adds non-linearity
- Output: probability (0 to 1)
- Mimics neuron firing

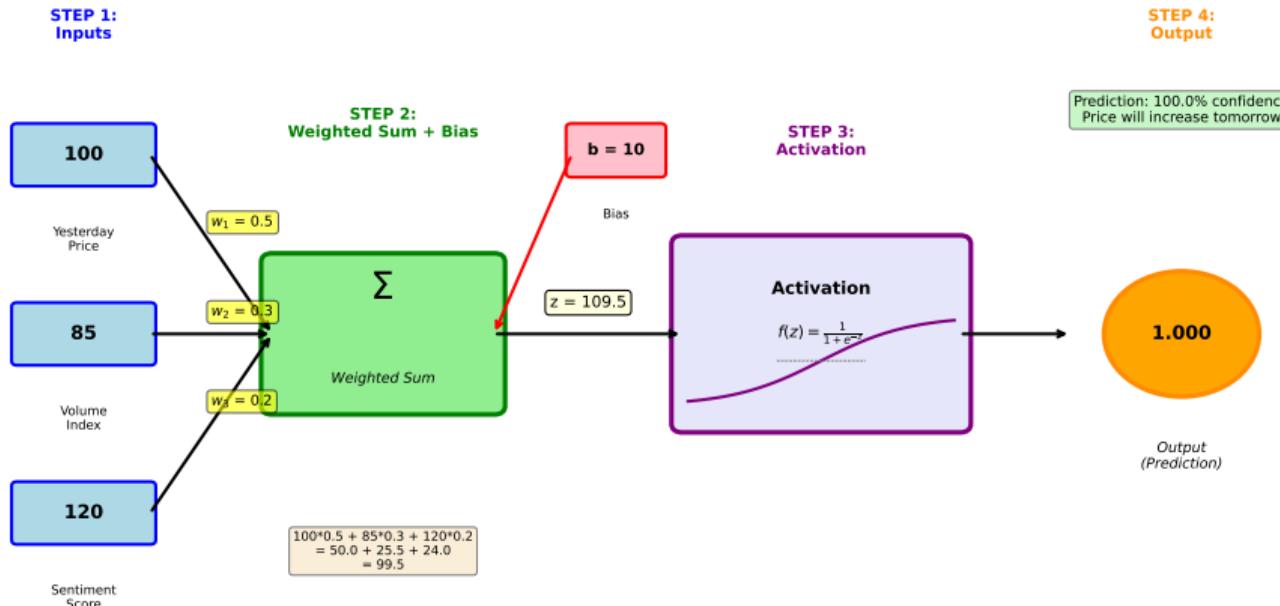
Complete:  $y = \sigma(\sum w_i x_i + b)$

---

Next: See a concrete example with real market numbers

# Single Neuron Computation: Step-by-Step Example

## How a Neuron Computes: Step-by-Step



Observe: How would changing the weights affect the final output probability?

### Think – Pair – Share

*What other business processes might benefit from 'learning from data' instead of following explicit rules?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 2: Building Blocks

Activation functions and their role in learning

[1] – [2] – [3] – [4] – [5]

# Activation Functions: Why Non-Linearity Matters

## The Problem

Without activation functions:

- Networks = linear regression
- Cannot learn complex patterns

## Three Common Functions

- **Sigmoid:** (0,1) for probabilities
- **ReLU:** Fast, efficient
- **Tanh:** Zero-centered (-1,1)

## Business Non-Linearity

1. Diminishing returns
2. Threshold effects
3. Saturation points
4. Network effects

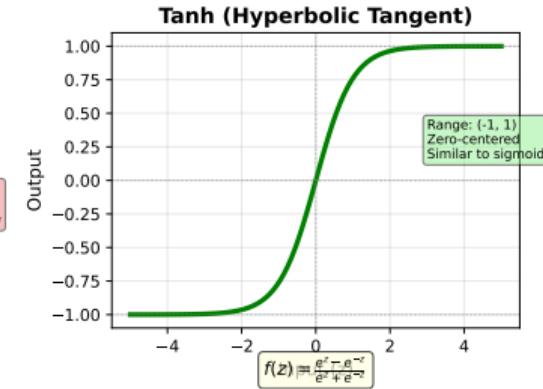
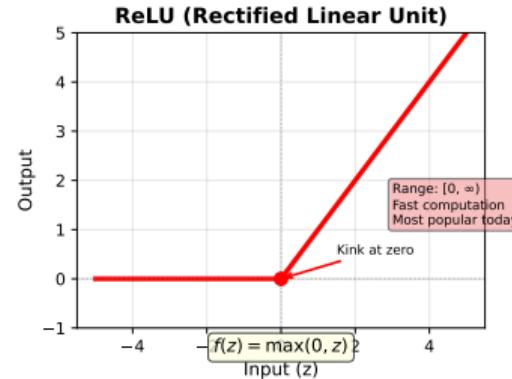
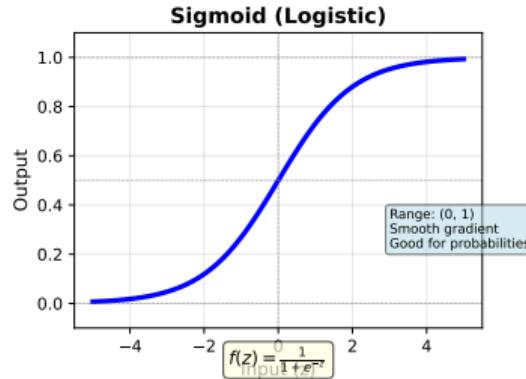
*Activation functions capture these patterns!*

---

Next: Visual comparison of these three activation functions

# Activation Functions: Visual Comparison

## Activation Functions: Adding Non-Linearity



Observe: Where does each function's output change most rapidly? Why does this matter?

# The Limitation: Why One Neuron Is Not Enough

## What One Neuron Can Do

- Single straight decision boundary
- Separate linearly separable patterns
- Simple rules only

**Analogy:** One rule for decisions

## What One Neuron Cannot Do

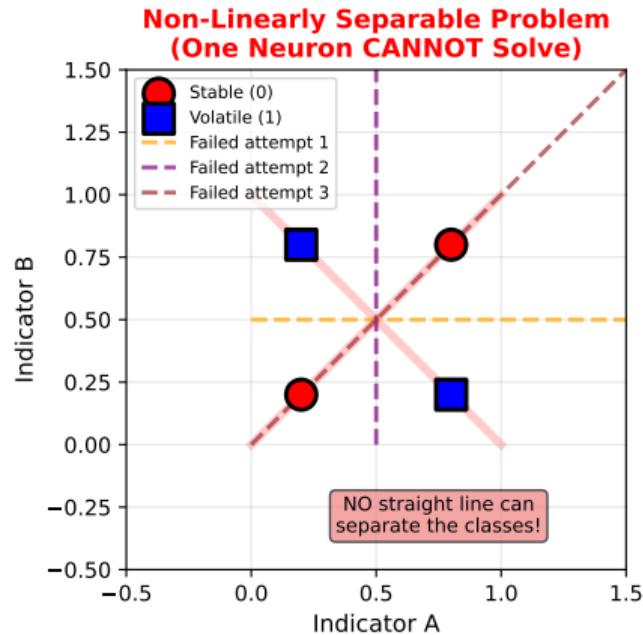
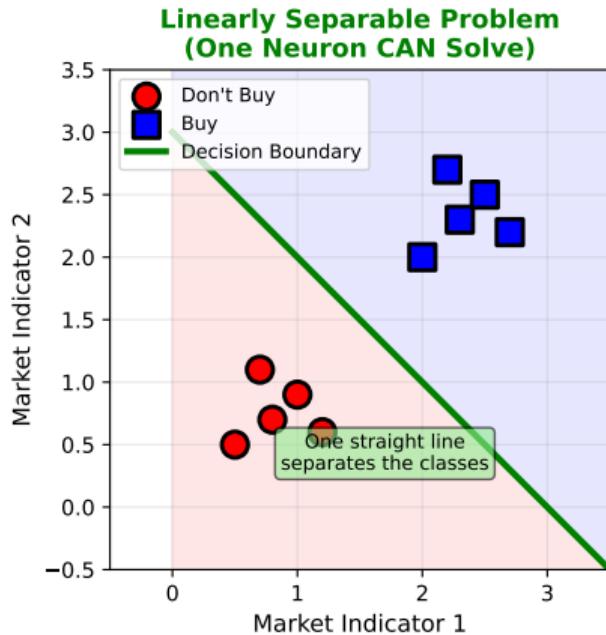
- Complex, curved boundaries
- XOR-like patterns
- Real-world market interactions

**Solution: Multiple Layers!**

---

Next: See the XOR problem that proves one neuron's limitation

## Why One Neuron Is Not Enough



**Solution: Use Multiple Layers (Hidden Layers) to Create Non-Linear Decision Boundaries**

Observe: Why is it impossible to draw a single straight line separating orange from blue?

### Think – Pair – Share

*Can you think of a business metric that shows diminishing returns or threshold effects?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 3: Network Architecture

Building layers of intelligence

[1] – [2] – **[3]** – [4] – [5]

## Multi-Layer Architecture

- **Input:** Raw features (no computation)
- **Hidden:** Pattern detection
- **Output:** Final prediction

**Result:** Buy/Sell decision

## Hierarchical Learning

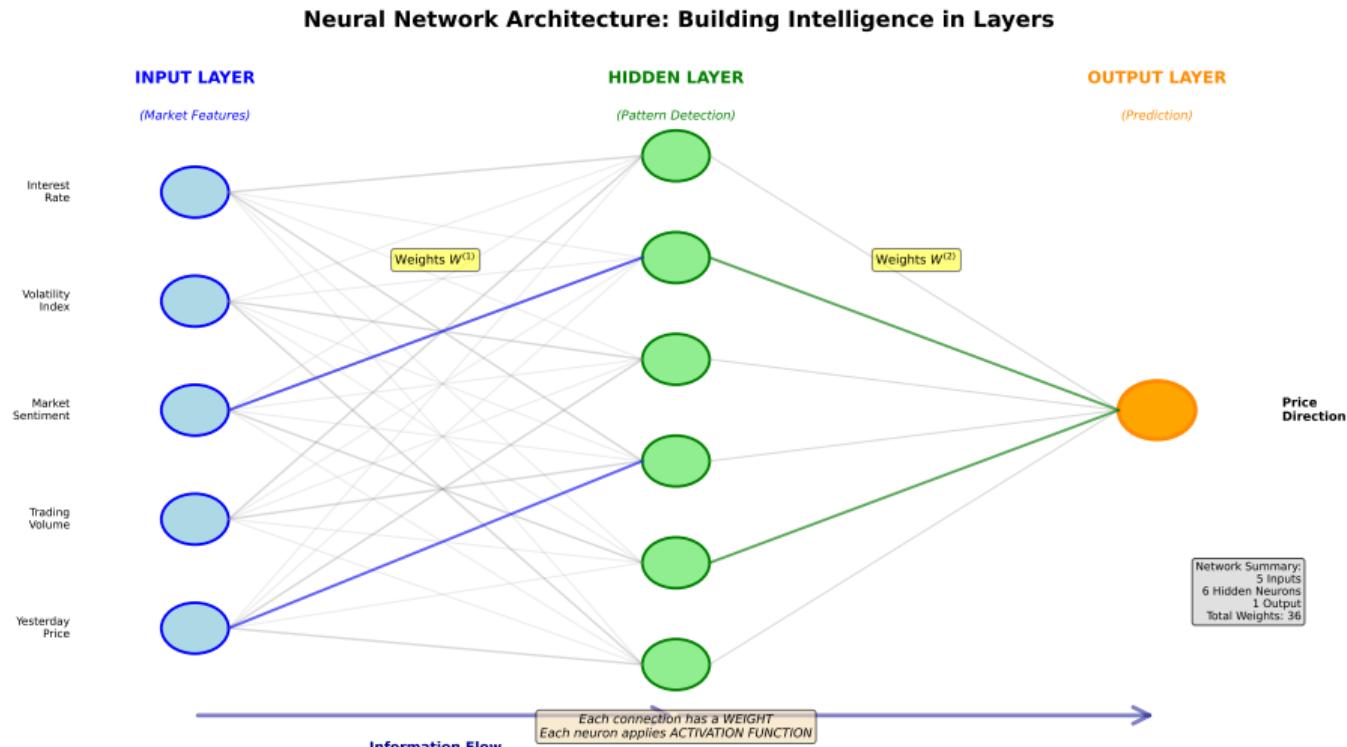
- **Layer 1:** Simple patterns
- **Layer 2:** Complex patterns
- **Layer 3:** Strategic decisions

Each layer builds on previous abstractions

---

Next: See the full network architecture with all connections

# Neural Network Architecture Diagram



Observe: Count the connections. Why are there 36 weights total?

# Forward Propagation: How Networks Make Predictions

## The Forward Pass

1. **Input:** Feed market features
2. **Hidden:**  $a = \sigma(Wx + b)$
3. **Output:**  $y = \sigma(Wa + b)$

All neurons compute in parallel!

## Example

Input: price=105.2, volume=0.75  
Output:  $y = 0.742$

### Interpretation:

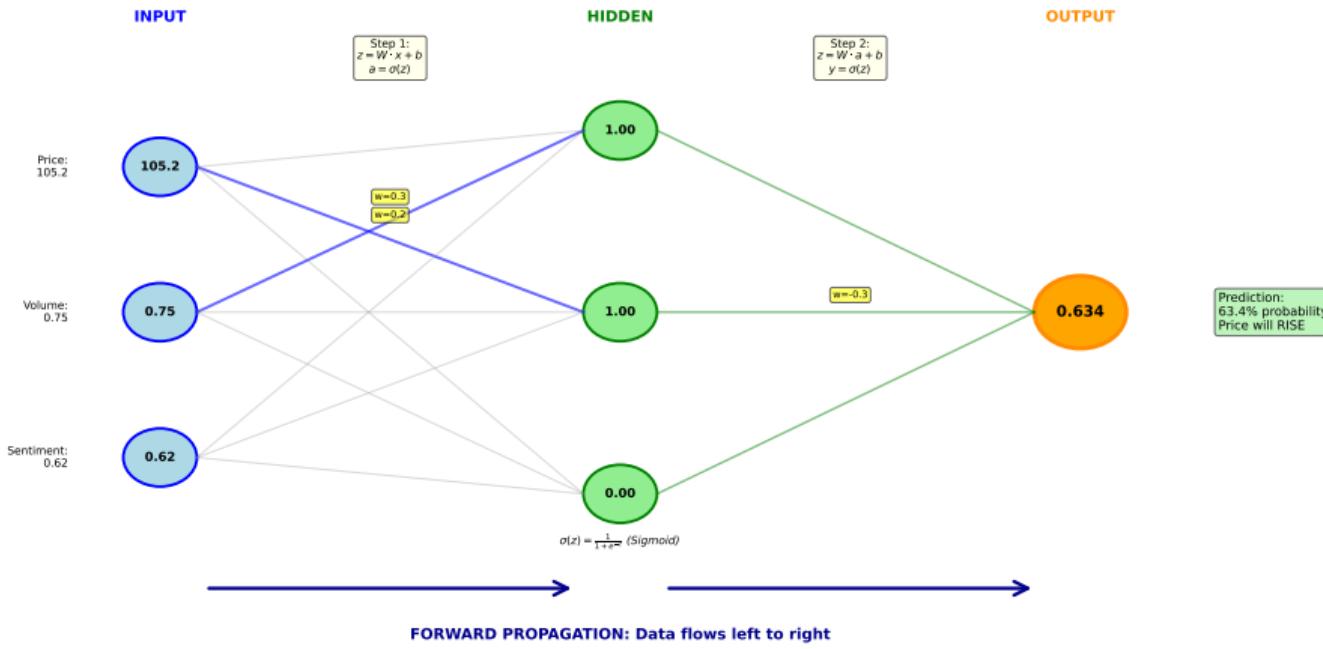
- 74.2% confidence price rises
- $y > 0.5$ : **BUY**
- $y < 0.5$ : **SELL**

---

Next: See forward propagation with actual numbers and calculations

# Forward Propagation: Detailed Example

## Forward Propagation: Making a Prediction



Observe: How do the hidden layer values combine to produce the final 0.742 output?

### Think – Pair – Share

*For your industry, what would be the 'inputs' and 'outputs' of a useful neural network?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 4: Learning Process

How networks learn from mistakes

[1] – [2] – [3] – **[4]** – [5]

## Learning Steps

1. Predict with random weights

2. Measure error:

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

3. Adjust weights:

$$w_{new} = w_{old} - \eta \nabla L$$

4. Repeat until convergence

## Example

Predicted: 55% rise, Actual: fell

$$\text{Error: } (0 - 0.55)^2 = 0.30$$

## Learning:

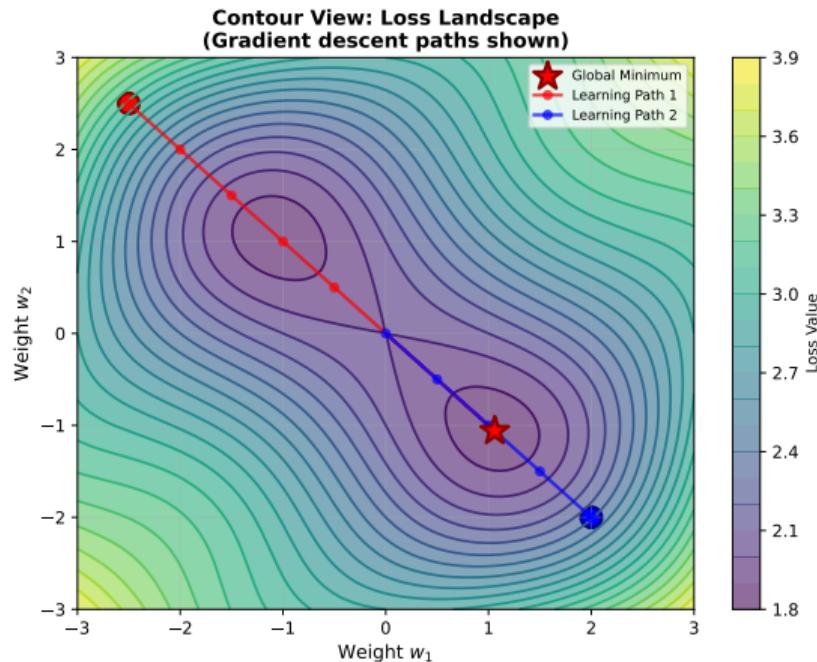
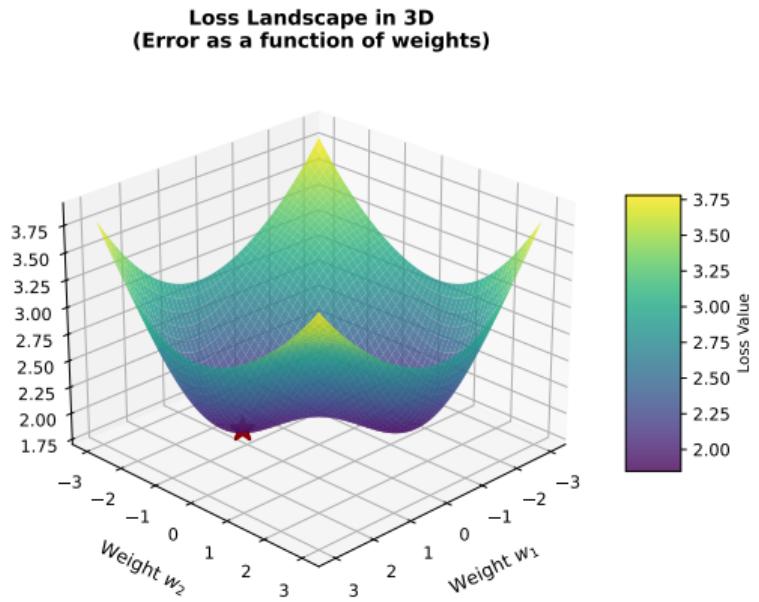
- Calculate gradient direction
- Move weights to reduce error

*Like a trader learning from mistakes*

---

Next: Visualize the loss landscape that we're trying to navigate

# Loss Landscape: The Error Surface



Goal: Find the weights that minimize the loss  
(The red star shows the optimal solution)

Observe: What happens if we start from different random initial weights?

## Algorithm

1. Calculate gradient (slope)
2. Step opposite direction
3. Repeat until convergence

## Learning Rate Trade-offs

- **Too small:** Slow
- **Too large:** Unstable
- **Just right:** Steady

## Business Analogy

Like a trader learning:

- Fast learning from obvious patterns
- Steady fine-tuning
- Convergence to optimal rules

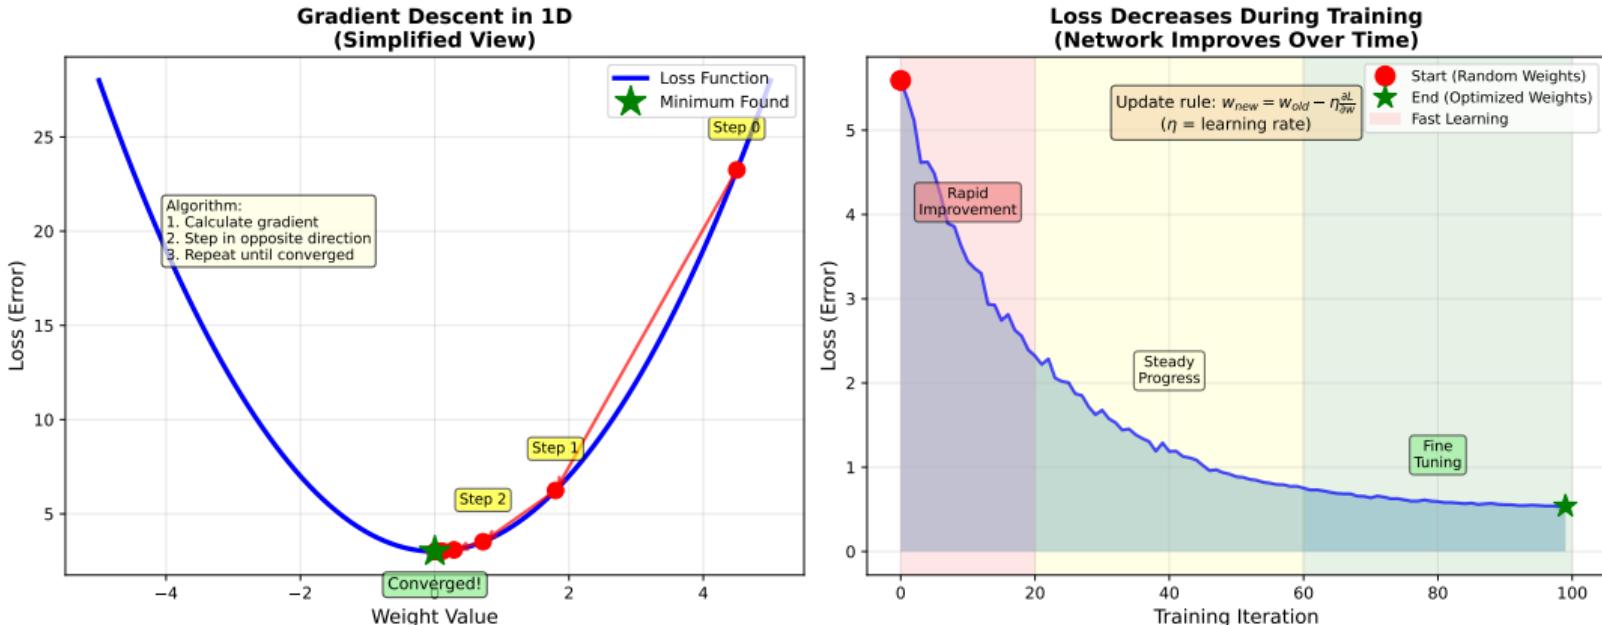
Gradient shows fastest error reduction

---

Next: See how loss decreases over training iterations

# Gradient Descent: Optimization in Action

## Gradient Descent: Learning by Stepping Downhill



Observe: How does the step size (learning rate) affect how quickly we reach the minimum?

### Think – Pair – Share

*How is gradient descent similar to how businesses optimize through trial and error?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 5: Application

Putting it all together with market prediction

[1] – [2] – [3] – [4] – **[5]**

# Putting It Together: Market Prediction Case Study

## Business Application

- **Goal:** Predict price direction
- **Data:** 60 days market data

## Input Features

1. Stock Price
2. Trading Volume
3. Market Sentiment
4. Volatility Index

## Target Variable

Binary: 1 = up, 0 = down  
Network outputs:  $p(\text{rise})$

## Setup

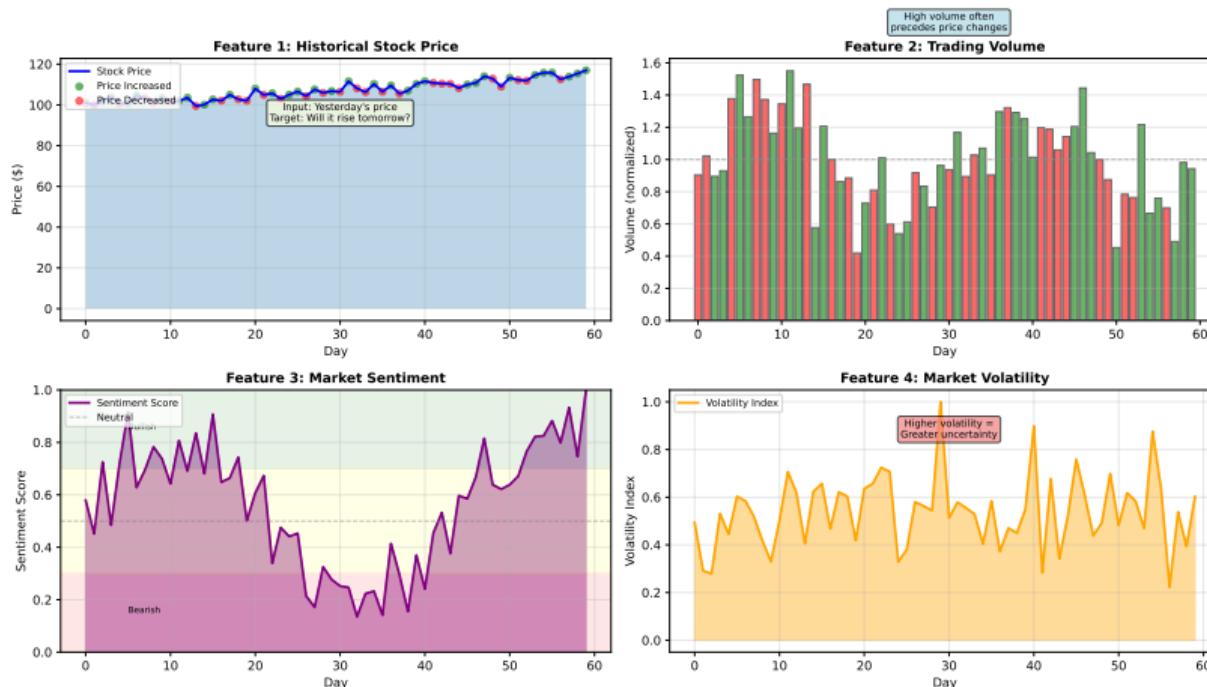
- Train: 45 days
- Test: 15 days
- Network: 4-6-1

---

Next: See the actual market data used for training

# Market Data: Input Features for Neural Network

Market Data: Input Features for Neural Network



Neural Network Input: All 4 features for each day | Output: Probability of price increase tomorrow

Observe: Which features seem most correlated with the price direction markers?

## The Experiment

- Before: Random weights (coin flip)
- After: Learned weights
- Test: 30 days unseen data

## Results

- **Before:** 50% accuracy
- **After:** 70% accuracy
- **Gain:** +20 points

## What Network Learned

- Volume + price + sentiment patterns
- Volatility indicates uncertainty
- Sentiment confirms trends

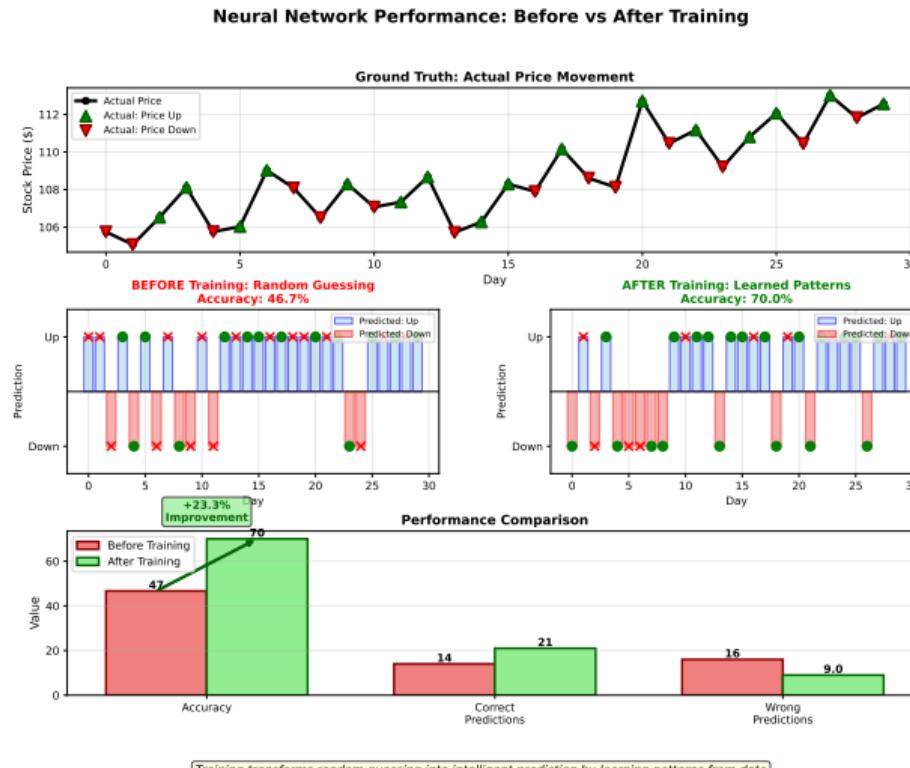
*Discovered from data alone!*

70% is good for markets (100% impossible)

---

Next: See detailed before/after comparison with prediction accuracy

# Prediction Results: Before vs After Training



Observe: Where does the trained model still make errors? What might explain these?

### Think – Pair – Share

*What data would you need to predict customer behavior in your domain?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Our Journey

1. Biological inspiration
2. Mathematical model
3. Activation functions
4. Network architecture
5. Forward propagation

## Learning Process

6. Loss functions
7. Gradient descent
8. Training iterations
9. Pattern discovery
10. Improved predictions

## Key Takeaway

Neural networks **learn patterns from data** rather than following explicit rules. This enables them to discover complex relationships that would be impossible to program manually.

---

Each concept was paired with a visualization to reinforce understanding

# When to Use Neural Networks

## Use Neural Networks When

- Large dataset (thousands+ examples)
- Complex patterns
- Difficult to specify rules
- Pattern recognition tasks
- Black-box acceptable

## Applications

Churn, fraud, recommendations, images, NLP

## Do NOT Use When

- Small dataset
- Simple relationships
- Need interpretability
- Rules are known
- Real-time constraints

## Alternatives

Regression, decision trees, expert systems

---

Choose the right tool - neural networks are powerful but not always appropriate

## Technical Limitations

- Data hungry
- Black box decisions
- Overfitting risk
- No guarantees
- Computational cost

## Ethical Concerns

- **Fairness:** Biased data leads to biased predictions
- **Transparency:** GDPR requires explanations
- **Accountability:** Who is responsible?
- **Impact:** Job displacement, market stability

*With great predictive power comes great responsibility!*

---

Always consider ethical implications before deploying AI systems

### Output Layer Gradient

$$\frac{\partial L}{\partial w^{(2)}} = (\hat{y} - y) \cdot \sigma'(z) \cdot a$$

### Hidden Layer Gradient

$$\frac{\partial L}{\partial w^{(1)}} = \delta^{(2)} \cdot w^{(2)} \cdot \sigma'(z^{(1)}) \cdot x$$

### Loss Functions

#### MSE (Regression):

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

#### Cross-Entropy (Classification):

$$L = - \sum [y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

---

Backpropagation efficiently computes how each weight contributed to the error

## Practical Tips

- Start simple (baseline first)
- Feature engineering matters
- Avoid overfitting (validation, dropout)
- Tune hyperparameters
- Monitor training curves

## Books

Goodfellow (Deep Learning), Nielsen, Geron

## Courses

- Andrew Ng (Coursera)
- Fast.ai
- MIT 6.S191

## Tools

PyTorch, TensorFlow, scikit-learn

## Practice

Kaggle, Yahoo Finance, UCI Repository

---

**Best way to learn: Build real projects with real data!**

# Appendix: Practice Problem for Business Students

## Design Challenge

You are a data scientist at a retail company.

**Problem:** Predict customer churn

**Data Available:**

- Demographics
- Purchase history
- Service interactions
- Website engagement

## Your Tasks

1. Design network architecture
2. Select input features
3. Choose activation functions
4. Select loss function
5. Define evaluation metrics
6. Identify ethical concerns
7. Plan stakeholder explanation

---

Discuss in groups - there's no single right answer!