

# Introduction to Neural Networks

## From Brain to Business: How Machines Learn to Predict

### Neural Networks for Business Applications

November 24, 2025

#### Learning Objectives

- **Explain** how biological neurons inspire artificial neural networks
- **Calculate** the output of an artificial neuron given inputs and weights
- **Design** a simple multilayer network architecture
- **Trace** information flow through forward propagation
- **Describe** how networks learn by minimizing prediction errors
- **Evaluate** when neural networks are appropriate for business

# The Prediction Challenge: Can We Predict Markets?

## The Business Question

Can we predict if a stock price will rise or fall tomorrow?

- Traditional: Statistical analysis, expert intuition
- Challenge: Markets are **complex, non-linear**
- Many factors: price, volume, sentiment, volatility

## The Limitation

Rule-based systems cannot capture all interactions

## Why This Matters

- Better investment decisions
- Risk management
- Portfolio optimization
- Automated trading strategies

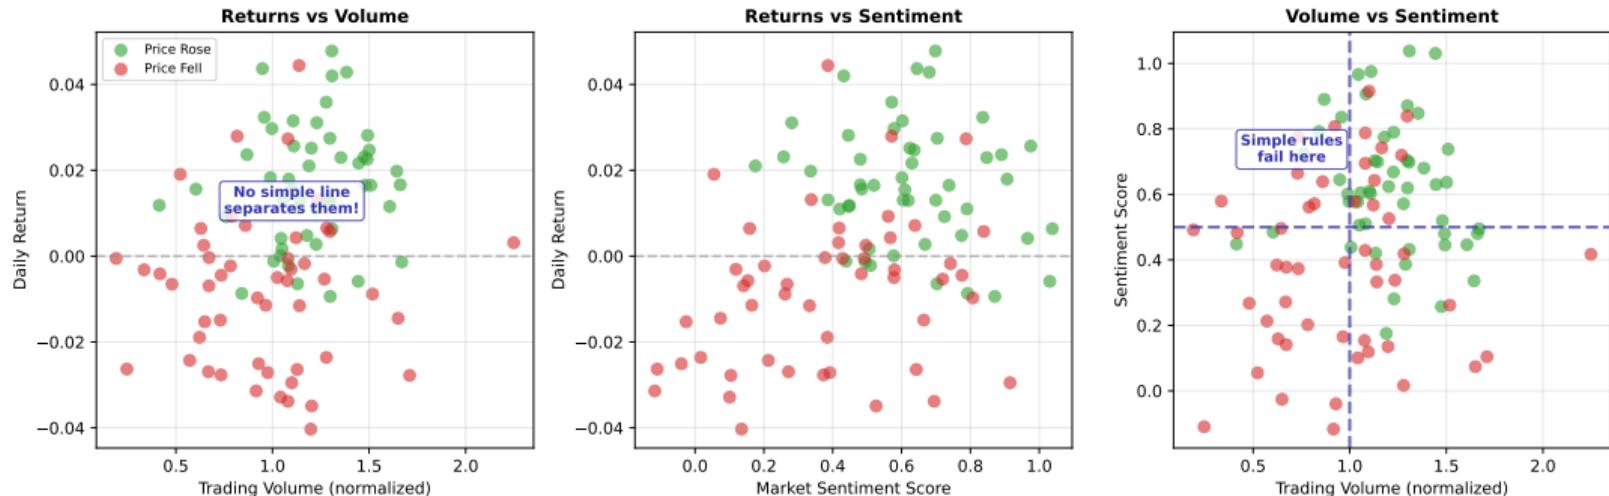
## What We Need

A system that learns patterns from data, not explicit rules

---

Our journey begins with understanding how nature solved similar prediction problems

# Why Simple Rules Fail: Market Data Complexity



Observe: Can you draw a single line that separates green (up) from red (down) in any panel?

## System Requirements

1. Process multiple inputs simultaneously
2. Learn patterns from historical data
3. Handle **non-linear** relationships
4. Improve predictions over time
5. Generalize to new conditions

## Key Insight

We need a system that learns, not one we program

## Inspiration from Nature

The human brain solves complex pattern recognition every day

## Brain Capabilities

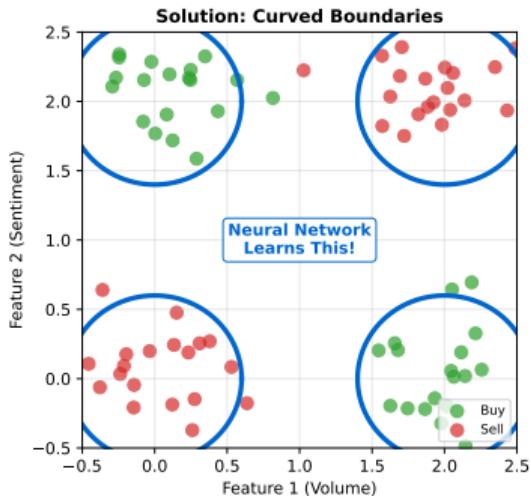
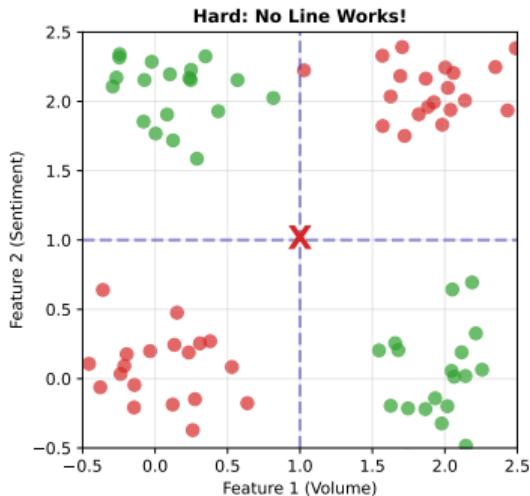
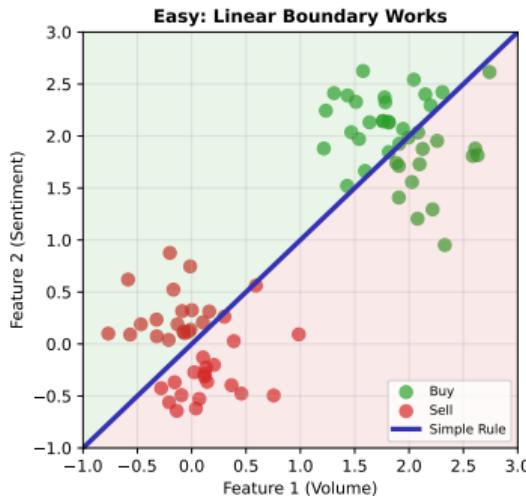
- Processes millions of inputs
- Learns from experience
- Handles ambiguity
- Generalizes to new situations

*Can we mimic this for business predictions?*

---

Next: Understanding biological neurons as the foundation

# The Goal: Learn Complex Decision Boundaries



Observe: The rightmost panel shows what neural networks can learn - curved boundaries that adapt to data

## Part 1: Foundations

From biological neurons to artificial intelligence

*Let's begin with the inspiration from nature*

[1] – [2] – [3] – [4] – [5]

## Biological Neuron Structure

- **Dendrites:** Receive signals
- **Soma:** Integrates weighted signals
- **Axon:** Transmits output
- **Synapses:** Variable connection strengths

## Key Principle

Fire when weighted sum exceeds threshold

## Business AI Insights

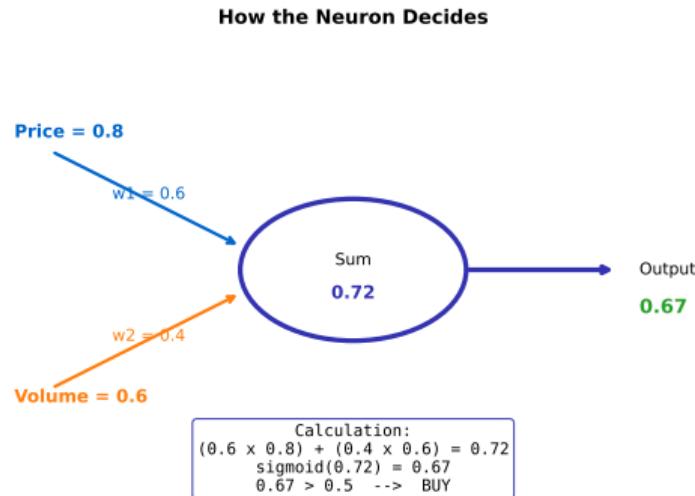
1. Multiple inputs combined
2. Weighted connections (importance)
3. Non-linear activation (thresholds)
4. Layered processing (abstraction)

*Mathematical models can learn the same way!*

---

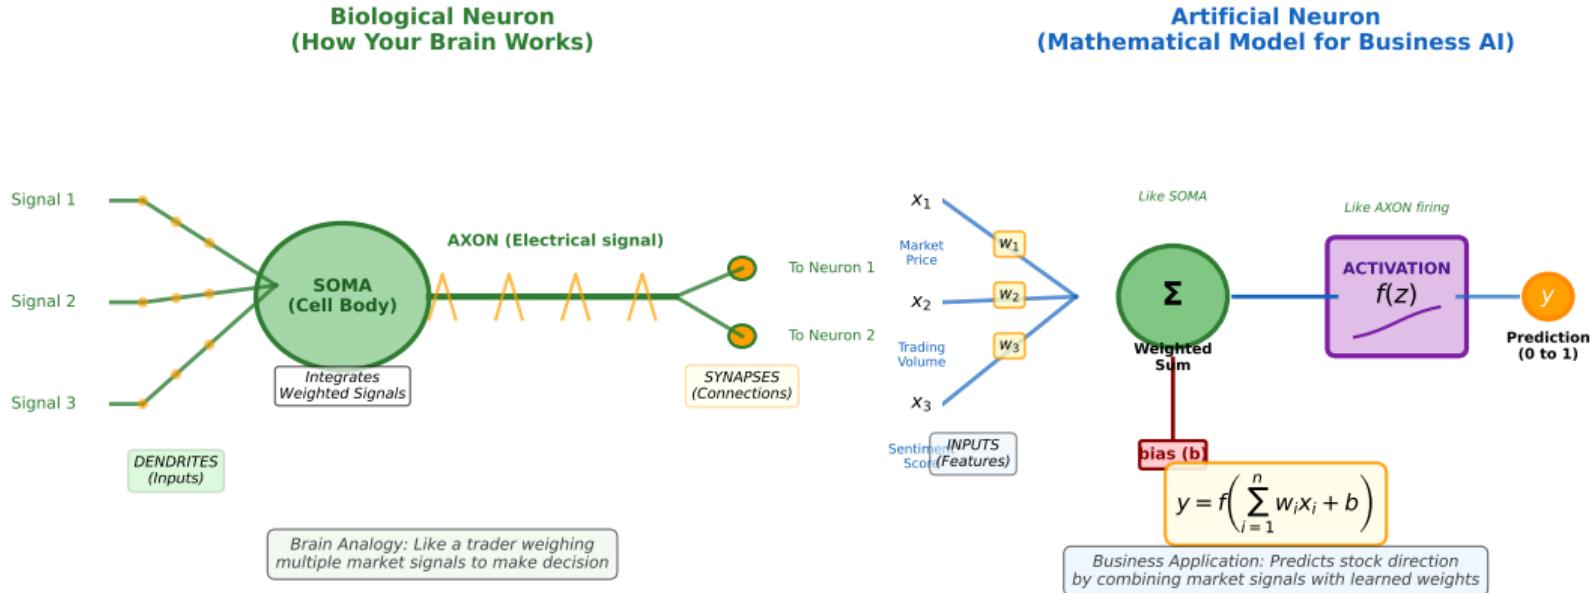
Next: See the visual comparison of biological vs artificial neurons

# From Concept to Computation: Neuron as Decision Maker



**Observe:** The decision boundary (purple line) divides the space into BUY and SELL zones based on weighted inputs

## From Biological Intelligence to Business AI



Observe: Which biological components map directly to mathematical operations?

## Step 1: Weighted Sum

$$z = \sum_{i=1}^n w_i x_i + b$$

- $x_i$ : Inputs (market data)
- $w_i$ : Weights (learned)
- $b$ : Bias (baseline)

## Step 2: Activation

$$y = \sigma(z) = \frac{1}{1 + e^{-z}}$$

- Adds non-linearity
- Output: probability (0 to 1)
- Mimics neuron firing

Complete:  $y = \sigma(\sum w_i x_i + b)$

---

Next: See a concrete example with real market numbers

# Practice: Calculate a Neuron's Output

## Given Values

- Inputs:  $x_1 = 1.2, x_2 = 0.8$
- Weights:  $w_1 = 0.3, w_2 = 0.5$
- Bias:  $b = -0.2$

## Step 1: Weighted Sum

$$z = w_1x_1 + w_2x_2 + b$$

$$z = (0.3)(1.2) + (0.5)(0.8) + (-0.2)$$

$$z = 0.36 + 0.40 - 0.20 = \mathbf{0.56}$$

## Step 2: Apply Sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-0.56}}$$

$$\sigma(0.56) = \frac{1}{1+0.571} = \frac{1}{1.571} = \mathbf{0.636}$$

## Interpretation

63.6% confidence: price will rise

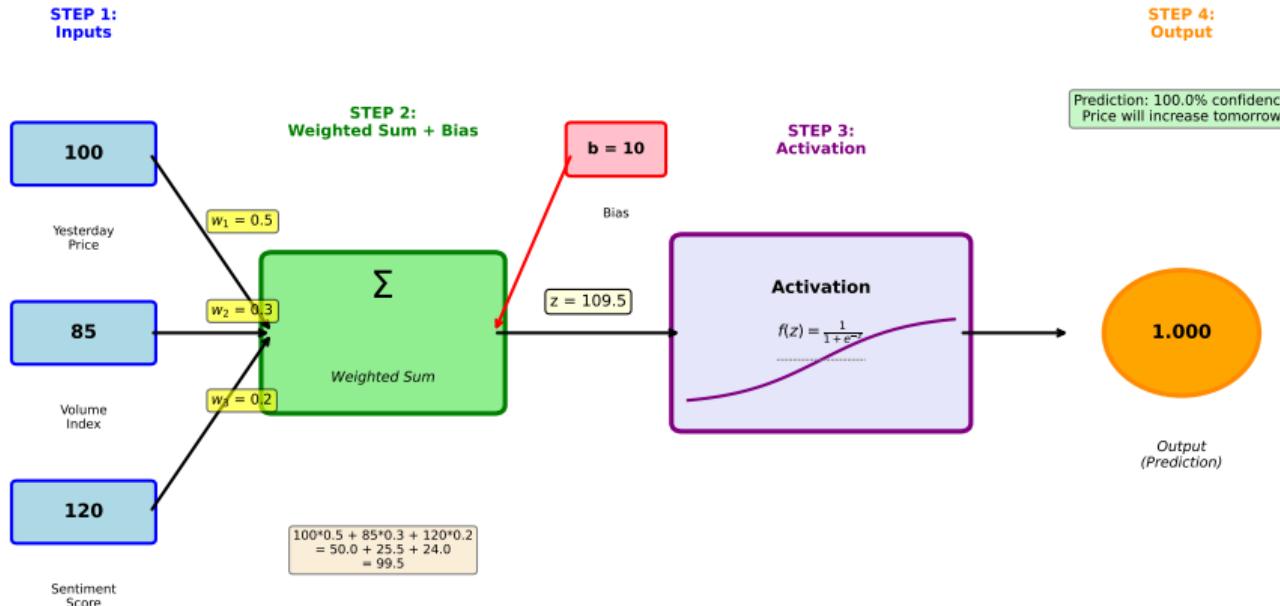
**Your Turn:** What if  $w_1 = 0.6$ ?

---

Work through this calculation – it's the foundation of all neural network predictions

# Single Neuron Computation: Step-by-Step Example

## How a Neuron Computes: Step-by-Step



Observe: How would changing the weights affect the final output probability?

### Think – Pair – Share

*What other business processes might benefit from 'learning from data' instead of following explicit rules?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 2: Building Blocks

Activation functions and their role in learning

*Now that we understand neurons, let's explore what makes them powerful*

[1] – [2] – [3] – [4] – [5]

# Activation Functions: Why Non-Linearity Matters

## The Problem

Without activation functions:

- Networks = linear regression
- Cannot learn complex patterns

## Three Common Functions

- **Sigmoid:** (0,1) for probabilities
- **ReLU:** Fast, efficient
- **Tanh:** Zero-centered (-1,1)

## Business Non-Linearity

1. Diminishing returns
2. Threshold effects
3. Saturation points
4. Network effects

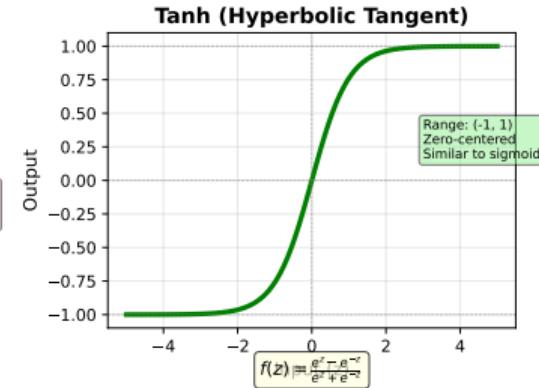
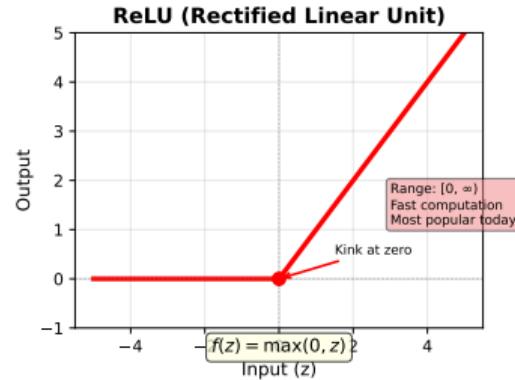
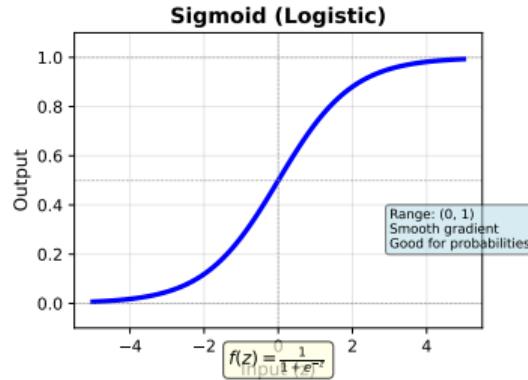
*Activation functions capture these patterns!*

---

Next: Visual comparison of these three activation functions

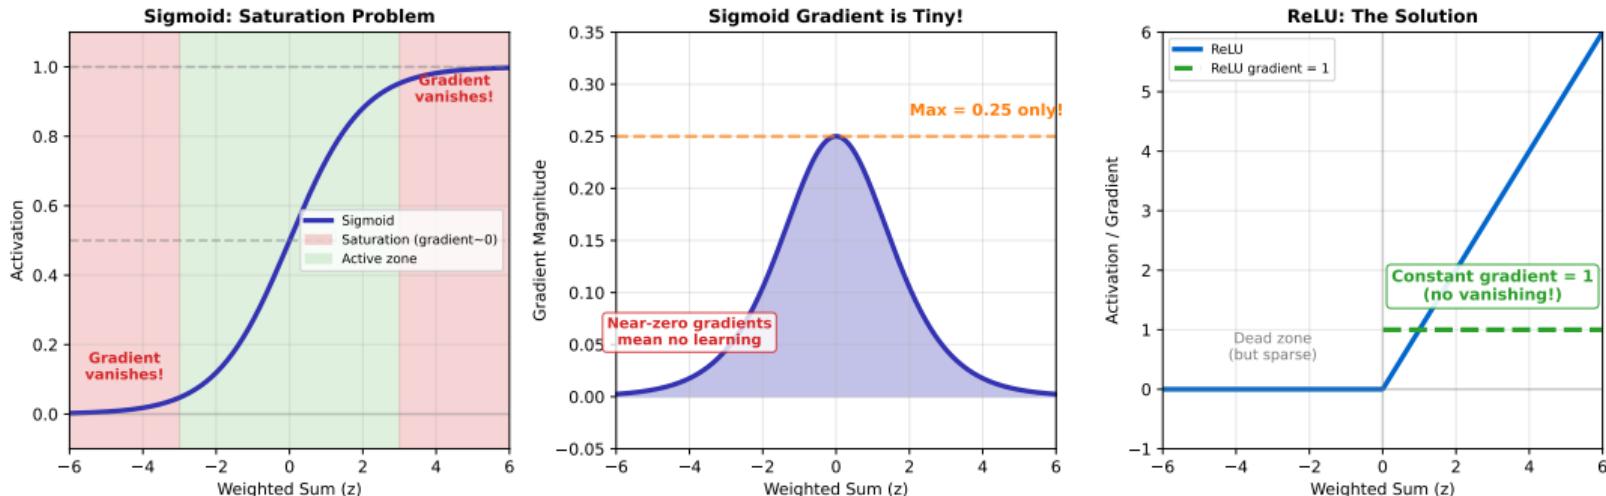
# Activation Functions: Visual Comparison

## Activation Functions: Adding Non-Linearity



Observe: Where does each function's output change most rapidly? Why does this matter?

# Advanced: The Vanishing Gradient Problem



Advanced insight: Sigmoid's tiny gradients in saturation zones slow learning – ReLU solves this in deep networks

# The Limitation: Why One Neuron Is Not Enough

## What One Neuron Can Do

- Single straight decision boundary
- Separate linearly separable patterns
- Simple rules only

**Analogy:** One rule for decisions

## What One Neuron Cannot Do

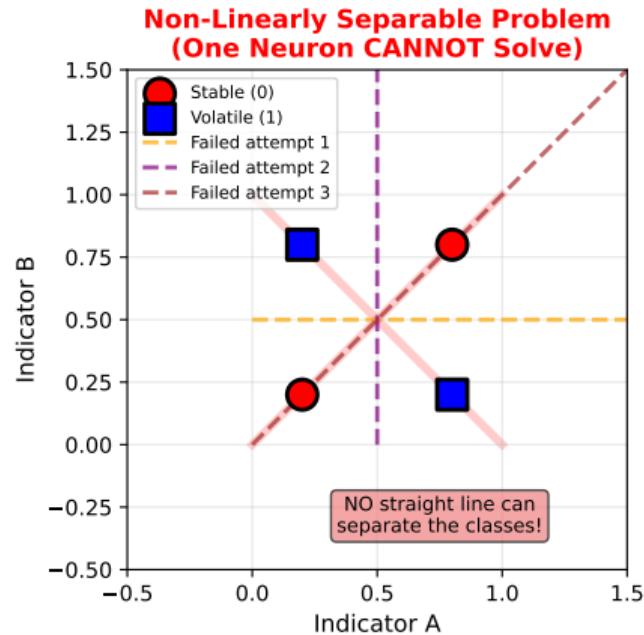
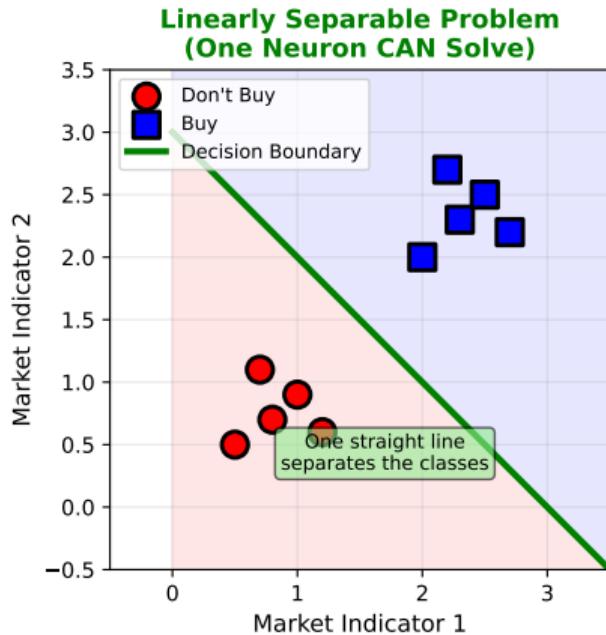
- Complex, curved boundaries
- XOR-like patterns
- Real-world market interactions

**Solution: Multiple Layers!**

---

Next: See the XOR problem that proves one neuron's limitation

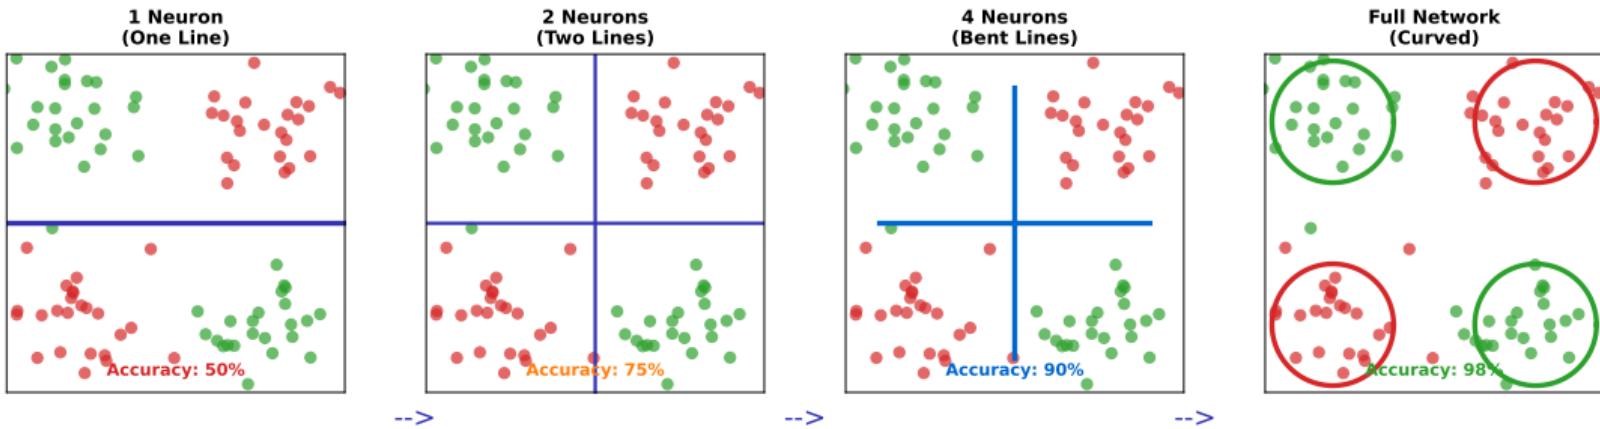
## Why One Neuron Is Not Enough



**Solution: Use Multiple Layers (Hidden Layers) to Create Non-Linear Decision Boundaries**

Observe: Why is it impossible to draw a single straight line separating orange from blue?

## Solution: How Adding Neurons Creates Curved Boundaries



**Key insight:** More neurons = more flexibility. Each neuron adds a decision line; combined, they form complex shapes

### Think – Pair – Share

*Can you think of a business metric that shows diminishing returns or threshold effects?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 3: Network Architecture

Building layers of intelligence

*With building blocks ready, let's construct full networks*

[1] – [2] – **[3]** – [4] – [5]

## Multi-Layer Architecture

- **Input:** Raw features (no computation)
- **Hidden:** Pattern detection
- **Output:** Final prediction

**Result:** Buy/Sell decision

## Hierarchical Learning

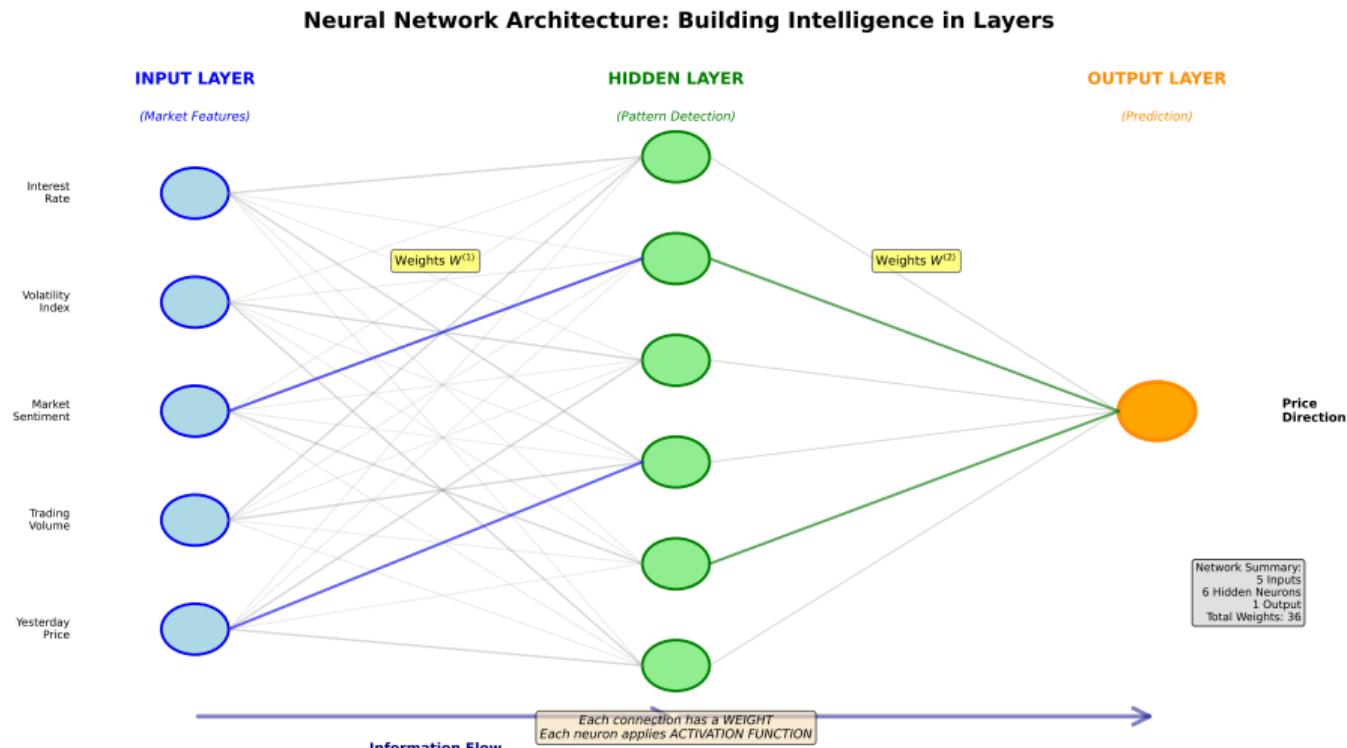
- **Layer 1:** Simple patterns
- **Layer 2:** Complex patterns
- **Layer 3:** Strategic decisions

Each layer builds on previous abstractions

---

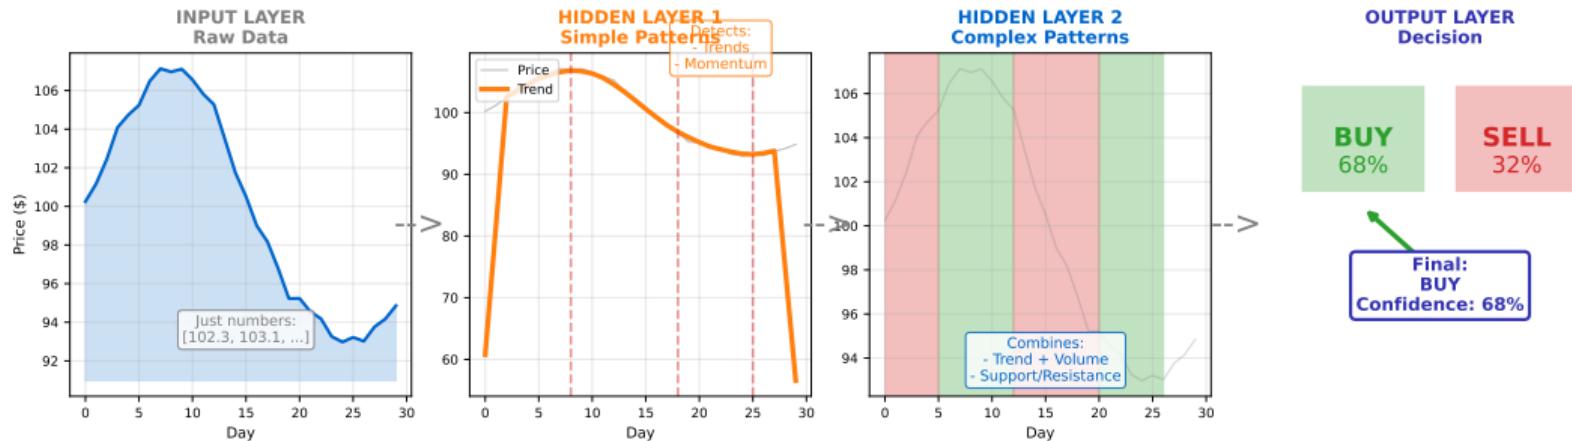
Next: See the full network architecture with all connections

# Neural Network Architecture Diagram



Observe: Count the connections. Why are there 36 weights total?

# What Each Layer “Sees”: Feature Hierarchy



Observe: Raw data transforms through layers into increasingly abstract representations until a decision emerges

# Forward Propagation: How Networks Make Predictions

## The Forward Pass

1. **Input:** Feed market features
2. **Hidden:**  $a = \sigma(Wx + b)$
3. **Output:**  $y = \sigma(Wa + b)$

All neurons compute in parallel!

## Example

Input: price=105.2, volume=0.75  
Output:  $y = 0.742$

### Interpretation:

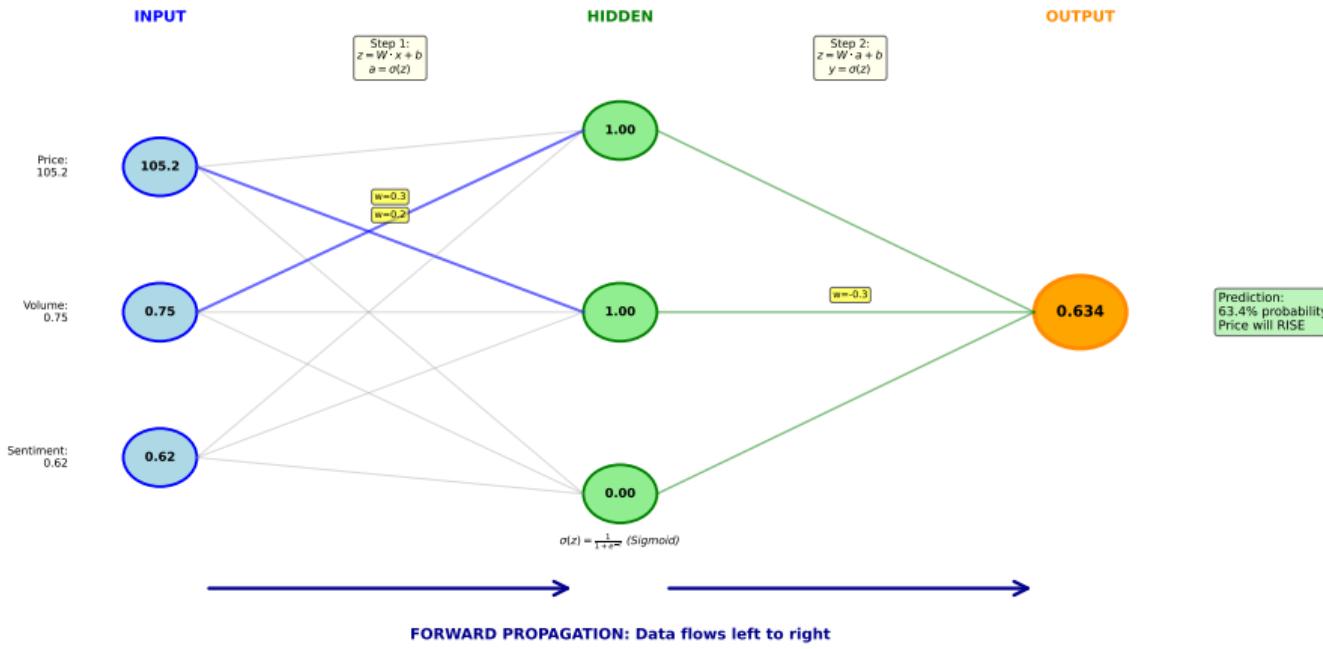
- 74.2% confidence price rises
- $y > 0.5$ : **BUY**
- $y < 0.5$ : **SELL**

---

Next: See forward propagation with actual numbers and calculations

# Forward Propagation: Detailed Example

## Forward Propagation: Making a Prediction



Observe: How do the hidden layer values combine to produce the final 0.742 output?

### Think – Pair – Share

*For your industry, what would be the 'inputs' and 'outputs' of a useful neural network?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 4: Learning Process

How networks learn from mistakes

*We can make predictions – now let's learn how to improve them*

[1] – [2] – [3] – **[4]** – [5]

## Learning Steps

1. Predict with random weights

2. Measure error:

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

3. Adjust weights:

$$w_{new} = w_{old} - \eta \nabla L$$

4. Repeat until convergence

## Example

Predicted: 55% rise, Actual: fell

$$\text{Error: } (0 - 0.55)^2 = 0.30$$

## Learning:

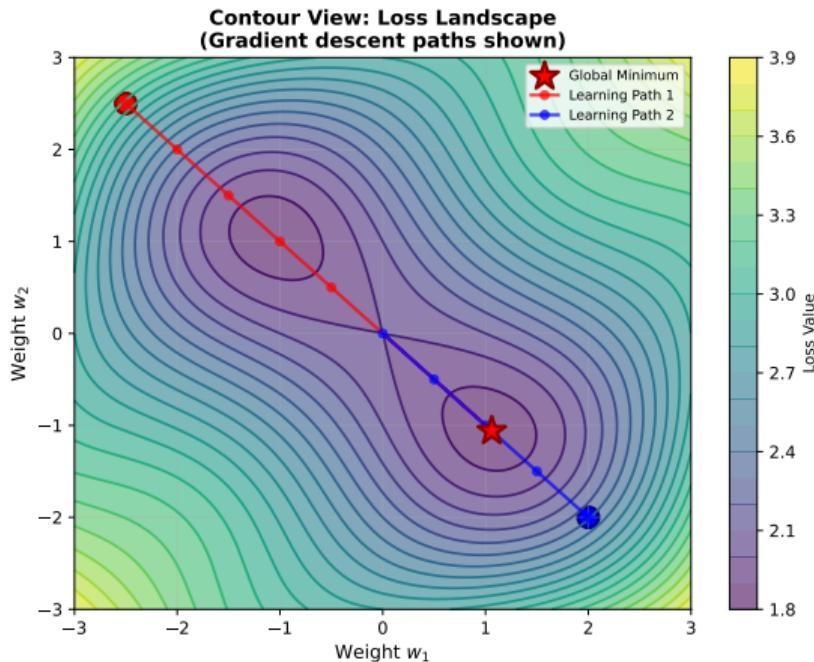
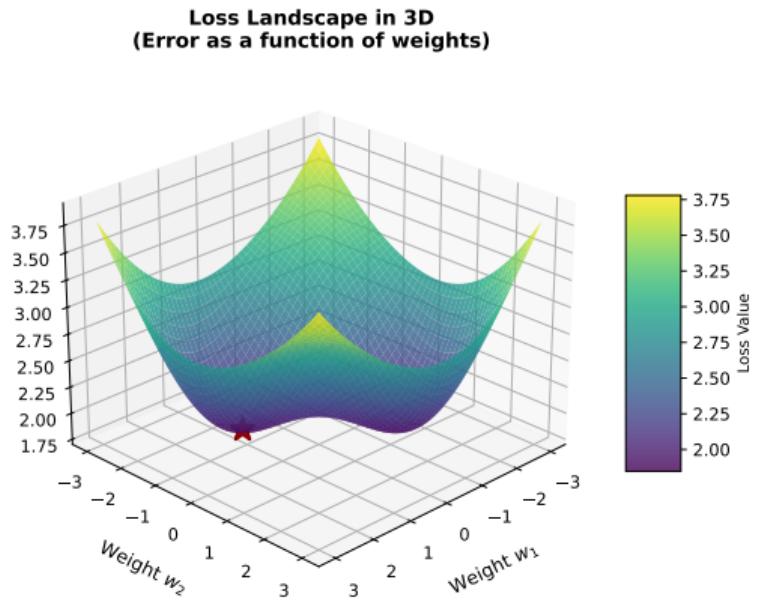
- Calculate gradient direction
- Move weights to reduce error

*Like a trader learning from mistakes*

---

Next: Visualize the loss landscape that we're trying to navigate

# Loss Landscape: The Error Surface



Goal: Find the weights that minimize the loss  
(The red star shows the optimal solution)

Observe: What happens if we start from different random initial weights?

## Algorithm

1. Calculate gradient (slope)
2. Step opposite direction
3. Repeat until convergence

## Learning Rate Trade-offs

- **Too small:** Slow
- **Too large:** Unstable
- **Just right:** Steady

## Business Analogy

Like a trader learning:

- Fast learning from obvious patterns
- Steady fine-tuning
- Convergence to optimal rules

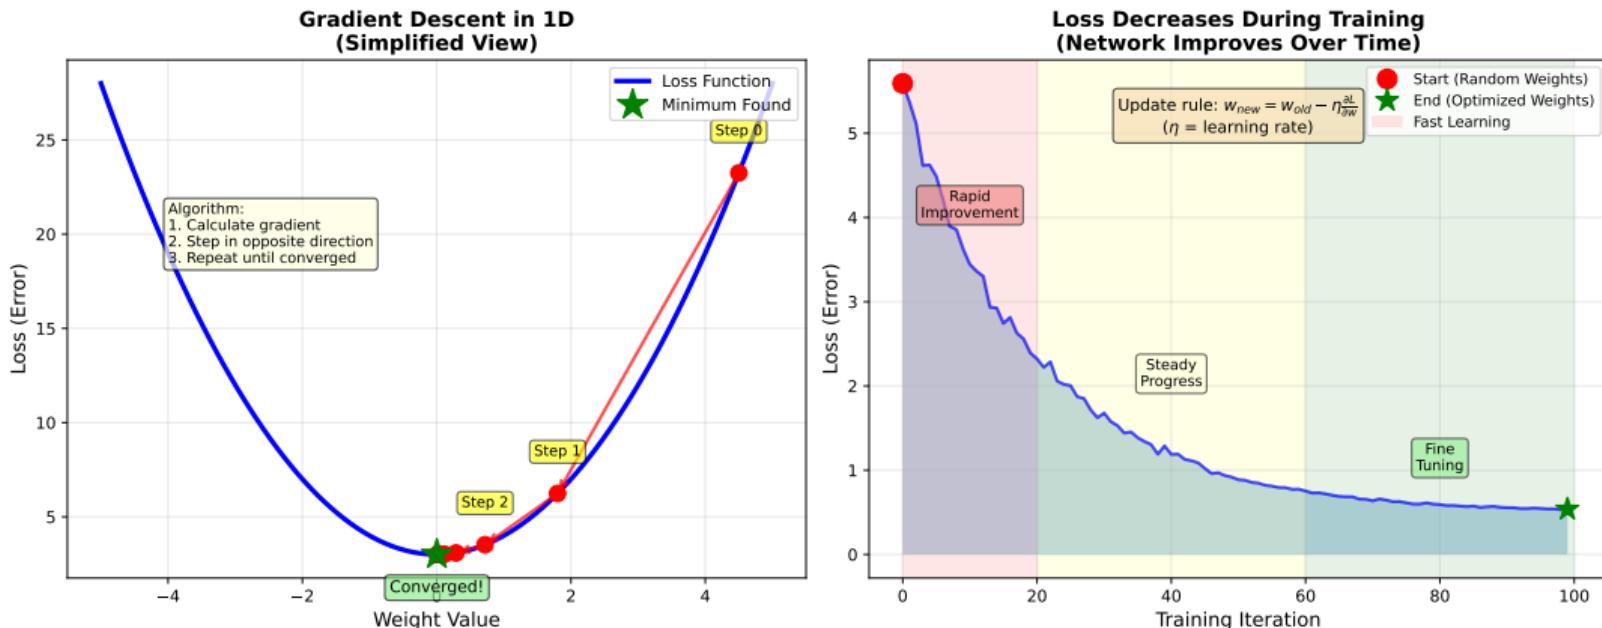
Gradient shows fastest error reduction

---

Next: See how loss decreases over training iterations

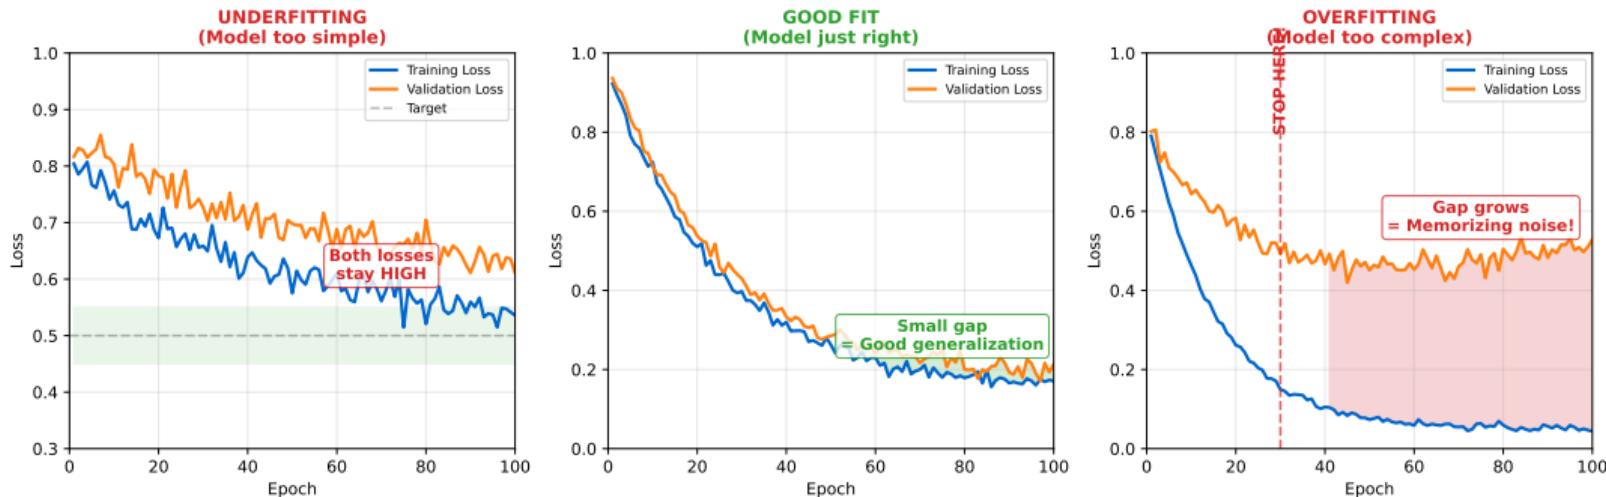
# Gradient Descent: Optimization in Action

## Gradient Descent: Learning by Stepping Downhill



Observe: How does the step size (learning rate) affect how quickly we reach the minimum?

# Critical Concept: Overfitting vs Underfitting



**Key practical skill:** Watch for diverging training/validation loss – that's when to stop training!

### Think – Pair – Share

*How is gradient descent similar to how businesses optimize through trial and error?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

## Part 5: Application

Putting it all together with market prediction

*Theory complete – let's apply everything to a real case*

[1] – [2] – [3] – [4] – **[5]**

# Putting It Together: Market Prediction Case Study

## Business Application

- **Goal:** Predict price direction
- **Data:** 60 days market data

## Input Features

1. Stock Price
2. Trading Volume
3. Market Sentiment
4. Volatility Index

## Target Variable

Binary: 1 = up, 0 = down  
Network outputs:  $p(\text{rise})$

## Setup

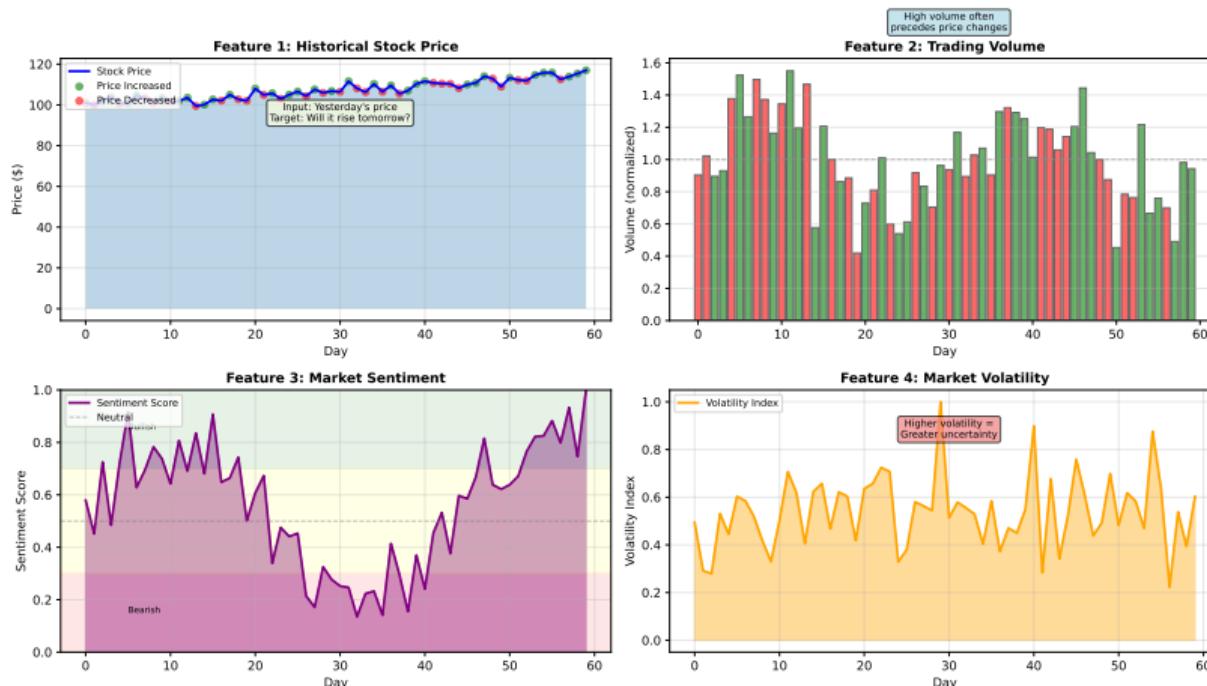
- Train: 45 days
- Test: 15 days
- Network: 4-6-1

---

Next: See the actual market data used for training

# Market Data: Input Features for Neural Network

Market Data: Input Features for Neural Network



Observe: Which features seem most correlated with the price direction markers?

# Training Results: Before vs After

## The Experiment

- Before: Random weights (coin flip)
- After: Learned weights
- Test: 30 days unseen data

## Results

- **Before:** 50% accuracy
- **After:** 70% accuracy
- **Gain:** +20 points

## What Network Learned

- Volume + price + sentiment patterns
- Volatility indicates uncertainty
- Sentiment confirms trends

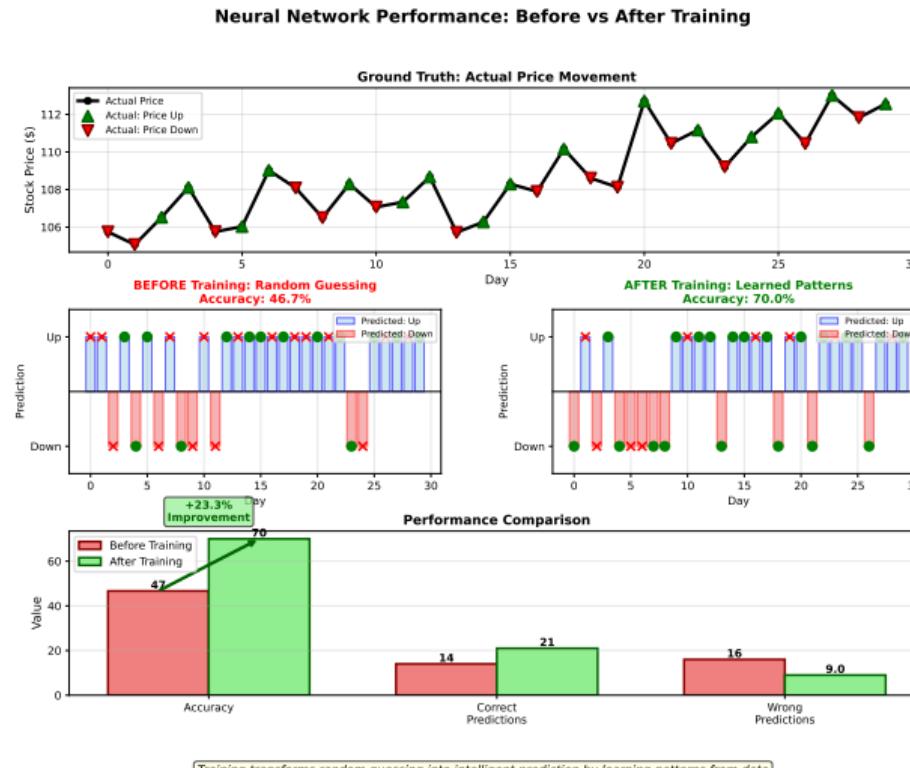
*Discovered from data alone!*

70% is good for markets (100% impossible)

---

Next: See detailed before/after comparison with prediction accuracy

# Prediction Results: Before vs After Training



Observe: Where does the trained model still make errors? What might explain these?

# Understanding Model Performance: Confusion Matrix

		Trading Confusion Matrix (100 test days)	
		False Positive	True Negative
Actual DOWN	False Positive	16	34
	True Negative	36	14
Actual UP	True Positive	36	14
	False Negative	16	34
Predicted UP		Predicted DOWN	



- Accuracy:** Overall correct predictions 70/100
- Precision:** When we say BUY, how often right? 36/52
- Recall:** Of all UP days, how many caught? 36/50
- F1 Score:** Balance of precision & recall harmonic mean

**Trading Insight:** 69% precision means ~1/3 of BUY signals are wrong!

**Business insight:** 70% accuracy means different things for trading – precision determines false BUY rate

### Think – Pair – Share

*What data would you need to predict customer behavior in your domain?*

1. Think (1 min)

Reflect individually on the question

2. Pair (2 min)

Discuss with a neighbor

3. Share (2 min)

Share insights with class

# The Business Case: Strategy Backtest Results



## PERFORMANCE COMPARISON

Metric	Buy & Hold	NN Strategy
Total Return	7.9%	<b>95.3%</b>
Sharpe Ratio	0.40	<b>3.76</b>
Max Drawdown	-16.8%	<b>-9.2%</b>
Win Rate	52%	<b>66%</b>

**Key Insight: 70% accuracy translates to significant alpha!**

(Backtest only - past performance does not guarantee future results)

**Bottom line: 70% accuracy translates to meaningful outperformance – this is why neural networks matter for business**

## Summary: Three Key Insights

### 1. Neurons Compute Weighted Sums

Each artificial neuron multiplies inputs by learned weights, adds a bias, and applies a non-linear activation function. This simple operation, repeated across layers, enables complex pattern recognition.

### 2. Networks Learn from Errors

Training uses gradient descent to minimize prediction errors. The network adjusts weights in the direction that reduces loss – like a trader learning from past mistakes.

### 3. Patterns Emerge from Data

Neural networks discover relationships we never explicitly programmed. They find what matters in the data, enabling predictions for complex, non-linear business problems.

These three principles underpin all deep learning – master them and you understand neural networks

# Quick Check: Test Your Understanding

**Q1: What does the activation function do?**

- (a) Stores the input data
- (b) **Adds non-linearity to enable complex patterns**
- (c) Calculates the learning rate

**Q2: Why do we need multiple layers?**

- (a) To make training faster
- (b) To use more data
- (c) **To learn hierarchical, complex patterns**

**Q3: What does gradient descent minimize?**

- (a) The number of neurons
- (b) **The prediction error (loss function)**
- (c) The training time

**Check Your Answers**

**Answer Key**

- Q1: (b) Non-linearity
- Q2: (c) Hierarchical patterns
- Q3: (b) Loss/error

**Scoring**

- 3/3: Excellent grasp!
- 2/3: Review that topic
- 1/3: Revisit core concepts

---

If any answer surprised you, go back and review that section

# When to Use Neural Networks

## Use Neural Networks When

- Large dataset (thousands+ examples)
- Complex patterns
- Difficult to specify rules
- Pattern recognition tasks
- Black-box acceptable

## Applications

Churn, fraud, recommendations, images, NLP

## Do NOT Use When

- Small dataset
- Simple relationships
- Need interpretability
- Rules are known
- Real-time constraints

## Alternatives

Regression, decision trees, expert systems

---

Choose the right tool - neural networks are powerful but not always appropriate

# Important Limitations & Ethical Responsibilities

## Technical Limitations

- Data hungry
- Black box decisions
- Overfitting risk
- No guarantees
- Computational cost

## Ethical Concerns

- **Fairness:** Biased data leads to biased predictions
- **Transparency:** GDPR requires explanations
- **Accountability:** Who is responsible?
- **Impact:** Job displacement, market stability

*With great predictive power comes great responsibility!*

---

Always consider ethical implications before deploying AI systems

### Output Layer Gradient

$$\frac{\partial L}{\partial w^{(2)}} = (\hat{y} - y) \cdot \sigma'(z) \cdot a$$

### Hidden Layer Gradient

$$\frac{\partial L}{\partial w^{(1)}} = \delta^{(2)} \cdot w^{(2)} \cdot \sigma'(z^{(1)}) \cdot x$$

### Loss Functions

#### MSE (Regression):

$$L = \frac{1}{n} \sum (y - \hat{y})^2$$

#### Cross-Entropy (Classification):

$$L = - \sum [y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

---

Backpropagation efficiently computes how each weight contributed to the error

## Practical Tips

- Start simple (baseline first)
- Feature engineering matters
- Avoid overfitting (validation, dropout)
- Tune hyperparameters
- Monitor training curves

## Books

Goodfellow (Deep Learning), Nielsen, Geron

## Courses

- Andrew Ng (Coursera)
- Fast.ai
- MIT 6.S191

## Tools

PyTorch, TensorFlow, scikit-learn

## Practice

Kaggle, Yahoo Finance, UCI Repository

---

**Best way to learn: Build real projects with real data!**

# Appendix: Practice Problem for Business Students

## Design Challenge

You are a data scientist at a retail company.

**Problem:** Predict customer churn

**Data Available:**

- Demographics
- Purchase history
- Service interactions
- Website engagement

## Your Tasks

1. Design network architecture
2. Select input features
3. Choose activation functions
4. Select loss function
5. Define evaluation metrics
6. Identify ethical concerns
7. Plan stakeholder explanation

---

Discuss in groups - there's no single right answer!