

# Mathematical Appendix

## Complete Derivations for Neural Networks

Neural Networks for Finance

BSc Lecture Series

November 26, 2025

---

**Full mathematical derivations for all four modules**

## A.1 Perceptron Convergence Theorem

### Theorem (Novikoff, 1962)

If the training data  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  is **linearly separable**, the perceptron learning algorithm converges in a finite number of updates.

#### What This Means:

- The algorithm will find a separating hyperplane
- It will stop updating weights (no more mistakes)
- Convergence is **guaranteed**, not probabilistic

#### What It Doesn't Say:

- Nothing about non-separable data
- Doesn't guarantee the "best" hyperplane
- Number of steps can be large

---

#### Referenced in Module 1

## A.2 Setup and Assumptions

**Data:** Training set  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  where  $y_i \in \{-1, +1\}$

### Assumption 1: Linear Separability

There exists  $\mathbf{w}^* \in \mathbb{R}^d$  with  $\|\mathbf{w}^*\| = 1$  and margin  $\gamma > 0$  such that:

$$y_i(\mathbf{w}^{*T}\mathbf{x}_i) \geq \gamma \quad \text{for all } i = 1, \dots, n$$

### Assumption 2: Bounded Data

All data points have bounded norm:

$$\|\mathbf{x}_i\| \leq R \quad \text{for all } i = 1, \dots, n$$

### Key Quantities:

- $\gamma$ : The “margin” - minimum distance from any point to the hyperplane
- $R$ : Maximum radius - the data lies in a ball of radius  $R$

---

### Defining the margin and bounded data

## A.3 Proof: Step 1 - Lower Bound on $\mathbf{w}^*{}^T \mathbf{w}^{(t)}$

**Goal:** Show the inner product  $\mathbf{w}^*{}^T \mathbf{w}^{(t)}$  grows with each mistake.

**Update Rule:** When mistake on  $(\mathbf{x}_i, y_i)$ :  $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + y_i \mathbf{x}_i$

**Computing the Inner Product:**

$$\begin{aligned}\mathbf{w}^*{}^T \mathbf{w}^{(t+1)} &= \mathbf{w}^*{}^T (\mathbf{w}^{(t)} + y_i \mathbf{x}_i) \\ &= \mathbf{w}^*{}^T \mathbf{w}^{(t)} + y_i (\mathbf{w}^*{}^T \mathbf{x}_i) \\ &\geq \mathbf{w}^*{}^T \mathbf{w}^{(t)} + \gamma \quad (\text{by separability assumption})\end{aligned}$$

**After  $t$  mistakes (starting from  $\mathbf{w}^{(0)} = \mathbf{0}$ ):**

$$\mathbf{w}^*{}^T \mathbf{w}^{(t)} \geq t\gamma$$

---

Showing the inner product grows

## A.4 Proof: Step 2 - Upper Bound on $\|\mathbf{w}^{(t)}\|^2$

**Goal:** Show the squared norm  $\|\mathbf{w}^{(t)}\|^2$  grows slowly.

**Computing the Squared Norm:**

$$\begin{aligned}\|\mathbf{w}^{(t+1)}\|^2 &= \|\mathbf{w}^{(t)} + y_i \mathbf{x}_i\|^2 \\&= \|\mathbf{w}^{(t)}\|^2 + 2y_i (\mathbf{w}^{(t)T} \mathbf{x}_i) + \|\mathbf{x}_i\|^2 \\&\leq \|\mathbf{w}^{(t)}\|^2 + 0 + R^2 \quad (\text{mistake means } y_i (\mathbf{w}^{(t)T} \mathbf{x}_i) \leq 0) \\&\leq \|\mathbf{w}^{(t)}\|^2 + R^2\end{aligned}$$

**After  $t$  mistakes (starting from  $\mathbf{w}^{(0)} = \mathbf{0}$ ):**

$$\|\mathbf{w}^{(t)}\|^2 \leq tR^2$$

---

Showing the norm is bounded

## A.5 Proof: Conclusion

**Combining the Two Bounds:**

From Step 1:  $\mathbf{w}^{*T}\mathbf{w}^{(t)} \geq t\gamma$

From Step 2:  $\|\mathbf{w}^{(t)}\|^2 \leq tR^2$ , so  $\|\mathbf{w}^{(t)}\| \leq \sqrt{t}R$

**Using Cauchy-Schwarz:**

$$t\gamma \leq \mathbf{w}^{*T}\mathbf{w}^{(t)} \leq \|\mathbf{w}^{*}\| \|\mathbf{w}^{(t)}\| = 1 \cdot \|\mathbf{w}^{(t)}\| \leq \sqrt{t}R$$

**Solving for  $t$ :**

$$t\gamma \leq \sqrt{t}R \implies \sqrt{t} \leq \frac{R}{\gamma} \implies t \leq \frac{R^2}{\gamma^2}$$

**Convergence Bound**

$$\text{Number of mistakes} \leq \left(\frac{R}{\gamma}\right)^2$$

**Maximum number of mistakes:**  $(R/\gamma)^2$

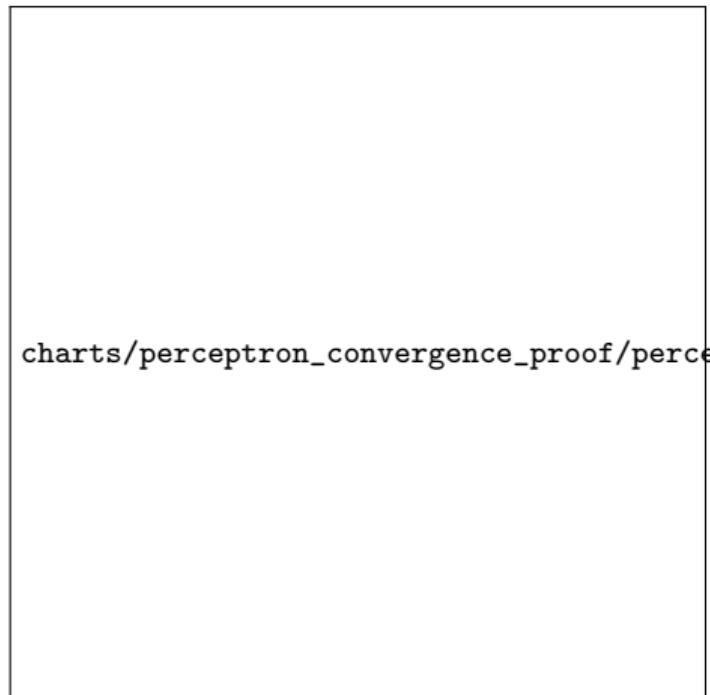
## A.6 Geometric Interpretation

### What the Bound $(R/\gamma)^2$ Tells Us:

- Large margin  $\gamma \rightarrow$  fewer mistakes
- Larger data spread  $R \rightarrow$  more mistakes
- Ratio matters, not absolute values

### Intuition:

- $\gamma =$  “wiggle room” for the hyperplane
- $R =$  how much ground to cover
- Easy problem: large  $\gamma$ , small  $R$
- Hard problem: small  $\gamma$ , large  $R$



charts/perceptron\_convergence\_proof/perceptro

Limitation: If  $\gamma = 0$  (not separable), bound is infinite  $\rightarrow$  no convergence guarantee.

What the proof means geometrically

## B.1 Backpropagation: Overview

**Goal:** Compute  $\frac{\partial \mathcal{L}}{\partial W_{jk}^{(l)}}$  and  $\frac{\partial \mathcal{L}}{\partial b_j^{(l)}}$  for all layers  $l$ .

### The Challenge:

- Loss depends on weights through many intermediate layers
- Naïve approach:  $O(n^2)$  operations per parameter
- Backprop achieves:  $O(n)$  total (same as forward pass!)

**Key Insight:** Reuse intermediate computations via the chain rule.

### What We Will Derive:

1. Error signal  $\delta^{(L)}$  at output layer
2. Recursion formula for  $\delta^{(l)}$  at hidden layers
3. Weight gradient  $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}}$
4. Bias gradient  $\frac{\partial \mathcal{L}}{\partial \mathbf{b}^{(l)}}$

---

### Referenced in Module 3

## B.2 Notation

**Network with  $L$  layers:**

- $\mathbf{h}^{(0)} = \mathbf{x}$ : Input (layer 0)
- $\mathbf{z}^{(l)} = \mathbf{W}^{(l)}\mathbf{h}^{(l-1)} + \mathbf{b}^{(l)}$ : Pre-activation (layer  $l$ )
- $\mathbf{h}^{(l)} = \phi(\mathbf{z}^{(l)})$ : Activation (layer  $l$ )
- $\hat{\mathbf{y}} = \mathbf{h}^{(L)}$ : Output (layer  $L$ )

**Dimensions:**

- $\mathbf{W}^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$ : Weight matrix
- $\mathbf{b}^{(l)} \in \mathbb{R}^{n_l}$ : Bias vector
- $n_l$ : Number of neurons in layer  $l$

**Error Signal (key quantity):**

$$\delta^{(l)} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(l)}} \in \mathbb{R}^{n_l}$$

---

Consistent notation for the derivation

## B.3 Chain Rule Review

**Univariate Chain Rule:**

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

**Multivariate Chain Rule:**

If  $L$  depends on  $z_1, \dots, z_n$ , each depending on  $w$ :

$$\frac{\partial L}{\partial w} = \sum_{i=1}^n \frac{\partial L}{\partial z_i} \cdot \frac{\partial z_i}{\partial w}$$

**Matrix Form:**

If  $\mathbf{z} = f(\mathbf{h})$  and  $L = L(\mathbf{z})$ :

$$\frac{\partial L}{\partial \mathbf{h}} = \left( \frac{\partial \mathbf{z}}{\partial \mathbf{h}} \right)^T \frac{\partial L}{\partial \mathbf{z}}$$

**Key for Backprop:** The Jacobian  $\frac{\partial \mathbf{z}}{\partial \mathbf{h}}$  connects layers.

---

The mathematical foundation of backpropagation

## B.4 Forward Pass Equations

**Layer-by-Layer Computation:**

**For each layer  $l = 1, 2, \dots, L$ :**

$$\text{Pre-activation: } \mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{h}^{(l-1)} + \mathbf{b}^{(l)}$$

$$\text{Activation: } \mathbf{h}^{(l)} = \phi^{(l)}(\mathbf{z}^{(l)})$$

**Element-wise:**

$$z_j^{(l)} = \sum_{k=1}^{n_{l-1}} W_{jk}^{(l)} h_k^{(l-1)} + b_j^{(l)}$$
$$h_j^{(l)} = \phi(z_j^{(l)})$$

**Final Output:**  $\hat{\mathbf{y}} = \mathbf{h}^{(L)}$

**Loss:**  $\mathcal{L} = \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$

---

**Computing outputs**

## B.5 Error Signal Definition

**Definition:** The error signal at layer  $l$  is:

$$\delta^{(l)} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(l)}}$$

### Why This Quantity?

- It captures “how much does the loss change if  $\mathbf{z}^{(l)}$  changes”
- All weight gradients can be expressed in terms of  $\delta^{(l)}$
- Can be computed recursively from layer to layer

### Strategy:

1. Compute  $\delta^{(L)}$  at output layer (depends on loss function)
2. Propagate backward:  $\delta^{(l)} = f(\delta^{(l+1)})$
3. Compute weight gradients:  $\frac{\partial \mathcal{L}}{\partial \mathbf{w}^{(l)}} = g(\delta^{(l)})$

---

The key quantity for backpropagation

## B.6 Output Layer Error Derivation

At the output layer  $L$ :

$$\delta^{(L)} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(L)}}$$

Using the Chain Rule:

$$\delta_j^{(L)} = \frac{\partial \mathcal{L}}{\partial z_j^{(L)}} = \sum_k \frac{\partial \mathcal{L}}{\partial h_k^{(L)}} \cdot \frac{\partial h_k^{(L)}}{\partial z_j^{(L)}}$$

For element-wise activation:

$$\frac{\partial h_k^{(L)}}{\partial z_j^{(L)}} = \phi'(z_j^{(L)}) \cdot \mathbf{1}_{j=k}$$

Result:

$$\delta_j^{(L)} = \frac{\partial \mathcal{L}}{\partial h_j^{(L)}} \cdot \phi'(z_j^{(L)})$$

In vector form:  $\delta^{(L)} = \nabla_{\mathbf{h}^{(L)}} \mathcal{L} \odot \phi'(\mathbf{z}^{(L)})$

---

Starting point for backward pass

## B.7 Special Case: MSE + Sigmoid

### Setup:

- Loss:  $\mathcal{L} = \frac{1}{2} \|\hat{\mathbf{y}} - \mathbf{y}\|^2 = \frac{1}{2} \sum_j (h_j^{(L)} - y_j)^2$
- Activation:  $\phi(z) = \sigma(z) = \frac{1}{1+e^{-z}}$
- Derivative:  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$

### Gradient of Loss:

$$\frac{\partial \mathcal{L}}{\partial h_j^{(L)}} = h_j^{(L)} - y_j = \hat{y}_j - y_j$$

### Output Layer Error:

$$\begin{aligned}\delta_j^{(L)} &= (h_j^{(L)} - y_j) \cdot \sigma(z_j^{(L)}) (1 - \sigma(z_j^{(L)})) \\ &= (\hat{y}_j - y_j) \cdot \hat{y}_j (1 - \hat{y}_j)\end{aligned}$$

**Vector form:**  $\delta^{(L)} = (\hat{\mathbf{y}} - \mathbf{y}) \odot \hat{\mathbf{y}} \odot (1 - \hat{\mathbf{y}})$

---

Complete derivation for common case

## B.8 Special Case: Cross-Entropy + Softmax

### Setup:

- Loss:  $\mathcal{L} = -\sum_j y_j \log(\hat{y}_j)$  (one-hot  $\mathbf{y}$ )
- Activation:  $\hat{y}_j = \text{softmax}(z_j^{(L)}) = \frac{e^{z_j^{(L)}}}{\sum_k e^{z_k^{(L)}}}$

### Combined Derivative (remarkably simple):

$$\delta_j^{(L)} = \frac{\partial \mathcal{L}}{\partial z_j^{(L)}} = \hat{y}_j - y_j$$

### Derivation Sketch:

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial z_j^{(L)}} &= -\sum_k y_k \frac{\partial \log \hat{y}_k}{\partial z_j^{(L)}} = -\sum_k y_k \frac{1}{\hat{y}_k} \frac{\partial \hat{y}_k}{\partial z_j^{(L)}} \\ &= -y_j(1 - \hat{y}_j) + \sum_{k \neq j} y_k \hat{y}_j = \hat{y}_j - y_j\end{aligned}$$

Vector form:  $\boldsymbol{\delta}^{(L)} = \hat{\mathbf{y}} - \mathbf{y}$

---

Complete derivation for classification - elegantly simple!

## B.9 Hidden Layer Error: Setup

**Goal:** Express  $\delta^{(l)}$  in terms of  $\delta^{(l+1)}$

**The Computational Graph:**

$$\mathbf{z}^{(l)} \rightarrow \mathbf{h}^{(l)} \rightarrow \mathbf{z}^{(l+1)} \rightarrow \mathbf{h}^{(l+1)} \rightarrow \dots \rightarrow \mathcal{L}$$

**Chain Rule:**

$$\delta_j^{(l)} = \frac{\partial \mathcal{L}}{\partial z_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{(l+1)}} \cdot \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}}$$

**Breaking It Down:**

1.  $\frac{\partial \mathcal{L}}{\partial z_k^{(l+1)}} = \delta_k^{(l+1)}$  (already computed)
2.  $\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}}$  needs to be derived

---

Applying the chain rule through layers

## B.10 Hidden Layer Error: Step 1

**Computing**  $\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}}$ :

**Recall:**

$$z_k^{(l+1)} = \sum_m W_{km}^{(l+1)} h_m^{(l)} + b_k^{(l+1)}$$

**Chain Rule:**

$$\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \sum_m W_{km}^{(l+1)} \frac{\partial h_m^{(l)}}{\partial z_j^{(l)}}$$

**For element-wise activation:**

$$\frac{\partial h_m^{(l)}}{\partial z_j^{(l)}} = \phi'(z_j^{(l)}) \cdot \mathbf{1}_{m=j}$$

**Therefore:**

$$\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = W_{kj}^{(l+1)} \cdot \phi'(z_j^{(l)})$$

---

**First component of the chain**

## B.11 Hidden Layer Error: Step 2

**Substituting Back:**

$$\begin{aligned}\delta_j^{(l)} &= \sum_k \delta_k^{(l+1)} \cdot \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} \\ &= \sum_k \delta_k^{(l+1)} \cdot W_{kj}^{(l+1)} \cdot \phi'(z_j^{(l)}) \\ &= \phi'(z_j^{(l)}) \cdot \sum_k W_{kj}^{(l+1)} \delta_k^{(l+1)} \\ &= \phi'(z_j^{(l)}) \cdot [\mathbf{W}^{(l+1)T} \boldsymbol{\delta}^{(l+1)}]_j\end{aligned}$$

**Interpretation:**

- $\mathbf{W}^{(l+1)T} \boldsymbol{\delta}^{(l+1)}$ : Error from next layer, weighted by connection strengths
- $\phi'(z_j^{(l)})$ : Scaled by local gradient of activation

---

**Second component of the chain**

## B.12 Hidden Layer Error: Activation Derivatives

### Common Activation Functions and Their Derivatives:

Activation	$\phi(z)$	$\phi'(z)$
Sigmoid	$\frac{1}{1+e^{-z}}$	$\phi(z)(1-\phi(z))$
Tanh	$\frac{e^z - e^{-z}}{e^z + e^{-z}}$	$1 - \phi(z)^2$
ReLU	$\max(0, z)$	$\mathbf{1}_{z>0}$
Leaky ReLU	$\max(\alpha z, z)$	$\alpha \mathbf{1}_{z\leq 0} + \mathbf{1}_{z>0}$

**Note:** ReLU derivative is 0 for  $z \leq 0$ , 1 for  $z > 0$

This can cause “dead neurons” if  $z$  is always negative.

---

Third component of the chain

## B.13 Hidden Layer Error: Complete Formula

### Backpropagation Recursion

$$\delta^{(l)} = (\mathbf{W}^{(l+1)T} \delta^{(l+1)}) \odot \phi'(\mathbf{z}^{(l)})$$

**Element-wise:**

$$\delta_j^{(l)} = \phi'(z_j^{(l)}) \cdot \sum_k W_{kj}^{(l+1)} \delta_k^{(l+1)}$$

**Interpretation:**

- **Matrix multiply:**  $\mathbf{W}^{(l+1)T} \delta^{(l+1)}$  - “pull back” error through weights
- **Element-wise multiply:**  $\odot \phi'(\mathbf{z}^{(l)})$  - scale by local gradient

**Why “Back” Propagation?**

- Errors flow backward:  $\delta^{(L)} \rightarrow \delta^{(L-1)} \rightarrow \dots \rightarrow \delta^{(1)}$
- Uses transpose of forward weights  $\mathbf{W}^T$

---

The backpropagation recursion

## B.14 Weight Gradient: Setup

**Goal:** Compute  $\frac{\partial \mathcal{L}}{\partial W_{jk}^{(l)}}$

**How  $W_{jk}^{(l)}$  Affects the Loss:**

$$W_{jk}^{(l)} \rightarrow z_j^{(l)} \rightarrow h_j^{(l)} \rightarrow \dots \rightarrow \mathcal{L}$$

**Chain Rule:**

$$\frac{\partial \mathcal{L}}{\partial W_{jk}^{(l)}} = \frac{\partial \mathcal{L}}{\partial z_j^{(l)}} \cdot \frac{\partial z_j^{(l)}}{\partial W_{jk}^{(l)}}$$

**We Already Know:**

$$\frac{\partial \mathcal{L}}{\partial z_j^{(l)}} = \delta_j^{(l)}$$

**Need to Compute:**

$$\frac{\partial z_j^{(l)}}{\partial W_{jk}^{(l)}} = ?$$

---

Computing the gradient with respect to weights

## B.15 Weight Gradient: Derivation

Recall:

$$z_j^{(l)} = \sum_m W_{jm}^{(l)} h_m^{(l-1)} + b_j^{(l)}$$

Partial Derivative:

$$\frac{\partial z_j^{(l)}}{\partial W_{jk}^{(l)}} = h_k^{(l-1)}$$

Therefore:

$$\frac{\partial \mathcal{L}}{\partial W_{jk}^{(l)}} = \delta_j^{(l)} \cdot h_k^{(l-1)}$$

Interpretation:

- Error at neuron  $j$  ( $\delta_j^{(l)}$ )
- Times activation of input  $k$  ( $h_k^{(l-1)}$ )
- “Blame” is proportional to both

---

Applying the chain rule to weights

## B.16 Weight Gradient: Complete Formula

### Weight Gradient Formula

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}} = \boldsymbol{\delta}^{(l)} (\mathbf{h}^{(l-1)})^T$$

#### Dimensions:

- $\boldsymbol{\delta}^{(l)} \in \mathbb{R}^{n_l}$  (column vector)
- $\mathbf{h}^{(l-1)} \in \mathbb{R}^{n_{l-1}}$  (column vector)
- $\boldsymbol{\delta}^{(l)}(\mathbf{h}^{(l-1)})^T \in \mathbb{R}^{n_l \times n_{l-1}}$  (outer product = matrix)

#### Element-wise:

$$\left[ \frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}} \right]_{jk} = \delta_j^{(l)} h_k^{(l-1)}$$

**Intuition:** The gradient is an outer product of error and activation.

---

The weight update formula

## B.17 Bias Gradient Derivation

**Goal:** Compute  $\frac{\partial \mathcal{L}}{\partial b_j^{(l)}}$

**Recall:**

$$z_j^{(l)} = \sum_m W_{jm}^{(l)} h_m^{(l-1)} + b_j^{(l)}$$

**Partial Derivative:**

$$\frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = 1$$

**Therefore:**

$$\frac{\partial \mathcal{L}}{\partial b_j^{(l)}} = \delta_j^{(l)} \cdot 1 = \delta_j^{(l)}$$

### Bias Gradient Formula

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}^{(l)}} = \boldsymbol{\delta}^{(l)}$$

**Note:** The bias gradient is simply the error signal itself!

Simpler than the weight gradient

## B.18 Complete Backpropagation Algorithm

**Input:** Training example  $(\mathbf{x}, \mathbf{y})$ , network weights  $\{\mathbf{W}^{(l)}, \mathbf{b}^{(l)}\}$

**Forward Pass:**

1. Set  $\mathbf{h}^{(0)} = \mathbf{x}$
2. For  $l = 1$  to  $L$ :
  - $\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{h}^{(l-1)} + \mathbf{b}^{(l)}$
  - $\mathbf{h}^{(l)} = \phi^{(l)}(\mathbf{z}^{(l)})$
3. Compute loss  $\mathcal{L}(\mathbf{h}^{(L)}, \mathbf{y})$

**Backward Pass:**

1. Compute  $\delta^{(L)} = \nabla_{\mathbf{h}^{(L)}} \mathcal{L} \odot \phi'^{(L)}(\mathbf{z}^{(L)})$
2. For  $l = L - 1$  to 1:
  - $\delta^{(l)} = (\mathbf{W}^{(l+1)T} \delta^{(l+1)}) \odot \phi'^{(l)}(\mathbf{z}^{(l)})$

**Gradients:**  $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}} = \delta^{(l)} (\mathbf{h}^{(l-1)})^T$ ,  $\frac{\partial \mathcal{L}}{\partial \mathbf{b}^{(l)}} = \delta^{(l)}$

---

The full algorithm

## B.19 Worked Example: 2-2-1 Network

**Network:** 2 inputs, 2 hidden neurons (sigmoid), 1 output (sigmoid), MSE loss

**Given:**  $\mathbf{x} = [1, 0.5]^T$ ,  $y = 1$ ,  $\mathbf{W}^{(1)} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$ ,  $\mathbf{W}^{(2)} = [0.5, 0.6]$

**Forward Pass:**

$$\mathbf{z}^{(1)} = \mathbf{W}^{(1)}\mathbf{x} = [0.2, 0.5]^T$$

$$\mathbf{h}^{(1)} = \sigma(\mathbf{z}^{(1)}) = [0.550, 0.622]^T$$

$$z^{(2)} = \mathbf{W}^{(2)}\mathbf{h}^{(1)} = 0.648$$

$$\hat{y} = \sigma(z^{(2)}) = 0.656$$

**Backward Pass:**

$$\delta^{(2)} = (\hat{y} - y) \cdot \hat{y}(1 - \hat{y}) = -0.078$$

$$\boldsymbol{\delta}^{(1)} = (\mathbf{W}^{(2)T}\delta^{(2)}) \odot \mathbf{h}^{(1)} \odot (1 - \mathbf{h}^{(1)}) = [-0.010, -0.011]^T$$

---

Following the numbers through the algorithm

## B.20 Computational Complexity Analysis

For a Network with  $L$  Layers:

Forward Pass:

- Layer  $I$ : Matrix-vector multiply  $\mathbf{W}^{(I)} \mathbf{h}^{(I-1)}$
- Cost:  $O(n_I \times n_{I-1})$  per layer
- Total:  $O(\sum_I n_I n_{I-1}) = O(W)$  where  $W = \text{total weights}$

Backward Pass:

- Layer  $I$ : Matrix-vector multiply  $\mathbf{W}^{(I+1)T} \delta^{(I+1)}$
- Same cost as forward:  $O(n_{I+1} \times n_I)$  per layer
- Total:  $O(W)$

Key Result:

### Backpropagation Complexity

Computing all gradients:  $O(W) = \text{same as one forward pass!}$

Naïve finite differences would cost  $O(W^2)$  - backprop is  $W \times$  faster.

Why backpropagation is efficient

## C.1 MSE from Maximum Likelihood

**Setup:** Regression problem where we model:

$$y = f_{\theta}(\mathbf{x}) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2)$$

**This Implies:**

$$p(y|\mathbf{x}, \theta) = \mathcal{N}(y; f_{\theta}(\mathbf{x}), \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - f_{\theta}(\mathbf{x}))^2}{2\sigma^2}\right)$$

**Log-Likelihood for Dataset**  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ :

$$\begin{aligned} \log p(\mathbf{y}|\mathbf{X}, \theta) &= \sum_{i=1}^n \log p(y_i|\mathbf{x}_i, \theta) \\ &= -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - f_{\theta}(\mathbf{x}_i))^2 \end{aligned}$$

---

**Why MSE is the natural choice for regression**

## C.2 MSE Derivation

**Maximum Likelihood Estimation:**

$$\hat{\theta}_{ML} = \arg \max_{\theta} \log p(\mathbf{y}|\mathbf{X}, \theta)$$

**Equivalently (dropping constants):**

$$\hat{\theta}_{ML} = \arg \min_{\theta} \sum_{i=1}^n (y_i - f_{\theta}(\mathbf{x}_i))^2$$

**This is Mean Squared Error (MSE):**

$$\mathcal{L}_{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

### Key Insight

MSE loss = Maximum likelihood under Gaussian noise assumption

The  $\frac{1}{n}$  is for averaging; the  $\frac{1}{2}$  often added is for cleaner gradients.

---

From Gaussian likelihood to squared error

## C.3 Cross-Entropy from Maximum Likelihood

**Setup:** Binary classification where:

$$p(y = 1 | \mathbf{x}, \theta) = f_{\theta}(\mathbf{x}) = \hat{y} \in [0, 1]$$

**Bernoulli Distribution:**

$$p(y | \mathbf{x}, \theta) = \hat{y}^y (1 - \hat{y})^{1-y}$$

**Log-Likelihood:**

$$\log p(y | \mathbf{x}, \theta) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

**Negative Log-Likelihood (what we minimize):**

$$\mathcal{L} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

This is the **Binary Cross-Entropy** loss!

---

Why cross-entropy is natural for classification

## C.4 Binary Cross-Entropy Derivation

**Loss Function:**

$$\mathcal{L}_{BCE} = -\frac{1}{n} \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

**Gradient with Respect to  $\hat{y}$ :**

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = -\frac{y}{\hat{y}} + \frac{1-y}{1-\hat{y}} = \frac{\hat{y}-y}{\hat{y}(1-\hat{y})}$$

**Combined with Sigmoid Output:**

If  $\hat{y} = \sigma(z)$ , then  $\frac{\partial \hat{y}}{\partial z} = \hat{y}(1 - \hat{y})$

$$\frac{\partial \mathcal{L}}{\partial z} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z} = \frac{\hat{y}-y}{\hat{y}(1-\hat{y})} \cdot \hat{y}(1-\hat{y}) = \hat{y} - y$$

**Beautifully simple gradient!**

---

The two-class case

## C.5 Categorical Cross-Entropy Derivation

**Setup:**  $K$ -class classification with one-hot encoding  $\mathbf{y} \in \{0, 1\}^K$

**Categorical Distribution:**

$$p(\mathbf{y}|\mathbf{x}, \theta) = \prod_{k=1}^K \hat{y}_k^{y_k}$$

**Negative Log-Likelihood:**

$$\mathcal{L}_{CE} = - \sum_{k=1}^K y_k \log \hat{y}_k$$

**For Single Sample (one-hot  $\mathbf{y}$  with  $y_c = 1$ ):**

$$\mathcal{L} = - \log \hat{y}_c$$

**Interpretation:** Minimize negative log of predicted probability for correct class.

---

The multi-class case

## C.6 Softmax Function Derivation

**Problem:** Convert raw scores  $\mathbf{z} \in \mathbb{R}^K$  to probabilities  $\hat{\mathbf{y}} \in [0, 1]^K$

**Requirements:**

- $\hat{y}_k \geq 0$  for all  $k$
- $\sum_k \hat{y}_k = 1$
- Larger  $z_k \rightarrow$  larger  $\hat{y}_k$

**Solution - Softmax:**

$$\hat{y}_k = \text{softmax}(z_k) = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}}$$

**Properties:**

- Exponential ensures positivity
- Normalization ensures sum to 1
- Invariant to adding constant:  $\text{softmax}(z_k + c) = \text{softmax}(z_k)$
- In limit:  $\text{softmax} \rightarrow \arg \max$  ("soft" version)

---

Converting scores to probabilities

## C.7 Softmax Gradient

Computing  $\frac{\partial \hat{y}_i}{\partial z_j}$ :

**Case 1:**  $i = j$

$$\frac{\partial \hat{y}_i}{\partial z_i} = \frac{e^{z_i} \cdot Z - e^{z_i} \cdot e^{z_i}}{Z^2} = \hat{y}_i - \hat{y}_i^2 = \hat{y}_i(1 - \hat{y}_i)$$

**Case 2:**  $i \neq j$

$$\frac{\partial \hat{y}_i}{\partial z_j} = \frac{0 - e^{z_i} \cdot e^{z_j}}{Z^2} = -\hat{y}_i \hat{y}_j$$

**Jacobian Matrix:**

$$\frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}} = \text{diag}(\hat{\mathbf{y}}) - \hat{\mathbf{y}}\hat{\mathbf{y}}^T$$

**Combined with Cross-Entropy:**  $\frac{\partial \mathcal{L}}{\partial z_j} = \hat{y}_j - y_j$  (same simple form!)

---

The derivative of softmax

## D.1 L2 Regularization: Bayesian View

### Bayesian Setup:

- Prior belief about weights:  $p(\theta)$
- Likelihood of data given weights:  $p(D|\theta)$
- Posterior:  $p(\theta|D) \propto p(D|\theta)p(\theta)$

### L2 Corresponds to Gaussian Prior:

$$p(\theta) = \mathcal{N}(\mathbf{0}, \sigma_\theta^2 \mathbf{I}) = \prod_i \frac{1}{\sqrt{2\pi\sigma_\theta^2}} \exp\left(-\frac{\theta_i^2}{2\sigma_\theta^2}\right)$$

### Log Prior:

$$\log p(\theta) = -\frac{1}{2\sigma_\theta^2} \sum_i \theta_i^2 + \text{const} = -\frac{1}{2\sigma_\theta^2} \|\theta\|_2^2 + \text{const}$$

**Interpretation:** We believe weights should be small (centered at 0).

---

Referenced in Module 4

## D.2 L2 from MAP Estimation

**Maximum A Posteriori (MAP) Estimation:**

$$\hat{\theta}_{MAP} = \arg \max_{\theta} p(\theta|D) = \arg \max_{\theta} p(D|\theta)p(\theta)$$

**Taking Logs:**

$$\hat{\theta}_{MAP} = \arg \max_{\theta} [\log p(D|\theta) + \log p(\theta)]$$

**With Gaussian Prior and Gaussian Likelihood (MSE):**

$$\hat{\theta}_{MAP} = \arg \min_{\theta} \left[ \frac{1}{2\sigma^2} \sum_i (y_i - f_{\theta}(\mathbf{x}_i))^2 + \frac{1}{2\sigma_{\theta}^2} \|\theta\|_2^2 \right]$$

**Defining**  $\lambda = \frac{\sigma^2}{\sigma_{\theta}^2}$ :

$$\hat{\theta}_{MAP} = \arg \min_{\theta} [\mathcal{L}_{MSE} + \lambda \|\theta\|_2^2]$$

This is exactly L2 regularization (weight decay)!

---

From Gaussian prior to weight decay

## D.3 L1 Regularization: Bayesian View

**L1 Corresponds to Laplace Prior:**

$$p(\theta) = \prod_i \frac{1}{2b} \exp\left(-\frac{|\theta_i|}{b}\right)$$

**Log Prior:**

$$\log p(\theta) = -\frac{1}{b} \sum_i |\theta_i| + \text{const} = -\frac{1}{b} \|\theta\|_1 + \text{const}$$

**Why Laplace Induces Sparsity:**

- Laplace has sharp peak at 0
- Higher probability mass near exactly 0
- MAP estimate tends to push weights to exactly 0

**Geometric View:**

- L2: Ball constraint (smooth, no corners)
- L1: Diamond constraint (corners at axes)
- Solution often lands on corners → sparse

---

**Why L1 induces sparsity**

## D.4 L1 from MAP Estimation

**MAP with Laplace Prior:**

$$\hat{\theta}_{MAP} = \arg \max_{\theta} \left[ \log p(D|\theta) - \frac{1}{b} \|\theta\|_1 \right]$$

**Equivalently:**

$$\hat{\theta}_{MAP} = \arg \min_{\theta} [\mathcal{L} + \lambda \|\theta\|_1]$$

where  $\lambda = \frac{\sigma^2}{b}$  (ratio of noise variance to prior scale).

**Why Some Weights Become Exactly Zero:**

- L1 gradient is  $\pm\lambda$  (constant magnitude)
- Even small weights get constant “push” toward 0
- Eventually cross zero and stay there
- L2 gradient is  $\lambda\theta$  (proportional to  $\theta$ )
- Push decreases as  $\theta \rightarrow 0$ , never reaches 0

---

**From Laplace prior to sparse solutions**

## D.5 Dropout as Approximate Bayesian Inference

### Gal and Ghahramani (2016) Result:

Dropout training approximates variational inference in a deep Gaussian process.

### Key Insight:

- Dropout = sampling from approximate posterior
- Each forward pass samples different “network”
- Prediction uncertainty = variance across samples

### Monte Carlo Dropout:

1. Keep dropout enabled at test time
2. Run multiple forward passes
3. Mean = prediction, Variance = uncertainty

### Practical Benefit:

- Free uncertainty estimates!
- No additional training required
- Useful for detecting out-of-distribution inputs

---

### A theoretical justification for dropout

## D.6 Dropout: Mathematical Analysis

### Dropout During Training:

$$\tilde{\mathbf{h}}^{(l)} = \mathbf{m}^{(l)} \odot \mathbf{h}^{(l)}, \quad m_j^{(l)} \sim \text{Bernoulli}(1 - p)$$

### Expected Value:

$$\mathbb{E}[\tilde{h}_j^{(l)}] = (1 - p) \cdot h_j^{(l)}$$

### At Test Time (Inverted Dropout):

Scale by  $(1 - p)$  during training:  $\tilde{\mathbf{h}}^{(l)} = \frac{1}{1-p} \mathbf{m}^{(l)} \odot \mathbf{h}^{(l)}$

No scaling needed at test time.

### Implicit Regularization Effect:

- Prevents co-adaptation of neurons
- Each neuron must be useful independently
- Equivalent to training exponentially many networks
- Final model = average of all subnetworks

---

Why random dropout improves generalization

## **End of Mathematical Appendix**

Complete Derivations for Neural Networks

For questions about derivations, consult:

Goodfellow, Bengio, Courville - "Deep Learning" (2016)  
Bishop - "Pattern Recognition and Machine Learning" (2006)  
Nielsen - "Neural Networks and Deep Learning" (online)