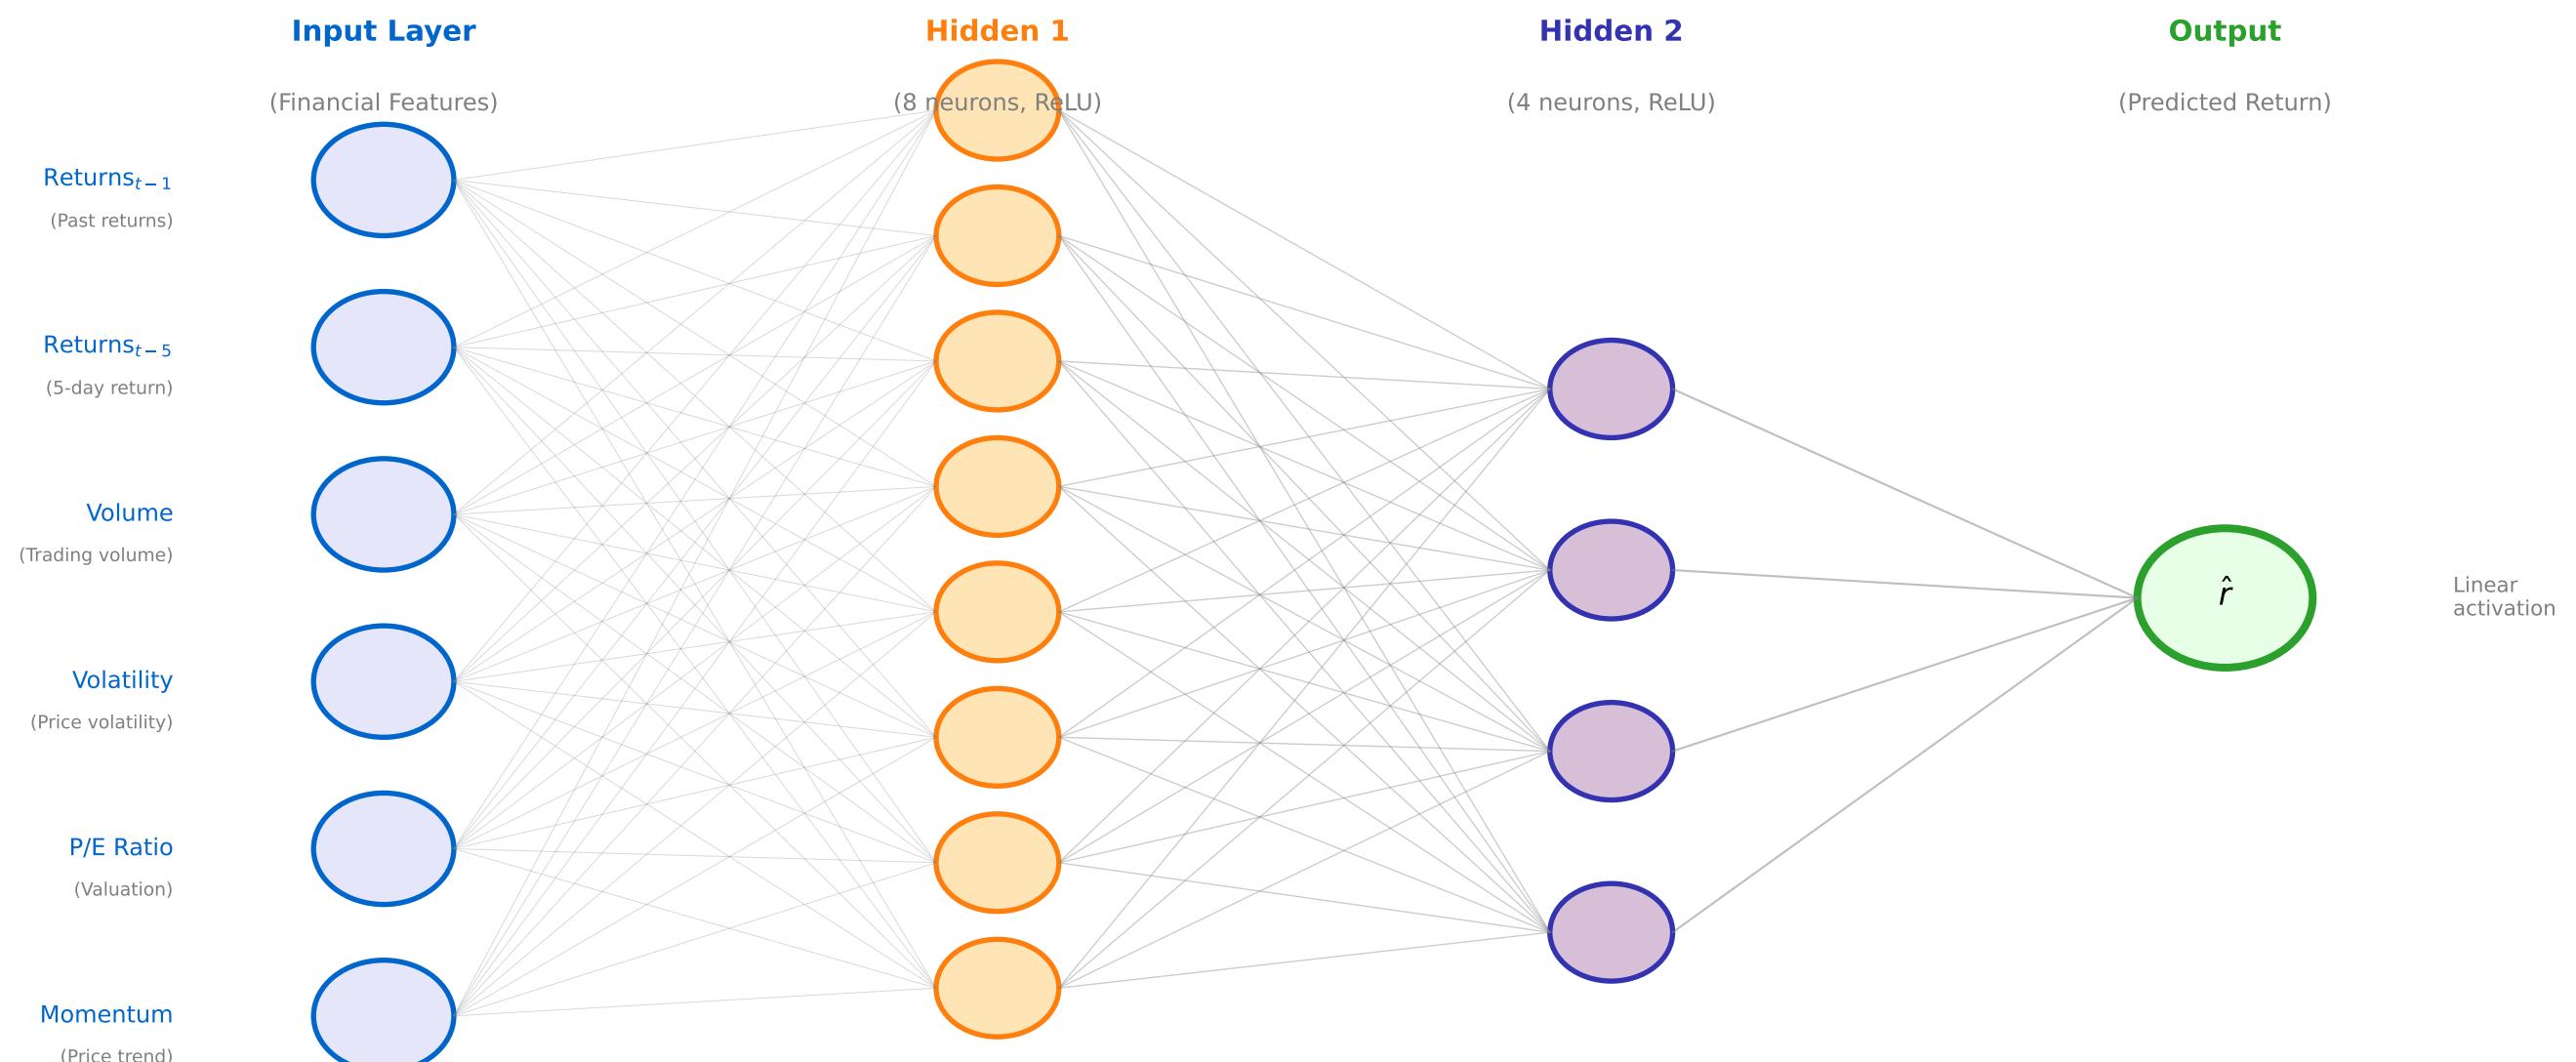


# MLP for Stock Return Prediction



Architecture Summary:  
- Input: 6 financial features (normalized)  
- Hidden 1: 8 neurons with ReLU  
- Hidden 2: 4 neurons with ReLU  
- Output: 1 (predicted return)  
- Loss: Mean Squared Error  
- Total parameters:  $6 \times 8 + 8 + 8 \times 4 + 4 + 4 \times 1 + 1 = 97$