

# Why Non-Linearity?

## The Core Problem

Without activation functions:

$$\mathbf{a}^{(1)} = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

$$\hat{\mathbf{y}} = \mathbf{W}^{(2)}\mathbf{a}^{(1)} + \mathbf{b}^{(2)}$$

Substituting:

$$\begin{aligned}\hat{\mathbf{y}} &= \mathbf{W}^{(2)}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) + \mathbf{b}^{(2)} \\ &= (\mathbf{W}^{(2)}\mathbf{W}^{(1)})\mathbf{x} + (\mathbf{W}^{(2)}\mathbf{b}^{(1)} + \mathbf{b}^{(2)}) \\ &= \mathbf{W}'\mathbf{x} + \mathbf{b}'\end{aligned}$$

**Result:** A single linear transformation!

## The Solution

Non-linear activation functions:

$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)})$$

where  $f$  is non-linear.

## Why This Works:

- Non-linearity breaks the collapse
- Composition of non-linear functions
- Can approximate any function

## Key Insight:

Non-linearity is what makes deep networks “deep” in a meaningful sense.

---

Non-linearity is essential for learning complex patterns

# Linear Networks Collapse

## Mathematical Proof

For any number of linear layers:

$$y = \mathbf{W}^{(L)} \mathbf{W}^{(L-1)} \dots \mathbf{W}^{(1)} \mathbf{x}$$

Since matrix multiplication is associative:

$$= (\mathbf{W}^{(L)} \mathbf{W}^{(L-1)} \dots \mathbf{W}^{(1)}) \mathbf{x}$$

$$= \mathbf{W}^{\text{eff}} \mathbf{x}$$

## Conclusion:

100 linear layers = 1 linear layer.


No benefit from depth without non-linearity.

---

Stacked linear layers = single linear layer

## Why Non-Linear Activations Are Essential

### Without Non-Linear Activation




Multiple linear layers = One linear layer

Math:

$$h = W_1 x$$
$$y = W_2 h$$
$$y = W_2(W_1 x)$$
$$y = (W_2 W_1)x$$
$$y = W_{\text{tot}}$$

Still just a linear transform!

### With Non-Linear Activation



Activation functions enable deep learning

Math:

$$h = \sigma(W_1 x)$$
$$y = W_2 h$$
$$y = W_2 \sigma(W_1 x)$$

Cannot simplify!

Non-linearity breaks collapse

Each layer adds expressiveness



...

linear-collapse\_proof

# The Sigmoid Function

## Definition

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

## Properties:

- Range:  $(0, 1)$
- Smooth and differentiable
- $\sigma(0) = 0.5$
- Symmetric:  $\sigma(-z) = 1 - \sigma(z)$


## Derivative:

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

## Use Cases:

- Binary classification (output)
- Probability interpretation
- Historical (hidden layers)

The classic activation: squashes to probability



.../..

sigmoid\_function

## Advantages

- + Bounded output (0, 1)
- + Smooth gradient
- + Probability interpretation
- + Historically important

## Disadvantages

### - Vanishing gradients

For  $|z| > 4$ :  $\sigma'(z) \approx 0$

Gradients become tiny

Deep networks can't learn

### - Not zero-centered

All positive outputs

Zig-zag weight updates

### - Computationally expensive

Requires exp function

Smooth and bounded, but gradients can vanish

## The Vanishing Gradient Problem

When  $z$  is very positive or negative:

$z$	$\sigma'(z)$
0	0.25
2	0.10
4	0.018
6	0.0025

Gradients shrink exponentially through layers!

**Result:** Early layers learn very slowly in deep networks.  
This limited deep learning until ReLU.

# The Tanh Function

## Definition

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} = 2\sigma(2z) - 1$$

## Properties:

- Range:  $(-1, 1)$
- Zero-centered
- $\tanh(0) = 0$
- Odd function:  $\tanh(-z) = -\tanh(z)$


## Derivative:

$$\tanh'(z) = 1 - \tanh^2(z)$$

## Advantage over Sigmoid:

Zero-centered outputs lead to more stable gradient updates.

Zero-centered: range  $(-1, 1)$



...

tanh.function

# ReLU: Rectified Linear Unit

## Definition

$$\text{ReLU}(z) = \max(0, z) = \begin{cases} z & z > 0 \\ 0 & z \leq 0 \end{cases}$$

## Properties:


- Range:  $[0, \infty)$
- Not bounded above
- Not differentiable at  $z = 0$
- Piecewise linear

## Derivative:

$$\text{ReLU}'(z) = \begin{cases} 1 & z > 0 \\ 0 & z \leq 0 \end{cases}$$

The Modern Default for hidden layers.

Simple but powerful: the modern default



...

## Advantages

### + No vanishing gradient

Gradient is 1 for  $z > 0$

Signal propagates through layers

### + Computationally cheap

Just comparison and assignment

No exponentials

6x faster than sigmoid

### + Sparse activation

Many neurons output 0

Efficient representation

### + Biological plausibility

Neurons can be “off”

## Disadvantages

### - “Dying ReLU” problem

If  $z < 0$  always: gradient = 0

Neuron never updates

Can “die” permanently

### - Not zero-centered

### - Unbounded (can explode)

## Variants:





- Leaky ReLU:  $\max(0.01z, z)$
- ELU:  $z$  if  $z > 0$ ,  $\alpha(e^z - 1)$  otherwise
- GELU: used in transformers

---

Cheap to compute, gradients don't vanish (for positive inputs)

# Activation Functions: Comparison

## Activation Functions: Function and Derivative



...

[activation\\_comparison](activation_comparison.html)

Different functions for different problems

*"Which activation function would you use for: (a) predicting stock returns, (b) buy/sell classification? Why?"*

Consider:

**(a) Stock Returns (Regression)**

- Output: continuous value
- Can be positive or negative
- Hidden: ReLU or tanh
- Output: **Linear (none)**
- Returns are unbounded

**(b) Buy/Sell (Classification)**

- Output: probability  $\in (0, 1)$
- Two mutually exclusive classes
- Hidden: ReLU
- Output: **Sigmoid**
- Or softmax for multi-class

---

Think-Pair-Share: 3 minutes

## Hidden Layer Guidelines

### Default: ReLU

- Works well in most cases
- Fast and stable

### If dying ReLU: Leaky ReLU

- Small negative slope
- Prevents dead neurons

### For RNNs: Tanh

- Bounded outputs help stability
- Zero-centered

## Output Layer Guidelines

Task	Activation
Binary class	Sigmoid
Multi-class	Softmax
Regression	Linear
Bounded regression	Sigmoid/tanh
Positive only	ReLU

### Finance Examples:

- Return prediction: Linear
- Direction prediction: Sigmoid
- Sector classification: Softmax
- Volatility: ReLU or Softplus

---

Output layer choice depends on your problem type

# The Fundamental Question

## How Powerful Are Neural Networks?

We've seen that MLPs can:

- Solve XOR (non-linear patterns)
- Combine features hierarchically
- Learn from data

## But a Deeper Question:

Are there functions that MLPs fundamentally *cannot* represent?

Or can they approximate *anything*?

## Why This Matters

### If MLPs are limited:

- Need to check if problem is solvable
- Architecture constraints matter
- Some patterns impossible

### If MLPs are universal:

- Architecture is not the bottleneck
- Challenges are elsewhere (data, training)
- Theoretical guarantee of capability

**Spoiler:** MLPs are universal approximators!

---

Just how powerful are neural networks?

## The Theorem (Informal)

A feedforward network with:

- One hidden layer
- Sufficient hidden neurons
- Non-linear activation (e.g., sigmoid)

can approximate any continuous function on a compact domain to arbitrary accuracy.

## Key Contributors:

- Cybenko (1989): sigmoid
- Hornik (1991): general activations
- Further extensions since

## Formal Statement

Let  $f : [0, 1]^n \rightarrow \mathbb{R}$  be continuous.

For any  $\epsilon > 0$ , there exists an MLP  $\hat{f}$  with:

$$|\hat{f}(\mathbf{x}) - f(\mathbf{x})| < \epsilon$$

for all  $\mathbf{x} \in [0, 1]^n$ .

## In Plain English:

No matter how complex the pattern, an MLP with enough hidden neurons can match it as closely as you want.

---

With enough hidden neurons, you can approximate any continuous function

# What Universal Approximation Means


## The Good News

- No function is “too complex”
- MLPs are theoretically complete
- Architecture is not the limit
- One hidden layer is enough (in theory)

## Visual Intuition:

Each hidden neuron contributes a “bump” or “step.” With enough bumps, you can approximate any shape.

Think of it like approximating a curve with many small line segments.



Universal Approximation Theorem (Cybenko, 1989): A feedforward network with a single hidden layer containing a finite number of neurons can approximate any continuous function on compact subsets of  $\mathbb{R}^n$

**More neurons = better approximation**

# What It Doesn't Mean

## Common Misconceptions

**"Any network can learn anything"**

No. Need enough neurons

- May need exponentially many

**"Training will find the solution"**

No. Theorem is about existence

- Says nothing about finding weights
- Optimization may fail

**"One layer is always enough"**

actually no.

- Deep networks often more efficient
- Fewer parameters for same accuracy

## The Gap: Existence vs Construction

**The theorem says:**

"A good approximation exists."

**It does NOT say:**

- How many neurons you need
- How to find the right weights
- How much data is required
- How long training takes
- Whether it will generalize

## Analogy:

"There exists a needle in this haystack" doesn't help you find it.

---

Existence of a solution does not mean we can find it

# Theory vs Practice

## Theoretical Guarantees

Universal approximation says:

- Given infinite neurons: perfect fit
- Given infinite data: find the function
- Given infinite compute: optimize

## Practical Reality

We have:

- Finite neurons: limited capacity
- Finite data: must generalize
- Finite compute: approximate solutions

## What Matters More in Practice

1. **Data quality and quantity**
  - More important than architecture
2. **Regularization**
  - Prevent overfitting
3. **Optimization**
  - Finding good weights
4. **Generalization**
  - Performance on new data

**Module 3** will address these practical challenges.

---

Universal approximation is necessary but not sufficient

## The Optimistic View

If markets have patterns, MLPs can learn them:

- Non-linear relationships? Possible.
- Complex interactions? Possible.
- Hidden factors? Possible.

## Theoretical Capability:

"An MLP could, in principle, capture any market pattern."

## The Realistic View

### Challenges Remain:

- Signal-to-noise ratio is low
- Markets are non-stationary
- Past patterns may not repeat
- Data is limited (especially for crashes)
- Overfitting is easy

### The EMH Counterargument:

If markets are efficient, there's nothing systematic to learn.

*Module 4 will explore this tension.*

---

In theory, yes. In practice, many challenges remain.

# Why Loss Functions?

## Learning Requires an Objective

To train a neural network, we need:

1. A way to measure errors
2. A number that decreases as we improve
3. A signal for weight updates

## The Loss Function:

$\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$  measures how wrong our predictions are.

## Goal of Training:

Find weights that minimize  $\mathcal{L}$ .

## Finance Analogy

### Profit & Loss (P&L):

- Measures trading performance
- Negative P&L = bad trades
- Optimize to maximize P&L

### Loss Function:

- Measures prediction errors
- High loss = bad predictions
- Optimize to minimize loss

**Note:** “Loss” is the opposite of “profit” – we minimize loss!

---

To learn, we must measure mistakes

# Mean Squared Error (MSE)

## Definition


$$\mathcal{L}_{\text{MSE}} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

## Properties:

- Always non-negative
- Zero only if perfect predictions
- Penalizes large errors heavily
- Differentiable everywhere

## Use Case:

- Regression problems
- Predicting continuous values
- Stock returns, prices, etc.



MSE Properties	
Convex:	Single global minimum for linear models
Differentiable:	Smooth gradients everywhere
Scale-sensitive:	Large errors penalized more (squared)
Outlier-sensitive:	Outliers amplify outlier impact Regression problems with normally distributed errors



.../.

The standard loss for predicting continuous values

# Cross-Entropy Loss

## Binary Cross-Entropy


$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

### Properties:

- For probability outputs
- Heavily penalizes confident wrong answers
- Connected to information theory

### Use Case:

- Classification problems
- Buy/sell decisions
- Any yes/no prediction



### The standard loss for classification



.../.

cross\_entropy\_visualization

# The Loss Landscape

## Loss as a Function of Weights

$$\mathcal{L}(\mathbf{W}, \mathbf{b})$$

For every choice of weights, there's a loss value.

### The Landscape:

- High regions: bad weights
- Low regions: good weights
- Global minimum: best weights
- Local minima: traps



### Training =

Finding the lowest point in this landscape.

---

Training = finding the lowest point in this landscape

Loss Landscape: Why Deep Networks Are Hard to Train



## Task-Specific Loss Functions

Task	Loss
Return prediction	MSE
Direction prediction	Cross-entropy
Volatility forecast	MSE
Multi-class sector	Categorical CE

## Beyond Standard Losses:

- Sharpe ratio optimization
- Asymmetric losses (penalize losses more than gains)
- Custom finance metrics

## Important Consideration

### MSE vs Business Metric:

A model with low MSE may still lose money!

### Example:

- Predict returns with 5% MSE
- But wrong on big moves
- Transaction costs eat profits
- Risk-adjusted return is poor

### Lesson:

Statistical accuracy  $\neq$  Trading profitability  
Module 4 explores this gap.

---

Different problems, different loss functions

## What We Learned

### 1. Historical Context

- AI Winter (1969-1982)
- Backprop renaissance (1986)
- Right idea + right time

### 2. MLP Architecture

- Hidden layers find patterns
- Matrix notation for computation
- Parameter counting

### 3. Activation Functions

- Non-linearity is essential
- ReLU for hidden, task-specific for output

### 4. Universal Approximation

- MLPs can learn any function
- But existence  $\neq$  construction

### 5. Loss Functions

- MSE for regression
- Cross-entropy for classification
- Loss landscape visualization

## The Big Picture:

We now have powerful architectures. But how do they *learn*?

---

From single perceptron to universal function approximator

*"We have the architecture. But how does it LEARN?"*

### The Missing Piece

We know:

- How to compute forward pass
- What loss functions measure
- That good weights exist

We don't know:

- How to find good weights
- How errors update weights
- How to avoid overfitting

**Mathematical details: See Appendix B (Backpropagation Derivation)**

### Coming in Module 3:

- Gradient descent (intuition)
- Backpropagation (the magic)
- Training dynamics
- Overfitting and regularization
- Practical training tips

**The Key:** Backpropagation – the algorithm that made deep learning possible.

---

Next: The magic of backpropagation