

# Multi-Layer Perceptron Architecture

## Neural Networks for Finance

Neural Networks for Finance

BSc Lecture Series

November 30, 2025

## Module 1 Summary

We learned that a single perceptron:

- Takes weighted inputs
- Applies a threshold
- Outputs a binary decision
- Can only draw **linear** boundaries

## The Perceptron Equation:

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right)$$

---

The perceptron: powerful but limited

## The Problem

The perceptron cannot solve XOR or any non-linearly separable problem.

## The AI Winter:

- Minsky-Papert (1969) critique
- Funding dried up
- “Neural networks don’t work”

## Today's Question:

What if we stack multiple perceptrons together?



# The XOR Problem Revisited

## Why One Line Isn't Enough

$x_1$	$x_2$	XOR
0	0	0
0	1	1
1	0	1
1	1	0

## The Geometry:

- Opposite corners have same label
- No single line can separate them
- We need *multiple* boundaries



.../.

xor\_solution\_ml

Some patterns require more than a single line

## Single Analyst (Perceptron)

One junior analyst screening stocks:

- Looks at a few metrics
- Applies simple rules
- Makes direct decisions
- Limited perspective

### Limitation:

"Buy if P/E < 15 AND momentum > 0"

This is a single linear rule.

## Investment Team (MLP)

A hierarchical team:

- Junior analysts find patterns
- Senior analysts synthesize
- CIO makes final call
- Complex reasoning emerges

### Capability:

"Consider value metrics, momentum signals, AND market regime together"

Multiple non-linear patterns.

**Key Insight:** Hierarchical processing enables complex pattern recognition.

---

A single analyst sees simple patterns. A team sees complex ones.

## What We'll Cover

### 1. Historical Context

- AI Winter survival
- Backprop rediscovery (1986)

### 2. MLP Architecture

- Intuition: The firm analogy
- Math: Matrix notation

### 3. Activation Functions

- Why non-linearity matters
- Sigmoid, Tanh, ReLU

### 4. Universal Approximation

- The fundamental theorem
- Implications and limits

### 5. Loss Functions

- MSE for regression
- Cross-entropy for classification

## Learning Objectives:

- Understand MLP architecture
- Master matrix notation
- Know when to use which activation
- Appreciate universal approximation

---

From single perceptron to universal function approximation

# The AI Winter (1969-1982)


## After Minsky-Papert

The neural network winter:

- Government funding cut
- Researchers moved to other fields
- “Connectionism is dead”
- Symbolic AI dominated

The Mood:

- Perceptrons can't solve XOR
- Multi-layer networks exist but...
- No efficient training algorithm
- Why bother?



After Minsky-Papert, neural network research nearly died

## Paul Werbos (1974)

PhD thesis at Harvard:

- Derived backpropagation
- For general non-linear systems
- Applied to neural networks
- **Largely ignored**

## Why Ignored?

- Published in economics, not CS
- AI winter was at its coldest
- No computational power to test
- No community to spread ideas

## Parallel Discoveries

### 1970s:

- Linnainmaa: automatic differentiation
- Control theory: similar ideas

### 1980s:

- Parker (1982): rediscovery
- LeCun (1985): independent work
- Rumelhart/Hinton/Williams (1986): fame

**Lesson:** Good ideas can be discovered multiple times before they "take off."

---

The key ideas existed but were ignored

# 1982: Hopfield Networks

## John Hopfield

A physicist (not AI researcher) revived interest:

- Connected neural networks to physics
- Energy-based formulation
- Published in PNAS (prestigious)
- Showed neural nets could store memories

## The Impact:

- Legitimized neural network research
- Attracted physicists to the field
- New mathematical tools
- Funding started returning

## Why Physics Helped

### Physics Connection:

- Neurons  $\leftrightarrow$  spins in magnets
- Learning  $\leftrightarrow$  energy minimization
- Networks  $\leftrightarrow$  statistical mechanics

### Finance Parallel:

Physicists would later apply similar ideas to:

- Option pricing
- Market dynamics
- Risk modeling
- Quantitative finance

---

John Hopfield: Physicist redisCOVERS neural networks

# 1986: The Backpropagation Paper

## The Paper That Changed Everything


Rumelhart, Hinton, Williams in Nature (1986):  
“Learning representations by back-propagating errors”

### Key Contributions:

- Clear algorithm presentation
- Demonstrated on real problems
- Published in high-impact journal
- Well-communicated to broad audience

### The Result:

Neural network renaissance begins.



Nature paper: “Learning representations by back-propagating errors”

# What Made 1986 Different?

## Werbos (1974)

- + Correct algorithm
- + General framework
- Wrong field (economics)
- No demonstrations
- No community
- No computers

## Lesson for Researchers:

Being right isn't enough. You need:

- The right timing
- The right communication
- The right audience
- The right technology

## Rumelhart et al. (1986)

- + Correct algorithm
- + Clear presentation
- + Compelling demos
- + High-profile venue (Nature)
- + Growing community
- + Computers available

---

The right idea at the right time with the right people

*"Backpropagation was discovered multiple times (1974, 1982, 1986). Why do some discoveries get ignored while others take off? What role did timing play?"*

Consider:

- Publication venue matters
- Community readiness
- Computational infrastructure
- Demonstration quality
- Today: transformers (2017) exploded
- LSTMs existed since 1997
- What changed?

---

Think-Pair-Share: 3 minutes

## After 1986

Neural networks were back:

- Funding returned
- New conferences (NIPS, now NeurIPS)
- “Connectionism” movement
- Real applications emerged

## Key Milestones:

- 1989: LeNet for digit recognition
- 1990s: Speech recognition
- 1990s: Financial applications begin

## But Challenges Remained

Not everything worked:

- Deep networks hard to train
- Vanishing gradients
- Limited compute power
- Another “winter” in 2000s

## True Revolution: 2012

AlexNet on ImageNet marked the deep learning era.  
(Module 4)

*But first, we need to understand the architecture...*

---

Neural networks are back - and this time they can learn

# The Investment Firm Analogy

## Hierarchical Decision Making

### Level 1: Junior Analysts (Hidden Layer 1)

- Look at raw data
- Find basic patterns
- “This looks like a value stock”
- “This has momentum”

### Level 2: Senior Analysts (Hidden Layer 2)

- Combine junior reports
- Higher-level synthesis
- “Value + momentum = quality”

### Level 3: CIO (Output Layer)

- Final buy/sell decision
- Combines all analyses
- Single decision point

### Key Properties:

1. Information flows upward
2. Each level adds abstraction
3. Later layers see patterns in patterns
4. Final layer integrates everything

This is an MLP!

---

## Hierarchical decision making

# Input Layer: The Data Gatherers

## The Input Layer

What it does:

- Receives raw data
- One neuron per feature
- No computation
- Just passes data forward

## In Finance:

- P/E ratio
- Momentum (returns)
- Volume
- Volatility
- Sector indicators
- Market cap

## Notation

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}$$

where:

- $n$  = number of features
- $x_i$  = value of feature  $i$

## Example (n=4):

$$\mathbf{x} = \begin{pmatrix} 15 \\ 0.08 \\ 1.2M \\ 0.25 \end{pmatrix} = \begin{pmatrix} \text{P/E} \\ \text{Return} \\ \text{Volume} \\ \text{Vol} \end{pmatrix}$$

---

The input layer receives raw information

# Hidden Layers: The Pattern Finders


## What Hidden Layers Do

They discover intermediate patterns:

- Not explicitly programmed
- Emerge from training
- Often uninterpretable
- But highly useful

## Each Hidden Neuron:

- Receives weighted inputs
- Applies activation function
- Outputs a single number
- “Detects” a specific pattern



.../.

hidden\_layer\_representation

“They see things in the data you didn’t explicitly ask for”

## Hypothetical Hidden Neurons

### Hidden Neuron 1: "Value Detector"

- Positive weight on low P/E
- Positive weight on high book value
- Activates for value stocks

### Hidden Neuron 2: "Momentum Detector"

- Positive weight on recent returns
- Positive weight on volume
- Activates for trending stocks

### Hidden Neuron 3: "Risk Detector"

- Positive weight on volatility
- Positive weight on debt
- Activates for risky stocks

## The Output Layer

Combines hidden neuron outputs:

$$\text{Buy} = f(w_1 \cdot \text{Value} + w_2 \cdot \text{Momentum} - w_3 \cdot \text{Risk})$$

### Key Insight:

We never told the network what "value" or "momentum" means. It *discovered* these concepts from data.

### Caveat:

Real hidden neurons may not be this interpretable. They might detect patterns we can't name.

---

Hidden neurons learn abstract concepts

## The Output Layer

Takes hidden representations and produces:

- Classification: probability of class
- Regression: continuous prediction
- Multiple outputs possible

### For Binary Classification:

Single output neuron with sigmoid:

$$\hat{y} = \sigma(w^T h + b)$$

Output  $\in (0, 1)$  interpreted as probability.

### For Regression:

Single output neuron with no activation (or linear):

$$\hat{y} = w^T h + b$$

Output is predicted value.

---

The output layer synthesizes everything into a decision

## Finance Examples

### Buy/Sell Classification:

- Output:  $P(\text{Buy})$
- If  $> 0.5$ : recommend Buy
- If  $< 0.5$ : recommend Sell


### Return Prediction:

- Output: predicted return
- Could be next-day, next-month
- Continuous value

### Multi-Class (Sector):

- Multiple output neurons
- Softmax activation
- Each output = probability of sector

# The Full MLP Architecture



.../.

mlp\_architecture\_2.3...

# Why Are They Called “Hidden”?

## We Don't Observe Them Directly

### Observable:

- Input layer: the features we provide
- Output layer: the prediction we get

### Hidden:

- Internal representations
- Not directly specified
- Learned automatically
- “Hidden” from us

## We Don't Tell Them What to Learn

### Traditional ML:

“Here are features: P/E, momentum, volume”  
We engineer the features.

### Deep Learning Philosophy:

“Here is raw data. Find useful patterns.”  
Network discovers features.

### Trade-off:

More automatic, but less interpretable.

---

Hidden layers discover features automatically

# How MLPs Solve XOR

## The Two-Hidden-Neuron Solution

### Hidden Neuron 1:

Learns: "Is it in the upper-right region?"

$$h_1 = \sigma(w_{11}x_1 + w_{12}x_2 + b_1)$$

### Hidden Neuron 2:



Learns: "Is it in the lower-left region?"

$$h_2 = \sigma(w_{21}x_1 + w_{22}x_2 + b_2)$$

### Output Neuron:

Combines: "If  $h_1$  XOR  $h_2$ , output 1"

Each hidden neuron draws *one* line. Together, they create a non-linear boundary.



---

Multiple decision boundaries working together

*"If hidden layers find features automatically, why do we still need feature engineering in finance?"*

Consider:

**Arguments for Feature Engineering:**

- Domain knowledge helps
- Less data needed
- More interpretable
- Faster training

**Arguments Against:**

- Human biases
- Miss non-obvious patterns
- Deep learning works on raw data
- ImageNet revolution

**Reality:** In finance, hybrid approaches often work best.

---

Think-Pair-Share: 3 minutes

# Universal Approximation: The Big Promise

## A Remarkable Theorem


With just *one* hidden layer and enough neurons, an MLP can approximate **any** continuous function to arbitrary accuracy.

### Implications:

- MLPs are universal function approximators
- No pattern is too complex (in theory)
- The architecture is not the bottleneck

### Caveats:

- “Enough neurons” may be exponential
- Finding the right weights is hard
- Theory vs practice gap



Universal Approximation Theorem (Cybenko, 1989): A feedforward network with a single hidden layer containing a finite number of neurons can approximate any continuous function on compact subsets of  $\mathbb{R}^n$

MLPs can learn ANY pattern (in theory)

# Now Let's Formalize

## What You Already Know

From the intuition section:

- Layers process sequentially
- Each layer transforms its input
- Hidden layers find patterns
- Output layer makes predictions

## What's Next

- Matrix notation for efficiency
- Precise forward pass equations
- Parameter counting
- Worked numerical examples

## Why Matrix Notation?

### Without Matrices:

Write  $n \times m$  separate equations for each weight.

### With Matrices:

$$\mathbf{h} = f(\mathbf{Wx} + \mathbf{b})$$

One equation captures everything.

### Benefits:

- Compact notation
- Efficient computation (GPUs)
- Easier to implement
- Clearer understanding

---

You understand the intuition. Let's write it precisely.

# Matrix Notation: Why Matrices?

## Single Neuron (Scalar)

$$h = f(w_1x_1 + w_2x_2 + w_3x_3 + b)$$

As Dot Product:

$$h = f(\mathbf{w}^T \mathbf{x} + b)$$

where  $\mathbf{w}, \mathbf{x} \in \mathbb{R}^3$



Multiple Neurons (Matrix):

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

where  $\mathbf{W} \in \mathbb{R}^{m \times n}$

Each row of  $\mathbf{W}$  is the weights for one hidden neuron.

Neural Network Layer as Matrix Multiplication



Matrices make neural network math elegant

# The Weight Matrix

## Weight Matrix $\mathbf{W}^{(l)}$

For layer  $l$ :

$$\mathbf{W}^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$$

where:

- $n_l$  = neurons in layer  $l$
- $n_{l-1}$  = neurons in layer  $l - 1$

## Entry $W_{ij}^{(l)}$ :

Weight from neuron  $j$  in layer  $l - 1$  to neuron  $i$  in layer  $l$ .

## Bias Vector $\mathbf{b}^{(l)}$

$$\mathbf{b}^{(l)} \in \mathbb{R}^{n_l}$$

One bias per neuron in layer  $l$ .

## Example: 4-3 Layer

Input: 4 neurons, Hidden: 3 neurons

$$\mathbf{W}^{(1)} = \begin{pmatrix} w_{11} & w_{12} & w_{13} & w_{14} \\ w_{21} & w_{22} & w_{23} & w_{24} \\ w_{31} & w_{32} & w_{33} & w_{34} \end{pmatrix}$$

Size:  $3 \times 4$  (12 weights)

$$\mathbf{b}^{(1)} \in \mathbb{R}^3 \text{ (3 biases)}$$

---

Each layer has its own weight matrix

# Forward Pass: Layer by Layer

## One Layer Computation

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}$$


$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)})$$

where:

- $\mathbf{z}^{(l)}$ : pre-activation (weighted sum)
- $\mathbf{a}^{(l)}$ : activation (after  $f$ )
- $\mathbf{a}^{(0)} = \mathbf{x}$ : input

## The Steps:

1. Matrix multiply:  $\mathbf{W}^{(l)} \mathbf{a}^{(l-1)}$
2. Add bias:  $+ \mathbf{b}^{(l)}$
3. Apply activation:  $f(\cdot)$



.../.

layer\_by\_layer\_computation

Computing outputs one layer at a time

# The Complete Forward Pass

For an L-Layer Network

Input:

$$\mathbf{a}^{(0)} = \mathbf{x}$$

Hidden Layers ( $l = 1, \dots, L - 1$ ):

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)}\mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}$$

$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)})$$

Output Layer:

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)}\mathbf{a}^{(L-1)} + \mathbf{b}^{(L)}$$

$$\hat{\mathbf{y}} = g(\mathbf{z}^{(L)})$$

where  $g$  may differ from  $f$ .

Example: 2-Layer Network

Layer 1 (hidden):

$$\mathbf{z}^{(1)} = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

$$\mathbf{a}^{(1)} = \text{ReLU}(\mathbf{z}^{(1)})$$

Layer 2 (output):

$$\mathbf{z}^{(2)} = \mathbf{W}^{(2)}\mathbf{a}^{(1)} + \mathbf{b}^{(2)}$$

$$\hat{\mathbf{y}} = \sigma(\mathbf{z}^{(2)})$$

Compact Form:

$$\hat{\mathbf{y}} = \sigma(\mathbf{W}^{(2)} \text{ReLU}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) + \mathbf{b}^{(2)})$$

---

Chaining layer computations together

## Dimension Checking

For  $\mathbf{z} = \mathbf{Wx} + \mathbf{b}$ :

$$\mathbf{W}: (n_{\text{out}} \times n_{\text{in}})$$

$$\mathbf{x}: (n_{\text{in}} \times 1)$$

$$\mathbf{Wx}: (n_{\text{out}} \times 1)$$

$$\mathbf{b}: (n_{\text{out}} \times 1)$$

$$\mathbf{z}: (n_{\text{out}} \times 1)$$

### Rule:

Inner dimensions must match.

$$(m \times n) \times (n \times p) = (m \times p)$$

## Example: 4-3-1 Network

### Layer 1:

- $\mathbf{W}^{(1)}: 3 \times 4$

- $\mathbf{x}: 4 \times 1$

- $\mathbf{z}^{(1)}: 3 \times 1$

### Layer 2:

- $\mathbf{W}^{(2)}: 1 \times 3$

- $\mathbf{a}^{(1)}: 3 \times 1$

- $\mathbf{z}^{(2)}: 1 \times 1$  (scalar)

**Common Error:** Transposed matrices. Always check dimensions!

---

Matrix dimensions must be compatible

## Worked Example: 2-3-1 Network

### Network Setup

Input:  $x = \begin{pmatrix} 0.5 \\ 0.8 \end{pmatrix}$

Layer 1 weights:

$$W^{(1)} = \begin{pmatrix} 0.2 & 0.4 \\ 0.3 & 0.1 \\ 0.5 & 0.2 \end{pmatrix}$$

$$b^{(1)} = \begin{pmatrix} 0.1 \\ -0.1 \\ 0.0 \end{pmatrix}$$

Layer 2 weights:

$$W^{(2)} = (0.6 \quad 0.3 \quad 0.4)$$

$$b^{(2)} = -0.2$$

### Forward Pass

#### Layer 1:

$$z^{(1)} = \begin{pmatrix} 0.2(0.5) + 0.4(0.8) + 0.1 \\ 0.3(0.5) + 0.1(0.8) - 0.1 \\ 0.5(0.5) + 0.2(0.8) + 0.0 \end{pmatrix} = \begin{pmatrix} 0.52 \\ 0.13 \\ 0.41 \end{pmatrix}$$

$$a^{(1)} = \text{ReLU}(z^{(1)}) = \begin{pmatrix} 0.52 \\ 0.13 \\ 0.41 \end{pmatrix}$$

#### Layer 2:

$$z^{(2)} = 0.6(0.52) + 0.3(0.13) + 0.4(0.41) - 0.2 = 0.315$$

$$\hat{y} = \sigma(0.315) = 0.578$$

**Output: 57.8% probability of class 1**

---

Following the numbers through the network

# Counting Parameters

## Parameters per Layer

For layer  $l$  with  $n_{l-1}$  inputs and  $n_l$  outputs:

**Weights:**  $n_l \times n_{l-1}$

**Biases:**  $n_l$

**Total:**  $n_l \times n_{l-1} + n_l = n_l(n_{l-1} + 1)$

## Network Total:

$$\text{Params} = \sum_{l=1}^L n_l(n_{l-1} + 1)$$

## Example: 4-10-5-1 Network

**Layer 1** ( $4 \rightarrow 10$ ):

$$10 \times 4 + 10 = 50$$

**Layer 2** ( $10 \rightarrow 5$ ):

$$5 \times 10 + 5 = 55$$

**Layer 3** ( $5 \rightarrow 1$ ):

$$1 \times 5 + 1 = 6$$

**Total: 111 parameters**

For 100 training samples: < 2 samples per parameter.  
Risk of overfitting!

---

How many weights does your network have?

*"A 4-10-5-1 network has how many parameters? Calculate and discuss: is this a lot or a little for stock prediction?"*

**Answer: 111 parameters**

**Consider:**

**Stock Data Context:**

- Daily data: ~252 days/year
- 10 years = 2,520 samples
- 111 params: 23 samples/param
- Seems okay...

**But Also Consider:**

- Financial regimes change
- Not all data equally relevant
- Need train/val/test split
- Model complexity vs data size

---

**Exercise: 3 minutes**

## A Realistic Setup


### Input Features (10):

- P/E, P/B, EV/EBITDA (value)
- 1m, 3m, 6m returns (momentum)
- 20d volatility (risk)
- Volume ratio (liquidity)
- Sector one-hot (2 features)

### Architecture:

- Hidden 1: 20 neurons (ReLU)
- Hidden 2: 10 neurons (ReLU)
- Output: 1 neuron (sigmoid)

**Total: 441 parameters**



Multiple factors combined through hidden layers