

## What is a Perceptron?

The simplest possible neural network:

- One artificial neuron
- Multiple inputs, one output
- Binary decision: Yes or No

Think of it as:

- A filter for data
- A simple classifier
- A linear decision maker

## Finance Application:

Stock screener that outputs “Buy” or “Don’t Buy” based on financial metrics.

---

A single perceptron is a stock screening filter

## Real-World Examples

### Email Spam Filter:

- Inputs: word frequencies
- Output: spam or not spam

### Loan Approval:


- Inputs: income, credit score, debt
- Output: approve or reject

### Stock Screening:

- Inputs: P/E, momentum, volume
- Output: buy or pass

All these are binary classification problems that a perceptron can solve (if the data is linearly separable).

## The Perceptron: Architecture



## Problem Setup

You want to build a simple stock screener:

- **Goal:** Decide Buy or Pass
- **Data:** Historical financial metrics
- **Method:** Perceptron classifier

## Available Features:

1. P/E Ratio (valuation)
2. 6-month momentum (%)
3. Average daily volume
4. Debt-to-Equity ratio
5. Earnings surprise (%)

## The Question

Given these features for a new stock, should we add it to our portfolio?

## Example Stock:

- P/E = 18
- Momentum = +12%
- Volume = 2M shares
- D/E = 0.8
- Surprise = +5%

## Traditional Approach:

Analyst manually weighs factors and decides.

## Perceptron Approach:

Learn the weights from historical winners/losers.

---

Given financial indicators, should we buy this stock?

# Inputs: The Raw Data

## What Are Inputs?

Each input  $x_i$  is a numerical feature:

- A measurement
- A statistic
- A signal

## In Finance:

- Price-based: returns, volatility
- Fundamental: P/E, ROE, debt ratios
- Technical: RSI, moving averages
- Sentiment: news scores, analyst ratings

## Key Requirement:

All inputs must be **numerical**. Categorical data needs encoding.

What data feeds into our decision?

## Notation

For a stock with  $n$  features:

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}$$

Example ( $n=3$ ):

$$\mathbf{x} = \begin{pmatrix} 18 \\ 0.12 \\ 0.8 \end{pmatrix} = \begin{pmatrix} \text{P/E} \\ \text{Momentum} \\ \text{D/E} \end{pmatrix}$$

**Note:** Features often need **normalization** (covered in Module 3).

# Weights: The Importance Factors

## What Are Weights?


Each weight  $w_i$  represents:

- Importance of input  $x_i$
- Direction of influence
- Learned from data

## Interpretation:

- $w_i > 0$ : Higher  $x_i$  pushes toward “Buy”
- $w_i < 0$ : Higher  $x_i$  pushes toward “Sell”
- $|w_i|$  large: Strong influence
- $|w_i|$  small: Weak influence

Weighted Sum: How Inputs Combine



Finance Example: Combining multiple factors to score a stock

$$z = w_1x_1 + w_2x_2 + w_3x_3 + b = 0.5 \times 0.8 + 0.3 \times 0.6 + 0.2 \times 0.4 + 0.1 = 0.76$$



.../..

weighted\_sum\_visualization

“Not all data is equally important” - weights encode importance

*"If you could only look at 3 metrics for a stock, which would you choose and why? How would you weight them?"*

Consider:

**Value Investor Might Choose:**

- P/E ratio ( $w = 0.5$ )
- Book value ( $w = 0.3$ )
- Dividend yield ( $w = 0.2$ )

**Growth Investor Might Choose:**

- Revenue growth ( $w = 0.5$ )
- Momentum ( $w = 0.3$ )
- Market share ( $w = 0.2$ )

**Key Insight:** Different investors would assign different weights. The perceptron *learns* these weights from historical performance.

---

Think-Pair-Share: 3 minutes

# The Weighted Sum: Adding Up Evidence

## Computing the Weighted Sum

$$z = \sum_{i=1}^n w_i x_i + b = w_1 x_1 + w_2 x_2 + \cdots + w_n x_n + b$$

### What This Means:

- Multiply each input by its weight
- Sum all the products
- Add the bias term  $b$
- Result: a single “score”

### The Bias $b$ :

- Shifts the decision threshold
- Like a “base rate” or prior
- Can be thought of as  $w_0 \cdot x_0$  where  $x_0 = 1$

---

Combine all weighted inputs into a single score

## Worked Example

### Inputs:

- $x_1 = 0.8$  (normalized P/E)
- $x_2 = 0.6$  (normalized momentum)

### Weights:

- $w_1 = 0.5$
- $w_2 = 0.7$
- $b = -0.3$

### Calculation:

$$\begin{aligned} z &= w_1 x_1 + w_2 x_2 + b \\ &= (0.5)(0.8) + (0.7)(0.6) + (-0.3) \\ &= 0.4 + 0.42 - 0.3 \\ &= \mathbf{0.52} \end{aligned}$$

## Analogy: The Voting Committee

### The Perceptron as a Committee

Member	Vote	Weight	Contribution
P/E analyst	+1	0.5	+0.5
Momentum	+1	0.7	+0.7
Bias (skeptic)	-1	0.3	-0.3
<b>Total</b>			<b>+0.9</b>

If Total > 0: Committee recommends **Buy**

#### Key Insight:

The perceptron is just a weighted voting system where the weights are learned from data.

Some votes count more than others

### Why This Works

#### Traditional Committee:

- Human experts set weights
- Based on experience/intuition
- May have biases
- Hard to scale

#### Perceptron Committee:

- Weights learned from data
- Based on historical performance
- Consistent application
- Scales to any volume

**Trade-off:** Data-driven weights may not capture regime changes or rare events.

# The Threshold: Making the Call

## The Activation Function

After computing  $z$ , we need a final decision.

### Step Function:

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$


### Interpretation:

- $z \geq 0$ : Evidence favors “Buy” → output 1
- $z < 0$ : Evidence favors “Sell” → output 0

### Why Step Function?

- Binary classification needs binary output
- Mimics neuron firing (all-or-nothing)
- Simple to implement

Above threshold = Buy, Below threshold = Sell



...

step\_functio

# The Complete Perceptron Flow

## The Pipeline


1. **Input:** Receive features  $x$
2. **Weight:** Multiply by  $w$
3. **Sum:** Add all products + bias
4. **Activate:** Apply step function
5. **Output:** Return prediction

## Compact Notation:

$$y = f(\mathbf{w}^T \mathbf{x} + b)$$

where  $\mathbf{w}^T \mathbf{x} = \sum_i w_i x_i$

## The Perceptron: Architecture



Inputs - $\downarrow$  Weights - $\downarrow$  Sum - $\downarrow$  Threshold - $\downarrow$  Decision

## What You Already Know

From the intuition section:

- Inputs are weighted
- Weights encode importance
- Sum is compared to threshold
- Output is binary

## What's Next

- Precise mathematical notation
- Geometric interpretation
- Foundation for learning algorithm

## Why Math Matters

### Without Math:

- “The network kind of learns”
- “Adjust weights somehow”
- “It works, probably”

### With Math:

- Precise learning rules
- Convergence guarantees
- Understanding of limitations

*The next 8 slides formalize what you already understand intuitively.*

---

You understand the intuition. Let's write it precisely.

# The Perceptron Equation

## Scalar Form

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right)$$

where  $f$  is the step function:

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

## Vector Form

$$y = f(\mathbf{w}^T \mathbf{x} + b)$$

where:

- $\mathbf{w} = (w_1, \dots, w_n)^T$
- $\mathbf{x} = (x_1, \dots, x_n)^T$

---

## The complete mathematical model

## Alternative Notation

We can absorb the bias into weights:

$$\tilde{\mathbf{w}} = \begin{pmatrix} b \\ w_1 \\ \vdots \\ w_n \end{pmatrix}, \quad \tilde{\mathbf{x}} = \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_n \end{pmatrix}$$

Then:

$$y = f(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}})$$

**Note:** This “bias trick” simplifies notation but they are equivalent.

# Unpacking the Mathematics

## Term by Term

Symbol	Meaning
$x_i$	Input feature $i$
$w_i$	Weight for feature $i$
$b$	Bias (threshold shift)
$z$	Weighted sum (pre-activation)
$f$	Activation function
$y$	Output prediction
$n$	Number of features

### Dimensions:

- $\mathbf{x} \in \mathbb{R}^n$
- $\mathbf{w} \in \mathbb{R}^n$
- $b, z, y \in \mathbb{R}$

## What Gets Learned?

### Learned (trainable):

- Weights  $w_1, \dots, w_n$
- Bias  $b$

### Fixed (architecture):

- Number of inputs  $n$
- Activation function  $f$

### Given (data):

- Input values  $x_1, \dots, x_n$
- Target labels (for training)

**Total Parameters:**  $n + 1$

(For a 3-feature perceptron: 4 parameters)

---

Each symbol has a meaning

# The Bias Term

## What Does Bias Do?

Without bias ( $b = 0$ ):

$$z = \mathbf{w}^T \mathbf{x}$$

The decision boundary passes through origin.

With bias ( $b \neq 0$ ):

$$z = \mathbf{w}^T \mathbf{x} + b$$

The decision boundary can be anywhere.

## Interpretation:

- $b > 0$ : Default toward “Buy”
- $b < 0$ : Default toward “Sell”
- Like a prior belief

## Finance Analogy

**Without Bias:**

“I have no opinion until I see data”

**With Positive Bias:**

“I’m generally bullish; you need to convince me to sell”

**With Negative Bias:**

“I’m skeptical by default; you need strong evidence to buy”

**Key Point:** Bias shifts the “bar” that evidence must clear. It’s learned from data just like weights.

---

Bias shifts the decision threshold

# The Step Activation Function

## Formal Definition

The Heaviside step function:

$$f(z) = \mathbf{1}_{z \geq 0} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$


## Properties:

- Output  $\in \{0, 1\}$
- Discontinuous at  $z = 0$
- Not differentiable (problem for gradient-based learning!)

## Variants:

- Sign function: outputs  $\{-1, +1\}$
- Same idea, different labels

Binary output: yes or no



.../.



**Preview:** The non-differentiability of the step function is why we'll need smoother activations (sigmoid, ReLU) in later modules.

# Geometric Interpretation: The Decision Boundary

## The Perceptron as a Hyperplane

The equation  $w^T x + b = 0$  defines a hyperplane:


- In 2D: a line
- In 3D: a plane
- In  $nD$ : a hyperplane

### Regions:

- $w^T x + b > 0$ : Class 1 (Buy)
- $w^T x + b < 0$ : Class 0 (Sell)
- $w^T x + b = 0$ : Decision boundary

### Weight Vector Direction:

$w$  is perpendicular to the decision boundary, pointing toward the positive class.



# Finance Example: Classifying Stocks

## Two-Feature Stock Screener

Features:

- $x_1$ : P/E ratio (normalized)
- $x_2$ : 6-month momentum (%)

Classes:


- Green: Outperformed (Buy)
- Red: Underperformed (Sell)

Goal:

Find  $w_1, w_2, b$  such that:

$$w_1 \cdot \text{P/E} + w_2 \cdot \text{Momentum} + b = 0$$

separates the classes.



Separating “good” stocks from “bad” stocks

# The Decision Boundary Formula

## In 2D: The Line Equation

From  $w_1x_1 + w_2x_2 + b = 0$ :

$$x_2 = -\frac{w_1}{w_2}x_1 - \frac{b}{w_2}$$

This is a line with:

- Slope:  $-\frac{w_1}{w_2}$
- Intercept:  $-\frac{b}{w_2}$


Example:

If  $w_1 = 2, w_2 = 1, b = -3$ :

$$x_2 = -2x_1 + 3$$

Stocks above this line: Buy

Stocks below this line: Sell



The line that separates buy from sell

# How Does the Perceptron Learn?

## The Learning Problem

Given:

- Training data:  $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^m$
- Each  $\mathbf{x}^{(i)}$ : feature vector
- Each  $y^{(i)} \in \{0, 1\}$ : true label

Find:

- Weights  $\mathbf{w}$
- Bias  $b$
- Such that predictions match labels

## The Approach:

Start with random weights, then iteratively adjust based on mistakes.

## The Core Idea

If prediction is correct:

Do nothing. Weights are fine.

If prediction is wrong:

Adjust weights to make this example more likely to be correct next time.

Repeat:

Keep cycling through training data until no mistakes (or convergence).

**Key Insight:** Learning = adjusting weights based on errors.

---

Learning = adjusting weights based on mistakes

## Two Types of Errors

**False Negative** ( $\hat{y} = 0, y = 1$ ):

- Predicted Sell, should be Buy
- The score  $z$  was too low
- Need to *increase* score for this  $x$
- Solution: Add  $x$  to  $w$

**False Positive** ( $\hat{y} = 1, y = 0$ ):

- Predicted Buy, should be Sell
- The score  $z$  was too high
- Need to *decrease* score for this  $x$
- Solution: Subtract  $x$  from  $w$

## Visual Intuition

**Before update:**

Point is on wrong side of boundary.

**After update:**

Boundary moves to include the point on the correct side.

**The Update Rule:**

$$w_{\text{new}} = w_{\text{old}} + (y - \hat{y}) \cdot x$$

**Check:**

- If  $y = 1, \hat{y} = 0$ : add  $x$
- If  $y = 0, \hat{y} = 1$ : subtract  $x$
- If  $y = \hat{y}$ : no change

---

Each mistake is a learning opportunity

## Why Adding $x$ Works


For a false negative (missed Buy):

- Current:  $\mathbf{w}^T \mathbf{x} + b < 0$
- After adding  $x$  to  $w$ :
- New score:  $(\mathbf{w} + \mathbf{x})^T \mathbf{x} + b$
- $= \mathbf{w}^T \mathbf{x} + \mathbf{x}^T \mathbf{x} + b$
- $= \mathbf{w}^T \mathbf{x} + \|\mathbf{x}\|^2 + b$

Since  $\|\mathbf{x}\|^2 > 0$ , the new score is higher!

**Geometrically:**

Adding  $x$  rotates the decision boundary toward classifying  $x$  correctly.



If wrong, move the boundary

## The Update Equations

For each training example  $(\mathbf{x}, y)$ :

**Weight update:**

$$\mathbf{w} \leftarrow \mathbf{w} + \eta(y - \hat{y})\mathbf{x}$$

**Bias update:**

$$b \leftarrow b + \eta(y - \hat{y})$$

where:

- $\eta > 0$  is the learning rate
- $\hat{y} = f(\mathbf{w}^T \mathbf{x} + b)$  is prediction
- $y$  is true label

## The Complete Algorithm

1. Initialize  $\mathbf{w} = \mathbf{0}$ ,  $b = 0$

2. repeat:

- a. For each  $(\mathbf{x}^{(i)}, y^{(i)})$  in training set:
- b. Compute  $\hat{y}^{(i)} = f(\mathbf{w}^T \mathbf{x}^{(i)} + b)$
- c. If  $\hat{y}^{(i)} \neq y^{(i)}$ :  
 $\mathbf{w} \leftarrow \mathbf{w} + \eta(y^{(i)} - \hat{y}^{(i)})\mathbf{x}^{(i)}$   
 $b \leftarrow b + \eta(y^{(i)} - \hat{y}^{(i)})$

3. until no errors (or max iterations)

---

The mathematical update rule

## What is $\eta$ ?

The learning rate controls step size:

- How much weights change per update
- Typical values: 0.01 to 1.0
- For perceptron: often  $\eta = 1$

## Effects:

### $\eta$ too small:

- Very slow learning
- Many iterations needed
- But stable

### $\eta$ too large:

- May overshoot
- Oscillate around solution
- But faster initially

---

Step size matters: too big or too small both cause problems

## For the Perceptron

### Good news:

For linearly separable data, the perceptron converges regardless of  $\eta > 0$ .

### Why?

The convergence theorem (next slides) guarantees finding a solution if one exists.

### In Practice:

$\eta = 1$  is common for perceptron. Learning rate matters more for:

- Gradient descent (Module 3)
- Non-separable data
- Multi-layer networks

# Worked Example: Stock Classification

## Setup

Two stocks, two features:

- $\mathbf{x}^{(1)} = (0.5, 0.8)$ ,  $y^{(1)} = 1$  (Buy)
- $\mathbf{x}^{(2)} = (0.2, 0.3)$ ,  $y^{(2)} = 0$  (Sell)

Initialize:  $\mathbf{w} = (0, 0)$ ,  $b = 0$ ,  $\eta = 1$

## Iteration 1: Example 1

- $z = 0 \cdot 0.5 + 0 \cdot 0.8 + 0 = 0$
- $\hat{y} = f(0) = 1$  (threshold at 0)
- $y = 1$ , correct! No update.

## Iteration 1: Example 2

- $z = 0$ ,  $\hat{y} = 1$
- $y = 0$ , wrong!
- $\mathbf{w} \leftarrow (0, 0) + 1(0 - 1)(0.2, 0.3) = (-0.2, -0.3)$
- $b \leftarrow 0 + 1(0 - 1) = -1$

## Iteration 2: Example 1

- $z = -0.2(0.5) - 0.3(0.8) - 1 = -1.34$
- $\hat{y} = 0$
- $y = 1$ , wrong!
- $\mathbf{w} \leftarrow (-0.2, -0.3) + (0.5, 0.8) = (0.3, 0.5)$
- $b \leftarrow -1 + 1 = 0$

## Iteration 2: Example 2

- $z = 0.3(0.2) + 0.5(0.3) + 0 = 0.21$
- $\hat{y} = 1$ ,  $y = 0$ , wrong!
- $\mathbf{w} \leftarrow (0.3, 0.5) - (0.2, 0.3) = (0.1, 0.2)$
- $b \leftarrow 0 - 1 = -1$

Continue until convergence...

---

Following the math with real numbers

# Convergence: Does It Always Work?

## The Perceptron Convergence Theorem

**Theorem (Rosenblatt, 1962):**

If the training data is **linearly separable**, the perceptron learning algorithm will find a separating hyperplane in a **finite** number of updates.


### Key Conditions:

- Data must be linearly separable
- Learning rate  $\eta > 0$
- Cycling through all examples

### Bound on Updates:

$$\text{mistakes} \leq \frac{R^2}{\gamma^2}$$

where  $R = \max \text{ norm}$ ,  $\gamma = \text{margin}$



The perceptron convergence theorem guarantees finding a solution IF one exists

*"What happens when data isn't linearly separable in financial markets? Can you think of examples?"*

Consider:

**Examples of Non-Separable Data:**

- High P/E growth stocks AND low P/E value stocks both outperform
- Medium-risk investments underperform both conservative and aggressive
- “Buy the rumor, sell the news” patterns

**What Happens to the Perceptron?**

- Never converges
- Oscillates forever
- Best we can do: minimize errors
- Need something more powerful...

**Foreshadowing:** This is exactly why we need **multi-layer** networks (Module 2).

---

Think-Pair-Share: 3 minutes

# The XOR Problem

## The Exclusive OR Function

$x_1$	$x_2$	XOR
0	0	0
0	1	1
1	0	1
1	1	0

### In Words:

Output is 1 if inputs are *different*, 0 if inputs are *same*.

### The Challenge:

Try to draw a single line that separates the 1s from the 0s...


Some patterns cannot be separated by a single line

### The XOR Problem: Why Single-Layer Perceptrons Fail

XOR Truth Table

$x_1$	$x_2$	XOR	Output
0	0	0 XOR 0	0
0	1	0 XOR 1	1
1	0	1 XOR 0	1
1	1	1 XOR 1	0

"Same inputs = 0, Different inputs = 1"



xor\_problem

# Why XOR Cannot Be Solved

## Geometric Impossibility

Perceptron decision boundary:

$$w_1x_1 + w_2x_2 + b = 0$$

This is always a **straight line**.

XOR requires:

A boundary that curves or has multiple segments.

Linear vs Non-Linear


Linearly Separable:

- AND, OR, NAND, NOR
- One line can separate

Not Linearly Separable:

- XOR, XNOR
- No single line works

No single hyperplane can separate XOR



linear\_vs\_nonlinear\_pattern

# 1969: The Critique That Changed Everything

## Minsky and Papert's Book

"Perceptrons: An Introduction to Computational Geometry" (1969)

### Key Arguments:

1. Single-layer perceptrons cannot compute XOR
2. Many important functions are non-linear
3. No known training algorithm for multi-layer networks
4. Scaling limitations

### The Impact:

The book was rigorous and influential. It convinced funding agencies that neural networks were a dead end.

## The Controversy

### Valid Points:

- Single layers are limited
- XOR problem is real
- No training algorithm existed (then)

### Overstated Points:

- "Neural networks can't work"
- Implied multi-layer networks wouldn't help
- Discouraged research for 15+ years

**Lesson:** Valid criticism of current methods shouldn't stop research into future improvements.

---

Marvin Minsky and Seymour Papert: "Perceptrons" book

# The First AI Winter Begins

## The Collapse


After 1969:

- Funding dried up
- Researchers left the field
- “Neural networks don't work”
- Symbolic AI took over

Duration: 1969 to ~1982

## What Survived:

- A few dedicated researchers
- Theoretical work continued quietly
- Hopfield networks (1982)
- Backpropagation (1986)



.../.

ai\_winter\_timeline

---

1969-1982: The dark ages of neural network research

# Module 1: Key Takeaways

## What We Learned

### 1. Historical Foundation


- McCulloch-Pitts (1943): neurons compute
- Hebb (1949): learning strengthens connections
- Rosenblatt (1958): perceptron learns

### 2. The Perceptron Model

- Weighted sum + threshold
- Linear decision boundary
- Learns from mistakes

### 3. Limitations

- Only linearly separable problems
- XOR is impossible
- Led to AI Winter



From biological inspiration to mathematical limitation

*“What if we stack multiple perceptrons?”*

### The Problem We Face

Single perceptrons can only solve linearly separable problems. Real financial data is rarely that simple.

### The Solution Preview:

- Add “hidden” layers
- Non-linear activation functions
- Multi-Layer Perceptrons (MLPs)

### Coming in Module 2:

- How XOR gets solved
- MLP architecture
- Activation functions (sigmoid, ReLU)
- Universal Approximation Theorem
- Loss functions

**Spoiler:** Adding just one hidden layer changes everything.

**Mathematical details for this module: See Appendix A (Perceptron Convergence Proof)**

---

Next: Solving XOR with Multi-Layer Perceptrons