

# Module 1: The Birth of Neural Computing

## From Biological Inspiration to the Perceptron (1943-1969)

Neural Networks for Finance

BSc Lecture Series

November 26, 2025

## How Does a Committee Make Decisions?

Imagine an investment committee evaluating a stock:

- **Analyst A:** "Strong earnings growth" (+1 vote)
- **Analyst B:** "High debt levels" (-1 vote)
- **Analyst C:** "Good momentum" (+1 vote)
- **Senior Partner:** "Market risk is elevated" (-2 votes)

## The Decision Process:

1. Gather evidence from each analyst
2. Weight opinions by seniority/expertise
3. Sum the weighted votes
4. If total > threshold: **Buy**

## Weighted Voting

Analyst	Vote	Weight
Analyst A	+1	1.0
Analyst B	-1	1.0
Analyst C	+1	1.0
Senior Partner	-1	2.0
<b>Weighted Sum</b>	<b>-1.0</b>	

**Decision: Don't Buy**

---

**Finance Hook:** This is exactly how a perceptron works!

# What If Machines Could Decide?

## The Central Question

In 1943, scientists asked:

*"Can we build a machine that learns to make decisions like a brain?"*

## Why This Matters for Finance:

- Humans are slow and biased
- Markets process millions of data points
- Pattern recognition at scale
- Consistent, emotionless decisions

## The Promise

If we could capture how neurons compute:

- Automatic stock screening
- Risk assessment at scale
- Pattern detection in market data
- Learning from historical decisions

## The Challenge

How do we translate biological processes into mathematical operations?

*This module tells the story of how scientists attempted this translation.*

---

The fundamental question that started neural network research

## The Complete Journey (4 Modules)

### 1. The Perceptron (Today)

- Single neuron foundations
- 1943-1969 history

### 2. Multi-Layer Perceptrons

- Stacking layers, activation functions

### 3. Training Neural Networks

- Backpropagation, optimization

### 4. Applications in Finance

- Stock prediction case study

## Today's Module Structure

### 1. Historical Context (1943-1969)

- McCulloch-Pitts, Hebb, Rosenblatt

### 2. Biological Inspiration

- From neurons to mathematics

### 3. The Perceptron

- Intuition, then math

### 4. Learning Algorithm

- How it adjusts weights

### 5. Limitations

- XOR problem, AI Winter

---

Your journey through neural network fundamentals

By the end of this module, you will be able to:

## 1. Understand biological inspiration

- How real neurons inspired artificial ones
- What we kept and what we simplified

## 2. Master the perceptron model

- Inputs, weights, sum, activation
- The decision-making unit

## 3. Interpret decision boundaries

- Geometric meaning of weights
- Linear separability concept

## 4. Apply the learning algorithm

- Weight update rule
- Convergence conditions

## 5. Recognize limitations

- XOR problem
- Why single layers are not enough

**Finance Connection:** Throughout, we'll use stock classification as our running example.

---

By the end of this module, you will be able to...

# 1943: The Mathematical Neuron

## Warren McCulloch & Walter Pitts

In 1943, a neurophysiologist and a logician asked:

*"Can we describe what neurons do using mathematics?"*

Their paper: "A Logical Calculus of Ideas Immanent in Nervous Activity"

### Key Insight:

- Neurons have binary states (fire or not)
- This is like TRUE/FALSE in logic
- Networks of neurons can compute any logical function

charts/mcculloch\_pitts\_diagram/mcculloch\_pitts\_dia

Warren McCulloch and Walter Pitts: "A Logical Calculus of Ideas Immanent in Nervous Activity"

## What McCulloch & Pitts Proposed

The brain performs computation through:

### 1. Binary Signals

- Neurons either fire (1) or don't (0)
- Like bits in a computer

### 2. Threshold Logic

- Sum of inputs exceeds threshold  $\rightarrow$  fire
- Otherwise  $\rightarrow$  stay quiet

### 3. Network Composition

- Complex behaviors from simple units
- AND, OR, NOT gates from neurons

## Logical Operations with Neurons

**AND Gate** (threshold = 2):

- Both inputs = 1  $\rightarrow$  output = 1
- Otherwise  $\rightarrow$  output = 0

**OR Gate** (threshold = 1):

- Any input = 1  $\rightarrow$  output = 1
- All inputs = 0  $\rightarrow$  output = 0

**Implication:** If neurons compute logic, and computers compute logic, then we can build artificial brains!

---

If neurons compute, can we build artificial ones?

# 1949: Hebbian Learning

## Donald Hebb's Insight

McCulloch-Pitts neurons were fixed. But how does the brain *learn*?

### Hebb's Rule (1949):

*"Neurons that fire together, wire together."*

#### In Plain Terms:

- If neuron A consistently activates neuron B
- The connection  $A \rightarrow B$  grows stronger
- Repeated patterns reinforce pathways

#### Finance Analogy:

An analyst who repeatedly identifies winning stocks gains more influence in the committee.

charts/hebb\_learning\_visualization/hebb\_learning\_v

Donald Hebb: "Neurons that fire together, wire together"

# 1958: The Perceptron is Born

## Frank Rosenblatt at Cornell

Combined McCulloch-Pitts neurons with Hebbian learning into a machine that could *learn from examples*.

### The Perceptron:

- A single artificial neuron
- Adjustable connection weights
- Learns to classify patterns
- Implemented in hardware (Mark I)

### Key Innovation:

Not just fixed logic gates, but a system that **learns** the right weights from training data.

charts/mark1\_perceptron\_diagram/mark1\_perceptron\_d

# The New York Times Headline

July 8, 1958 - The New York Times

*"New Navy Device Learns By Doing; Psychologist Shows Embryo of Computer Designed to Read and Grow Wiser"*

## The Promises Made:

- Machines that recognize faces
- Automatic translation of languages
- Systems that “perceive” like humans
- The Navy predicted: walking, talking, self-reproducing machines

## The Reality:

The perceptron could classify simple patterns, but the gap between promise and capability was vast.

## Lessons for Today

### Sound Familiar?

- “AI will replace all jobs”
- “Machines will be smarter than humans by 20XX”
- “This changes everything”

### Pattern:

1. Genuine breakthrough
2. Media amplification
3. Overpromising
4. Disappointment
5. “AI Winter”

*History repeats...*

---

“New Navy Device Learns By Doing” - The hype cycle begins

## Timeline: The Early Years

[charts/timeline\\_1943\\_1969/timeline\\_1943\\_1969.pdf](#)

*"The perceptron was funded by the US Navy for military applications. How does funding source shape research direction? Are there parallels in modern AI development?"*

Consider:

- Military vs. commercial vs. academic funding
- What problems get prioritized?
- Open vs. closed research
- Today: Tech giants fund most AI research
- Government initiatives (CHIPS Act, etc.)
- Startup ecosystem influence

---

Think-Pair-Share: 3 minutes

## Anatomy of a Real Neuron

### 1. Dendrites (Input)

- Tree-like branches
- Receive signals from other neurons
- Thousands of connections

### 2. Cell Body (Soma) (Processing)

- Integrates incoming signals
- Contains the nucleus
- Determines if neuron fires

### 3. Axon (Output)

- Long fiber carrying output signal
- Connects to other neurons
- All-or-nothing signal

## How It Works

1. Signals arrive at dendrites
2. Soma sums the inputs
3. If sum exceeds threshold: neuron **fires**
4. Action potential travels down axon
5. Signal reaches next neurons

## Key Numbers:

- Human brain: ~86 billion neurons
- Each neuron: ~7,000 connections
- Total synapses: ~100 trillion

---

Dendrites receive, soma processes, axon transmits

## Mathematical Abstraction

### 1. Inputs ( $x_1, x_2, \dots, x_n$ )

- Numerical values (features)
- Replace dendrites

### 2. Weights ( $w_1, w_2, \dots, w_n$ )

- Importance of each input
- Replace synapse strength

### 3. Weighted Sum

- $z = \sum_{i=1}^n w_i x_i + b$
- Replace soma integration

### 4. Activation Function

- $y = f(z)$
- Replace firing decision

## The Complete Model

$$y = f\left(\sum_{i=1}^n w_i x_i + b\right)$$

### Components:

- $x_i$ : Input features
- $w_i$ : Learnable weights
- $b$ : Bias (threshold adjustment)
- $f$ : Activation function
- $y$ : Output (prediction)

**Key Point:** The weights are what the network *learns*.

---

From biology to mathematics: the abstraction trade-off

## Biological vs. Artificial: Side by Side

[charts/biological\\_vs\\_artificial\\_neuron/biological\\_vs\\_artificial\\_neuron.pdf](#)

## A Financial Analyst as a Neuron

Biology	Finance
Dendrites	Market data feeds
Synapses	Data reliability weights
Soma	Analyst's judgment
Threshold	Conviction level
Axon	"Buy" recommendation

### The Process:

1. Receive multiple data points
2. Weight by source quality
3. Aggregate into overall view
4. If conviction > threshold: recommend

## Example: Stock Screening

### Inputs (Data):

- $x_1$ : P/E ratio = 15
- $x_2$ : Revenue growth = 20%
- $x_3$ : Debt/Equity = 0.5

### Weights (Importance):

- $w_1 = 0.3$  (value focus)
- $w_2 = 0.5$  (growth priority)
- $w_3 = -0.2$  (debt penalty)

### Decision:

$$z = 0.3(15) + 0.5(20) - 0.2(0.5) = 14.4$$

If  $z > 10$ : **Buy**

---

Inputs (data) -  $\downarrow$  Weights (importance) -  $\downarrow$  Decision (output)

## Benefits of Simplification

### 1. Mathematical Tractability

- We can write equations
- Analyze behavior formally
- Prove theorems

### 2. Computability

- Easy to implement in code
- Fast computation
- Scales to millions of units

### 3. Trainability

- Can adjust weights systematically
- Gradient-based optimization
- Learn from data

## What We Can Now Do

- Define learning algorithms
- Compute exact outputs
- Train on historical data
- Make predictions on new data
- Analyze decision boundaries

## Scale Comparison:

	Brain	GPU
Operations/sec	$10^{16}$	$10^{15}$
Power	20W	300W
Training time	Years	Hours

Different trade-offs, different capabilities.

---

Simplification enables computation

## Biological Complexity We Ignored

### 1. Temporal Dynamics

- Real neurons have timing
- Spike patterns carry information
- We use static activations

### 2. Structural Complexity

- Dendrites have local computation
- Different neuron types
- We use uniform units

### 3. Neurochemistry

- Neurotransmitters vary
- Modulatory systems
- We use simple multiplication

## Implications

### What ANNs Cannot Do (Well):

- Energy efficiency of brain
- One-shot learning
- Continuous adaptation
- Common sense reasoning

### The Trade-off:



*Artificial neurons are inspired by biology, not copies of it.*

---

The brain does far more than our models capture

## What is a Perceptron?

The simplest possible neural network:

- One artificial neuron
- Multiple inputs, one output
- Binary decision: Yes or No

Think of it as:

- A filter for data
- A simple classifier
- A linear decision maker

## Finance Application:

Stock screener that outputs “Buy” or “Don’t Buy” based on financial metrics.

---

A single perceptron is a stock screening filter

## Real-World Examples

### Email Spam Filter:

- Inputs: word frequencies
- Output: spam or not spam

### Loan Approval:

- Inputs: income, credit score, debt
- Output: approve or reject

### Stock Screening:

- Inputs: P/E, momentum, volume
- Output: buy or pass

All these are binary classification problems that a perceptron can solve (if the data is linearly separable).

`charts/perceptron_architecture/perceptron_architecture.pdf`

## Problem Setup

You want to build a simple stock screener:

- **Goal:** Decide Buy or Pass
- **Data:** Historical financial metrics
- **Method:** Perceptron classifier

## Available Features:

1. P/E Ratio (valuation)
2. 6-month momentum (%)
3. Average daily volume
4. Debt-to-Equity ratio
5. Earnings surprise (%)

## The Question

Given these features for a new stock, should we add it to our portfolio?

## Example Stock:

- P/E = 18
- Momentum = +12%
- Volume = 2M shares
- D/E = 0.8
- Surprise = +5%

## Traditional Approach:

Analyst manually weighs factors and decides.

## Perceptron Approach:

Learn the weights from historical winners/losers.

---

Given financial indicators, should we buy this stock?

# Inputs: The Raw Data

## What Are Inputs?

Each input  $x_i$  is a numerical feature:

- A measurement
- A statistic
- A signal

## In Finance:

- Price-based: returns, volatility
- Fundamental: P/E, ROE, debt ratios
- Technical: RSI, moving averages
- Sentiment: news scores, analyst ratings

## Key Requirement:

All inputs must be **numerical**. Categorical data needs encoding.

---

What data feeds into our decision?

## Notation

For a stock with  $n$  features:

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}$$

Example ( $n=3$ ):

$$\mathbf{x} = \begin{pmatrix} 18 \\ 0.12 \\ 0.8 \end{pmatrix} = \begin{pmatrix} \text{P/E} \\ \text{Momentum} \\ \text{D/E} \end{pmatrix}$$

**Note:** Features often need **normalization** (covered in Module 3).

## What Are Weights?

Each weight  $w_i$  represents:

- Importance of input  $x_i$
- Direction of influence
- Learned from data

## Interpretation:

- $w_i > 0$ : Higher  $x_i$  pushes toward "Buy"
- $w_i < 0$ : Higher  $x_i$  pushes toward "Sell"
- $|w_i|$  large: Strong influence
- $|w_i|$  small: Weak influence

charts/weighted\_sum\_visualization/weighted\_sum\_vis

"Not all data is equally important" - weights encode importance

*"If you could only look at 3 metrics for a stock, which would you choose and why? How would you weight them?"*

Consider:

**Value Investor Might Choose:**

- P/E ratio ( $w = 0.5$ )
- Book value ( $w = 0.3$ )
- Dividend yield ( $w = 0.2$ )

**Growth Investor Might Choose:**

- Revenue growth ( $w = 0.5$ )
- Momentum ( $w = 0.3$ )
- Market share ( $w = 0.2$ )

**Key Insight:** Different investors would assign different weights. The perceptron *learns* these weights from historical performance.

---

Think-Pair-Share: 3 minutes

# The Weighted Sum: Adding Up Evidence

## Computing the Weighted Sum

$$z = \sum_{i=1}^n w_i x_i + b = w_1 x_1 + w_2 x_2 + \cdots + w_n x_n + b$$

### What This Means:

- Multiply each input by its weight
- Sum all the products
- Add the bias term  $b$
- Result: a single “score”

### The Bias $b$ :

- Shifts the decision threshold
- Like a “base rate” or prior
- Can be thought of as  $w_0 \cdot x_0$  where  $x_0 = 1$

---

Combine all weighted inputs into a single score

## Worked Example

### Inputs:

- $x_1 = 0.8$  (normalized P/E)
- $x_2 = 0.6$  (normalized momentum)

### Weights:

- $w_1 = 0.5$
- $w_2 = 0.7$
- $b = -0.3$

### Calculation:

$$\begin{aligned} z &= w_1 x_1 + w_2 x_2 + b \\ &= (0.5)(0.8) + (0.7)(0.6) + (-0.3) \\ &= 0.4 + 0.42 - 0.3 \\ &= \mathbf{0.52} \end{aligned}$$

## Analogy: The Voting Committee

### The Perceptron as a Committee

Member	Vote	Weight	Contribution
P/E analyst	+1	0.5	+0.5
Momentum	+1	0.7	+0.7
Bias (skeptic)	-1	0.3	-0.3
<b>Total</b>			<b>+0.9</b>

If Total > 0: Committee recommends **Buy**

#### Key Insight:

The perceptron is just a weighted voting system where the weights are learned from data.

Some votes count more than others

### Why This Works

#### Traditional Committee:

- Human experts set weights
- Based on experience/intuition
- May have biases
- Hard to scale

#### Perceptron Committee:

- Weights learned from data
- Based on historical performance
- Consistent application
- Scales to any volume

**Trade-off:** Data-driven weights may not capture regime changes or rare events.

# The Threshold: Making the Call

## The Activation Function

After computing  $z$ , we need a final decision.

### Step Function:

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

### Interpretation:


- $z \geq 0$ : Evidence favors "Buy" → output 1
- $z < 0$ : Evidence favors "Sell" → output 0

### Why Step Function?

- Binary classification needs binary output
- Mimics neuron firing (all-or-nothing)
- Simple to implement

---

Above threshold = Buy, Below threshold = Sell



## The Pipeline

1. **Input:** Receive features  $\mathbf{x}$
2. **Weight:** Multiply by  $\mathbf{w}$
3. **Sum:** Add all products + bias
4. **Activate:** Apply step function
5. **Output:** Return prediction

## Compact Notation:

$$y = f(\mathbf{w}^T \mathbf{x} + b)$$

where  $\mathbf{w}^T \mathbf{x} = \sum_i w_i x_i$

[charts/perceptron\\_architecture/perceptron\\_architecture.pdf](charts/perceptron_architecture/perceptron_architecture.pdf)

## What You Already Know

From the intuition section:

- Inputs are weighted
- Weights encode importance
- Sum is compared to threshold
- Output is binary

## What's Next

- Precise mathematical notation
- Geometric interpretation
- Foundation for learning algorithm

## Why Math Matters

### Without Math:

- “The network kind of learns”
- “Adjust weights somehow”
- “It works, probably”

### With Math:

- Precise learning rules
- Convergence guarantees
- Understanding of limitations

*The next 8 slides formalize what you already understand intuitively.*

---

You understand the intuition. Let's write it precisely.

# The Perceptron Equation

## Scalar Form

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right)$$

where  $f$  is the step function:

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

## Vector Form

$$y = f(\mathbf{w}^T \mathbf{x} + b)$$

where:

- $\mathbf{w} = (w_1, \dots, w_n)^T$
- $\mathbf{x} = (x_1, \dots, x_n)^T$

---

## The complete mathematical model

## Alternative Notation

We can absorb the bias into weights:

$$\tilde{\mathbf{w}} = \begin{pmatrix} b \\ w_1 \\ \vdots \\ w_n \end{pmatrix}, \quad \tilde{\mathbf{x}} = \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_n \end{pmatrix}$$

Then:

$$y = f(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}})$$

**Note:** This “bias trick” simplifies notation but they are equivalent.

# Unpacking the Mathematics

## Term by Term

Symbol	Meaning
$x_i$	Input feature $i$
$w_i$	Weight for feature $i$
$b$	Bias (threshold shift)
$z$	Weighted sum (pre-activation)
$f$	Activation function
$y$	Output prediction
$n$	Number of features

### Dimensions:

- $\mathbf{x} \in \mathbb{R}^n$
- $\mathbf{w} \in \mathbb{R}^n$
- $b, z, y \in \mathbb{R}$

Each symbol has a meaning

## What Gets Learned?

### Learned (trainable):

- Weights  $w_1, \dots, w_n$
- Bias  $b$

### Fixed (architecture):

- Number of inputs  $n$
- Activation function  $f$

### Given (data):

- Input values  $x_1, \dots, x_n$
- Target labels (for training)

**Total Parameters:**  $n + 1$

(For a 3-feature perceptron: 4 parameters)

# The Bias Term

## What Does Bias Do?

Without bias ( $b = 0$ ):

$$z = \mathbf{w}^T \mathbf{x}$$

The decision boundary passes through origin.

With bias ( $b \neq 0$ ):

$$z = \mathbf{w}^T \mathbf{x} + b$$

The decision boundary can be anywhere.

## Interpretation:

- $b > 0$ : Default toward "Buy"
- $b < 0$ : Default toward "Sell"
- Like a prior belief

---

Bias shifts the decision threshold

## Finance Analogy

### Without Bias:

"I have no opinion until I see data"

### With Positive Bias:

"I'm generally bullish; you need to convince me to sell"

### With Negative Bias:

"I'm skeptical by default; you need strong evidence to buy"

**Key Point:** Bias shifts the "bar" that evidence must clear. It's learned from data just like weights.

# The Step Activation Function

## Formal Definition

The Heaviside step function:


$$f(z) = \mathbf{1}_{z \geq 0} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Properties:

- Output  $\in \{0, 1\}$
- Discontinuous at  $z = 0$
- Not differentiable (problem for gradient-based learning!)

## Variants:

- Sign function: outputs  $\{-1, +1\}$
- Same idea, different labels



**Preview:** The non-differentiability of the step function is why we'll need smoother activations (sigmoid, ReLU) in later modules.

---

Binary output: yes or no

## The Perceptron as a Hyperplane

The equation  $\mathbf{w}^T \mathbf{x} + b = 0$  defines a hyperplane:


- In 2D: a line
- In 3D: a plane
- In  $n$ D: a hyperplane

## Regions:

- $\mathbf{w}^T \mathbf{x} + b > 0$ : Class 1 (Buy)
- $\mathbf{w}^T \mathbf{x} + b < 0$ : Class 0 (Sell)
- $\mathbf{w}^T \mathbf{x} + b = 0$ : Decision boundary

## Weight Vector Direction:

$\mathbf{w}$  is perpendicular to the decision boundary, pointing toward the positive class.



`charts/decision_boundary_2d/decision_boundary_2d.pdf`

# Finance Example: Classifying Stocks

## Two-Feature Stock Screener

Features:

- $x_1$ : P/E ratio (normalized)
- $x_2$ : 6-month momentum (%)

Classes:

- **Green**: Outperformed (Buy)
- **Red**: Underperformed (Sell)

Goal:

Find  $w_1, w_2, b$  such that:

$$w_1 \cdot \text{P/E} + w_2 \cdot \text{Momentum} + b = 0$$

separates the classes.



charts/stock\_features\_scatter/stock\_features\_scatter.p

## In 2D: The Line Equation

From  $w_1x_1 + w_2x_2 + b = 0$ :

$$x_2 = -\frac{w_1}{w_2}x_1 - \frac{b}{w_2}$$

This is a line with:

- Slope:  $-\frac{w_1}{w_2}$
- Intercept:  $-\frac{b}{w_2}$

Example:

If  $w_1 = 2, w_2 = 1, b = -3$ :

$$x_2 = -2x_1 + 3$$

Stocks above this line: Buy

Stocks below this line: Sell

charts/finance\_decision\_boundary/finance\_decision\_b

# How Does the Perceptron Learn?

## The Learning Problem

Given:

- Training data:  $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^m$
- Each  $\mathbf{x}^{(i)}$ : feature vector
- Each  $y^{(i)} \in \{0, 1\}$ : true label

Find:

- Weights  $\mathbf{w}$
- Bias  $b$
- Such that predictions match labels

## The Approach:

Start with random weights, then iteratively adjust based on mistakes.

## The Core Idea

If prediction is correct:

Do nothing. Weights are fine.

If prediction is wrong:

Adjust weights to make this example more likely to be correct next time.

Repeat:

Keep cycling through training data until no mistakes (or convergence).

**Key Insight:** Learning = adjusting weights based on errors.

---

Learning = adjusting weights based on mistakes

## Two Types of Errors

**False Negative** ( $\hat{y} = 0, y = 1$ ):

- Predicted Sell, should be Buy
- The score  $z$  was too low
- Need to *increase* score for this  $x$
- Solution: Add  $x$  to  $w$

**False Positive** ( $\hat{y} = 1, y = 0$ ):

- Predicted Buy, should be Sell
- The score  $z$  was too high
- Need to *decrease* score for this  $x$
- Solution: Subtract  $x$  from  $w$

## Visual Intuition

**Before update:**

Point is on wrong side of boundary.

**After update:**

Boundary moves to include the point on the correct side.

**The Update Rule:**

$$w_{\text{new}} = w_{\text{old}} + (y - \hat{y}) \cdot x$$

**Check:**

- If  $y = 1, \hat{y} = 0$ : add  $x$
- If  $y = 0, \hat{y} = 1$ : subtract  $x$
- If  $y = \hat{y}$ : no change

---

Each mistake is a learning opportunity

## Why Adding $x$ Works


For a false negative (missed Buy):

- Current:  $\mathbf{w}^T \mathbf{x} + b < 0$
- After adding  $x$  to  $\mathbf{w}$ :
- New score:  $(\mathbf{w} + x)^T \mathbf{x} + b$
- $= \mathbf{w}^T \mathbf{x} + x^T \mathbf{x} + b$
- $= \mathbf{w}^T \mathbf{x} + \|x\|^2 + b$

Since  $\|x\|^2 > 0$ , the new score is higher!

## Geometrically:

Adding  $x$  rotates the decision boundary toward classifying  $x$  correctly.



charts/perceptron\_learning\_animation/perceptron\_learnin

## The Update Equations

For each training example  $(\mathbf{x}, y)$ :

**Weight update:**

$$\mathbf{w} \leftarrow \mathbf{w} + \eta(y - \hat{y})\mathbf{x}$$

**Bias update:**

$$b \leftarrow b + \eta(y - \hat{y})$$

where:

- $\eta > 0$  is the learning rate
- $\hat{y} = f(\mathbf{w}^T \mathbf{x} + b)$  is prediction
- $y$  is true label

## The Complete Algorithm

1. Initialize  $\mathbf{w} = \mathbf{0}$ ,  $b = 0$

2. repeat:

- a. For each  $(\mathbf{x}^{(i)}, y^{(i)})$  in training set:
- b. Compute  $\hat{y}^{(i)} = f(\mathbf{w}^T \mathbf{x}^{(i)} + b)$
- c. If  $\hat{y}^{(i)} \neq y^{(i)}$ :  
 $\mathbf{w} \leftarrow \mathbf{w} + \eta(y^{(i)} - \hat{y}^{(i)})\mathbf{x}^{(i)}$   
 $b \leftarrow b + \eta(y^{(i)} - \hat{y}^{(i)})$

3. until no errors (or max iterations)

---

The mathematical update rule

## What is $\eta$ ?

The learning rate controls step size:

- How much weights change per update
- Typical values: 0.01 to 1.0
- For perceptron: often  $\eta = 1$

## Effects:

### $\eta$ too small:

- Very slow learning
- Many iterations needed
- But stable

### $\eta$ too large:

- May overshoot
- Oscillate around solution
- But faster initially

---

Step size matters: too big or too small both cause problems

## For the Perceptron

### Good news:

For linearly separable data, the perceptron converges regardless of  $\eta > 0$ .

### Why?

The convergence theorem (next slides) guarantees finding a solution if one exists.

### In Practice:

$\eta = 1$  is common for perceptron. Learning rate matters more for:

- Gradient descent (Module 3)
- Non-separable data
- Multi-layer networks

## Setup

Two stocks, two features:

- $\mathbf{x}^{(1)} = (0.5, 0.8)$ ,  $y^{(1)} = 1$  (Buy)
- $\mathbf{x}^{(2)} = (0.2, 0.3)$ ,  $y^{(2)} = 0$  (Sell)

Initialize:  $\mathbf{w} = (0, 0)$ ,  $b = 0$ ,  $\eta = 1$

## Iteration 1: Example 1

- $z = 0 \cdot 0.5 + 0 \cdot 0.8 + 0 = 0$
- $\hat{y} = f(0) = 1$  (threshold at 0)
- $y = 1$ , correct! No update.

## Iteration 1: Example 2

- $z = 0$ ,  $\hat{y} = 1$
- $y = 0$ , wrong!
- $\mathbf{w} \leftarrow (0, 0) + 1(0 - 1)(0.2, 0.3) = (-0.2, -0.3)$
- $b \leftarrow 0 + 1(0 - 1) = -1$

## Iteration 2: Example 1

- $z = -0.2(0.5) - 0.3(0.8) - 1 = -1.34$
- $\hat{y} = 0$
- $y = 1$ , wrong!
- $\mathbf{w} \leftarrow (-0.2, -0.3) + (0.5, 0.8) = (0.3, 0.5)$
- $b \leftarrow -1 + 1 = 0$

## Iteration 2: Example 2

- $z = 0.3(0.2) + 0.5(0.3) + 0 = 0.21$
- $\hat{y} = 1$ ,  $y = 0$ , wrong!
- $\mathbf{w} \leftarrow (0.3, 0.5) - (0.2, 0.3) = (0.1, 0.2)$
- $b \leftarrow 0 - 1 = -1$

Continue until convergence...

---

Following the math with real numbers

# Convergence: Does It Always Work?

## The Perceptron Convergence Theorem

**Theorem (Rosenblatt, 1962):**

If the training data is **linearly separable**, the perceptron learning algorithm will find a separating hyperplane in a **finite** number of updates.

### Key Conditions:

- Data must be linearly separable
- Learning rate  $\eta > 0$
- Cycling through all examples

### Bound on Updates:

$$\text{mistakes} \leq \frac{R^2}{\gamma^2}$$

where  $R = \text{max norm}$ ,  $\gamma = \text{margin}$

[charts/convergence\\_plot/convergence\\_plot.pdf](#)

*"What happens when data isn't linearly separable in financial markets? Can you think of examples?"*

Consider:

**Examples of Non-Separable Data:**

- High P/E growth stocks AND low P/E value stocks both outperform
- Medium-risk investments underperform both conservative and aggressive
- “Buy the rumor, sell the news” patterns

**What Happens to the Perceptron?**

- Never converges
- Oscillates forever
- Best we can do: minimize errors
- Need something more powerful...

**Foreshadowing:** This is exactly why we need **multi-layer** networks (Module 2).

---

Think-Pair-Share: 3 minutes

# The XOR Problem

## The Exclusive OR Function

$x_1$	$x_2$	XOR
0	0	0
0	1	1
1	0	1
1	1	0

### In Words:

Output is 1 if inputs are *different*, 0 if inputs are *same*.

### The Challenge:

Try to draw a single line that separates the 1s from the 0s...


Some patterns cannot be separated by a single line

### The XOR Problem: Why Single-Layer Perceptrons Fail

XOR Truth Table

$x_1$	$x_2$	XOR	Output
0	0	0 XOR 0	0
0	1	0 XOR 1	1
1	0	1 XOR 0	1
1	1	1 XOR 1	0

"Same inputs = 0, Different inputs = 1"



## Geometric Impossibility

Perceptron decision boundary:

$$w_1x_1 + w_2x_2 + b = 0$$

This is always a **straight line**.

XOR requires:

A boundary that curves or has multiple segments.

Linear vs Non-Linear

Linearly Separable:

- AND, OR, NAND, NOR
- One line can separate

Not Linearly Separable:

- XOR, XNOR
- No single line works



charts/linear\_vs\_nonlinear\_patterns/linear\_vs\_nonlinear

# 1969: The Critique That Changed Everything

## Minsky and Papert's Book

"Perceptrons: An Introduction to Computational Geometry" (1969)

### Key Arguments:

1. Single-layer perceptrons cannot compute XOR
2. Many important functions are non-linear
3. No known training algorithm for multi-layer networks
4. Scaling limitations

### The Impact:

The book was rigorous and influential. It convinced funding agencies that neural networks were a dead end.

## The Controversy

### Valid Points:

- Single layers are limited
- XOR problem is real
- No training algorithm existed (then)

### Overstated Points:

- "Neural networks can't work"
- Implied multi-layer networks wouldn't help
- Discouraged research for 15+ years

**Lesson:** Valid criticism of current methods shouldn't stop research into future improvements.

---

Marvin Minsky and Seymour Papert: "Perceptrons" book

# The First AI Winter Begins

## The Collapse

After 1969:

- Funding dried up
- Researchers left the field
- “Neural networks don't work”
- Symbolic AI took over

**Duration:** 1969 to ~1982

## What Survived:

- A few dedicated researchers
- Theoretical work continued quietly
- Hopfield networks (1982)
- Backpropagation (1986)

`charts/ai_winter_timeline/ai_winter_timeline.pdf`

## What We Learned

### 1. Historical Foundation

- McCulloch-Pitts (1943): neurons compute
- Hebb (1949): learning strengthens connections
- Rosenblatt (1958): perceptron learns

### 2. The Perceptron Model

- Weighted sum + threshold
- Linear decision boundary
- Learns from mistakes

### 3. Limitations

- Only linearly separable problems
- XOR is impossible
- Led to AI Winter

charts/module1\_summary\_diagram/module1\_summary\_diagram

*“What if we stack multiple perceptrons?”*

### The Problem We Face

Single perceptrons can only solve linearly separable problems. Real financial data is rarely that simple.

### The Solution Preview:

- Add “hidden” layers
- Non-linear activation functions
- Multi-Layer Perceptrons (MLPs)

### Coming in Module 2:

- How XOR gets solved
- MLP architecture
- Activation functions (sigmoid, ReLU)
- Universal Approximation Theorem
- Loss functions

**Spoiler:** Adding just one hidden layer changes everything.

**Mathematical details for this module: See Appendix A (Perceptron Convergence Proof)**

---

Next: Solving XOR with Multi-Layer Perceptrons