

Module 3: Learning from Mistakes

Training Neural Networks with Backpropagation (1986-2012)

Neural Networks for Finance

BSc Lecture Series

November 26, 2025

"We have the architecture. How does it LEARN?"

What We Know:

- MLP architecture (Module 2)
- Forward pass computation
- Loss functions measure error
- Good weights exist (universal approximation)

What We Don't Know:

- How to find good weights
- How errors guide updates
- Why training sometimes fails
- How to avoid overfitting

This module bridges the gap from architecture to learning.

The fundamental challenge of neural network training

How Traders Improve

A trader's learning process:

1. Make a trade (forward pass)
2. Wait for P&L (loss function)
3. Analyze what went wrong (gradient)
4. Adjust strategy (weight update)
5. Repeat thousands of times (epochs)

Key Insight:

Mistakes are information. Each error tells you how to adjust.

How does a trader improve? By analyzing what went wrong.

Neural Network Training

Trading	Neural Net
Trade execution	Forward pass
P&L calculation	Loss function
Post-trade analysis	Backpropagation
Strategy adjustment	Weight update
Experience	Training epochs

Both learn by **iteratively correcting mistakes**.

Today's Journey

1. Loss Functions (Review)

- Measuring prediction error
- MSE intuition

2. Gradient Descent

- Finding the minimum
- Learning rate tuning

3. Backpropagation

- Credit assignment
- Chain rule in action

4. Training Dynamics

- Batch vs. stochastic
- Epochs and convergence

5. Overfitting

- The enemy of generalization
- The backtest trap

Learning Objectives:

- Understand gradient descent intuitively
- Grasp backpropagation as “blame assignment”
- Recognize and prevent overfitting

From measuring error to updating weights

The Challenge

Given:

- Training data: $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^m$
- Network architecture
- Loss function \mathcal{L}

Find:

- Weights \mathbf{W} and biases \mathbf{b}
- That minimize \mathcal{L}
- And generalize to new data

Scale of the Problem:

A 4-10-5-1 network: 111 parameters

A ResNet-50: 25 million parameters

Why Is This Hard?

Dimensionality:

- Millions of weights to tune
- Exponentially many combinations
- Can't try them all

Non-Convexity:

- Many local minima
- Saddle points
- Flat regions

The Solution:

Gradient-based optimization

"Move downhill in weight space"

Thousands of weights to tune - how do we find the right values?

Yann LeCun at Bell Labs

First commercially deployed neural network:

- Handwritten digit recognition
- Used by US Postal Service
- Read millions of checks
- Proved neural nets could work

Key Innovations:

- Convolutional architecture
- Shared weights
- Backprop through convolutions

charts/timeline_1986_2012/timeline_1986_2012.pdf

The Discovery

Sepp Hochreiter (1991) identified why deep networks fail:

The Problem:

- Gradients multiply through layers
- Sigmoid derivative: max 0.25
- Through 10 layers: $0.25^{10} \approx 10^{-6}$
- Early layers learn nothing

Symptoms:

- Later layers learn quickly
- Early layers stuck at random
- Network never converges

Why Sigmoid Causes Problems

For sigmoid: $\sigma'(z) = \sigma(z)(1 - \sigma(z))$

Maximum value: $\sigma'(0) = 0.25$

Layers	Max Gradient
1	0.25
5	10^{-3}
10	10^{-6}
20	10^{-12}

Implication: Deep networks seemed impossible until ReLU (2010).

Deep networks couldn't learn - gradients disappeared

Long Short-Term Memory

Hochreiter & Schmidhuber solution:

- Designed for sequences
- Explicit “memory” cells
- Gating mechanisms
- Gradients can flow unchanged

Key Innovation:

The “constant error carousel” – a path where gradients don't decay.

Applications:

- Speech recognition
- Machine translation
- Time series prediction

Finance Relevance

LSTMs became popular for:

- Stock price prediction
- Volatility forecasting
- Sentiment analysis
- Algorithmic trading

Why LSTM for Finance?

- Financial data is sequential
- Long-term dependencies matter
- Regime changes persist

Note: Now largely replaced by Transformers (2017).

2012: The ImageNet Moment

AlexNet Wins ImageNet

Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton:

- 15.3% error rate
- Second place: 26.2%
- **40% relative improvement**
- Used GPUs for training

What Made It Work:

1. ReLU activation (not sigmoid)
2. Dropout regularization
3. GPU training (60x faster)
4. Large dataset (1.2M images)
5. Data augmentation

Why This Was Different

Previous Attempts:

- Shallow networks
- Hand-crafted features
- Small datasets
- CPU training

AlexNet:

- 8 layers deep
- Learned features
- Massive data
- GPU parallelism

The Result: Deep learning became the dominant paradigm. Every major AI company pivoted.

AlexNet: Deep learning proves its superiority

What Changed Between 1990 and 2012?

The Ingredients for Success

1. Big Data

- ImageNet: 1.2M labeled images
- Internet made data collection possible
- 1990: thousands of samples

2. Compute Power

- GPUs: 100x speedup
- Moore's law compounding
- Training in days, not years

3. Algorithmic Improvements

- ReLU: no vanishing gradients
- Dropout: better generalization
- Batch normalization (2015)

4. Open Research Culture

- arXiv preprints
- Open-source frameworks
- Reproducibility

Key Insight: The core ideas from 1986 worked – they just needed scale and engineering.

Big data + GPUs + ReLU + dropout = breakthrough

What Does “Wrong” Mean?

Quantifying Prediction Error

We need a function that:

- Takes predictions and labels
- Returns a single number
- Higher = worse predictions
- Differentiable (for gradients)

The Loss Function:

$$\mathcal{L}(\hat{y}, y)$$

Properties We Want:

- $\mathcal{L} \geq 0$ (non-negative)
- $\mathcal{L} = 0$ iff perfect prediction
- Smooth (for optimization)

We need a way to measure how wrong our predictions are

Different Tasks, Different Losses

Task	Loss
Regression	MSE
Binary classification	Cross-entropy
Multi-class	Categorical CE
Ranking	Hinge loss

Finance Examples:

- Return prediction: MSE
- Buy/sell: Binary CE
- Sector classification: Categorical CE

P&L as a Loss Function

For traders:

- P&L = realized gain/loss
- Negative P&L = bad trades
- Goal: maximize P&L

Connection to ML Loss:

- ML loss = prediction error
- Higher loss = worse model
- Goal: minimize loss

Key Difference:

P&L is a *performance metric*.

ML loss is an *optimization target*.

They may not align perfectly!

P&L is the loss function of trading

When P&L \neq Loss

A model might have:

- Low MSE (accurate predictions)
- But low P&L (wrong on big moves)

Or:

- High MSE (noisy predictions)
- But high P&L (right when it matters)

Implication:

Consider using custom loss functions that better align with trading goals.

Module 4 explores this tension.

Total Loss Over Dataset

For m training examples:

$$\mathcal{L}(\mathbf{W}) = \frac{1}{m} \sum_{i=1}^m \ell(\hat{y}^{(i)}, y^{(i)})$$

where:

- ℓ : loss per example
- $\hat{y}^{(i)} = f(\mathbf{x}^{(i)}; \mathbf{W})$: prediction
- $y^{(i)}$: true label
- \mathbf{W} : all network weights

Goal:

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \mathcal{L}(\mathbf{W})$$

Why Average?

Sum vs Average:

- Sum: scales with dataset size
- Average: comparable across datasets
- Gradient magnitude consistent

The Optimization Landscape:

$\mathcal{L}(\mathbf{W})$ defines a surface over weight space.

- High regions: bad weights
- Low regions: good weights
- We seek the lowest point

The loss function quantifies prediction error

The Formula

$$\mathcal{L}_{MSE} = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - \hat{y}^{(i)})^2$$

In Words:

1. Compute error: $y - \hat{y}$
2. Square it: $(y - \hat{y})^2$
3. Average over all samples

Why Squaring?

- Makes all errors positive
- Penalizes large errors heavily
- Mathematically convenient

Example

y	\hat{y}	$(y - \hat{y})^2$
5%	3%	4
-2%	1%	9
8%	7%	1
MSE		4.67

Units: (percentage points)²

RMSE: $\sqrt{MSE} = 2.16\%$

"On average, we're off by about 2%"

"How far off were we, on average?"

Squared Errors as Areas

Each error $(y - \hat{y})^2$ is the area of a square with side length $|y - \hat{y}|$.

MSE = Average Square Area

Why This Matters:

- Error of 4 is 16x worse than error of 1
- Large errors dominate
- Outliers have huge impact

Alternative: MAE

Mean Absolute Error:

$$\mathcal{L}_{MAE} = \frac{1}{m} \sum |y - \hat{y}|$$

More robust to outliers.

[charts/mse_visualization/mse_visualization.pdf](#)

Worked Example

Predictions for 5 Stocks:

Stock	\hat{y}	y	Error ²
AAPL	+5%	+2%	9
MSFT	+3%	+4%	1
GOOG	-1%	+2%	9
AMZN	+4%	+4%	0
META	+2%	-3%	25
MSE		8.8	

RMSE = 2.97%

Interpretation

"On average, our return predictions are off by about 3 percentage points."

Is This Good?

Depends on context:

- Market daily vol: ~1%
- 3% RMSE = 3 std devs
- Not very predictive

Reality Check:

Even small predictability (RMSE slightly < volatility) can be valuable in trading.

Worked example with stock returns

"Why might we want to penalize large errors more than small ones in stock prediction?"

Consider:

Arguments For (Use MSE):

- Big errors are costlier
- Crashes matter more than small moves
- Position sizing affected
- Risk management

Arguments Against (Use MAE):

- Markets have fat tails
- Outliers can dominate MSE
- May optimize for rare events
- Robustness to noise

Reality: Many practitioners use MAE or Huber loss (combines both) for financial applications.

Think-Pair-Share: 3 minutes

Loss as a Function of Weights

$$\mathcal{L}(\mathbf{W})$$

For every choice of weights, there's a loss value.

In 2D (two weights):

A surface we can visualize.

In High Dimensions:

A hypersurface we navigate blindly.

Features:

- Global minimum (best)
- Local minima (traps)
- Saddle points
- Flat regions (plateaus)



`charts/loss_landscape_3d/loss_landscape_3d.pdf`

The Challenge

Find:

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \mathcal{L}(\mathbf{W})$$

Difficulties:

- Millions of dimensions
- Non-convex landscape
- No closed-form solution
- Can't try all possibilities

We Need:

An *iterative* algorithm that gradually improves weights.

Possible Approaches

Random Search:

- Try random weights
- Keep best so far
- **Hopelessly slow**

Grid Search:

- Try all combinations
- 10^{100} possibilities
- **Impossible**

Gradient-Based:

- Use local slope information
- Move toward improvement
- **Tractable!**

How do we find the weights that minimize loss?

The Blind Hiker Analogy

The Scenario

Imagine you're:

- Blindfolded
- On a mountainside
- Trying to reach the valley
- Can only feel the local slope

What Would You Do?

1. Feel the ground around you
2. Determine which way is downhill
3. Take a step in that direction
4. Repeat until you reach a valley

Neural Network Translation

Hiker	Network
Position	Weights \mathbf{W}
Altitude	Loss \mathcal{L}
Slope	Gradient $\nabla \mathcal{L}$
Step	Weight update
Valley	Minimum loss

Key Insight:

We don't need to see the whole landscape. Local slope is enough!

"You're blindfolded on a mountain. How do you find the valley?"

Answer: Feel the Slope

The Strategy

1. Compute the slope (gradient)
2. Move opposite to the slope
3. Repeat until convergence

Why Opposite?

- Gradient points uphill
- We want to go downhill
- Move in negative gradient direction

The Update Rule:

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \nabla_{\mathbf{W}} \mathcal{L}$$

Move in the direction that goes down

Gradient Descent Algorithm

1. Initialize \mathbf{W} randomly
2. **repeat**:
 - a. Compute loss $\mathcal{L}(\mathbf{W})$
 - b. Compute gradient $\nabla \mathcal{L}$
 - c. Update: $\mathbf{W} \leftarrow \mathbf{W} - \eta \nabla \mathcal{L}$
3. **until** convergence

η = learning rate (step size)

What Is the Gradient?

The gradient $\nabla \mathcal{L}$ is a vector of partial derivatives:

$$\nabla_w \mathcal{L} = \begin{pmatrix} \frac{\partial \mathcal{L}}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}}{\partial w_n} \end{pmatrix}$$

Each Component:

$\frac{\partial \mathcal{L}}{\partial w_i}$ = How much does loss change if we change w_i slightly?

The gradient tells us which way is “up”

Properties

Direction:

- Points toward steepest increase
- $-\nabla \mathcal{L}$ points toward steepest decrease

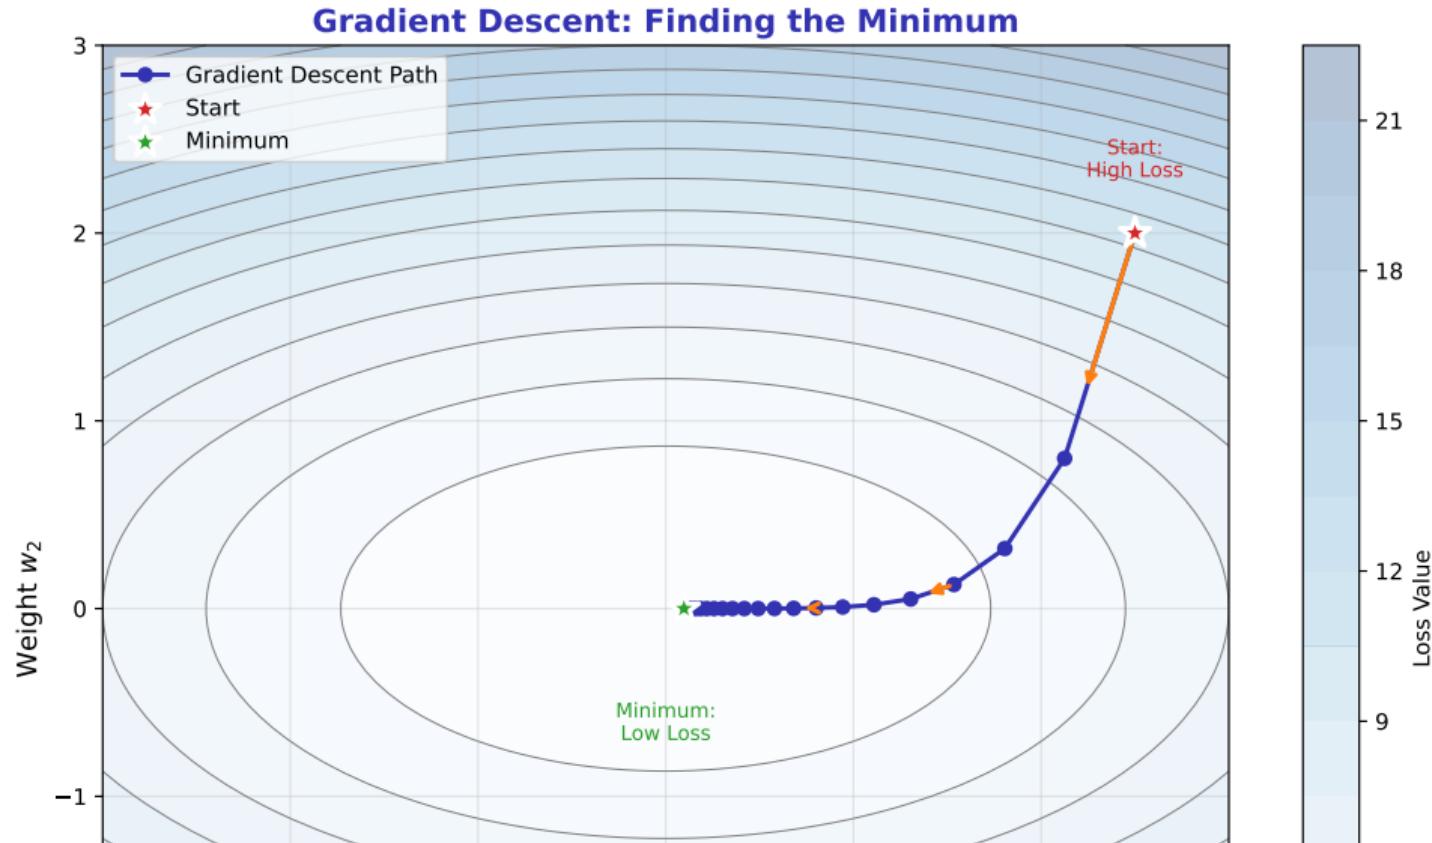
Magnitude:

- $\|\nabla \mathcal{L}\|$ = slope steepness
- Near minimum: gradient ≈ 0

At a Minimum:

$$\nabla \mathcal{L} = \mathbf{0}$$

No direction goes further down.



Portfolio Adjustment

Similar iterative process:

1. Evaluate current portfolio
2. Estimate sensitivities ("greeks")
3. Adjust positions to reduce risk
4. Repeat periodically

Delta Hedging:

- Measure option delta
- Adjust stock position
- Move toward neutral

Similar to iterative portfolio rebalancing

Comparison

GD	Portfolio
Loss	Risk/Variance
Weights	Positions
Gradient	Sensitivities
Learning rate	Trading aggressiveness
Convergence	Optimal allocation

Key Difference:

Markets change continuously. Portfolios must adapt.
Neural networks train once (mostly).

The Learning Rate: Step Size

The Hyperparameter η

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \nabla \mathcal{L}$$

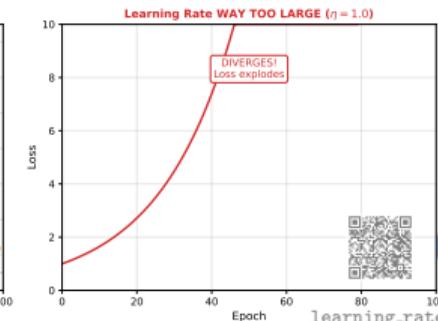
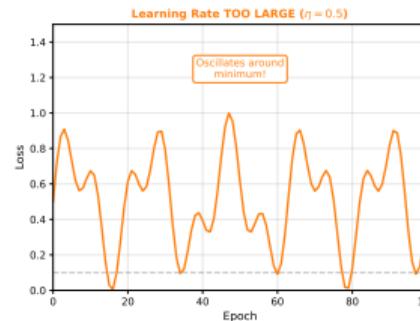
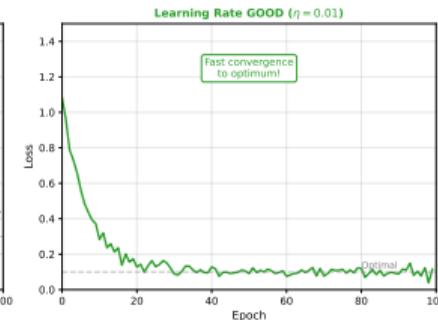
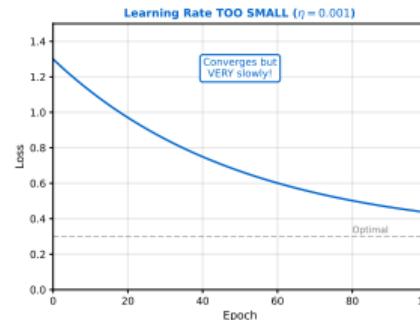
η Controls:

- Size of each weight update
- Speed of convergence
- Stability of training

Typical Values:

- 10^{-4} to 10^{-1}
- Often starts at 0.01 or 0.001
- May decrease during training

Learning Rate: The Most Important Hyperparameter



Rule of thumb: Start with 0.001-0.01 and adjust based on training dynamics

Learning rate controls how far we move each step

The Problem

When η is too large:

- Steps overshoot the minimum
- May jump to worse regions
- Loss oscillates or explodes
- Training diverges

Symptoms:

- Loss goes up, not down
- Loss becomes NaN
- Weights grow very large
- Erratic training curves

Finance Analogy

Overtrading:

- Adjusting positions too aggressively
- Chasing every signal
- Transaction costs accumulate
- Portfolio becomes unstable

Solution:

Reduce learning rate until stable.

Rule of Thumb: If loss explodes, halve η .

Too big = overshoot the minimum

The Problem

When η is too small:

- Steps are tiny
- Progress is slow
- May get stuck in flat regions
- Training takes forever

Symptoms:

- Loss decreases very slowly
- Many epochs with little improvement
- May stop before reaching minimum
- Wasted computation

Finance Analogy

Underreacting:

- Ignoring market signals
- Missing opportunities
- Portfolio drifts from target
- Slow adaptation to regime changes

Solution:

Increase learning rate or use adaptive methods.

Modern Practice: Adaptive optimizers (Adam, RMSprop) adjust η automatically.

Too small = converge too slowly

"In trading, what's analogous to learning rate? What happens if you adjust positions too aggressively or too conservatively?"

Consider:

Position Sizing:

- How much to trade per signal
- Kelly criterion vs. fractional Kelly
- Risk management constraints

Rebalancing Frequency:

- How often to adjust
- Transaction cost vs. tracking error
- Market impact considerations

Key Insight: Both trading and ML require balancing responsiveness against stability.

Think-Pair-Share: 3 minutes

The Challenge

We know:

- The output was wrong
- We need to update weights
- There are thousands of weights

The Question:

Which weights caused the error?

Credit Assignment:

Attributing output error to individual weights deep in the network.

Why Is This Hard?

Direct Attribution:

- Output layer weights: clear influence
- Hidden layer weights: indirect
- Early layers: very indirect

The Chain of Influence:

$$w_1 \rightarrow h_1 \rightarrow h_2 \rightarrow \dots \rightarrow \hat{y} \rightarrow \mathcal{L}$$

Each weight affects the loss through many intermediate steps.

The output was wrong. Which weights caused it?

Attribution in Trading

A portfolio lost money. Why?

1. Macro call wrong?
2. Sector allocation off?
3. Stock selection bad?
4. Timing poor?
5. Execution costly?

Performance Attribution:

- Decompose returns by factor
- Trace P&L to decisions
- Learn which calls were wrong

Neural Network Attribution

Trading	Neural Net
Macro view	Early layers
Sector allocation	Hidden layers
Stock picks	Later layers
Final trades	Output
P&L	Loss

Backpropagation is the neural network's performance attribution algorithm.

"Which decisions led to this P&L?"

Backpropagation: Blame Assignment

The Algorithm

Backpropagation computes $\frac{\partial \mathcal{L}}{\partial w}$ for every weight w in the network.

Key Idea:

Work backward from output to input, propagating error attribution.

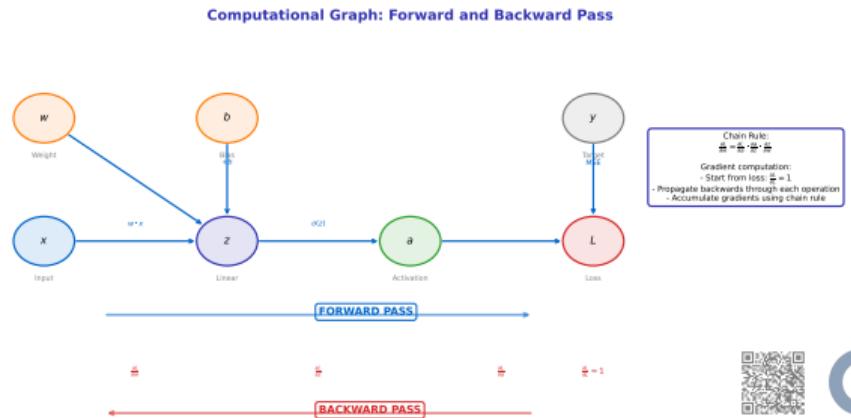
Two Passes:

1. **Forward Pass:** Compute outputs
2. **Backward Pass:** Compute gradients

Efficiency:

Computes ALL gradients in time proportional to one forward pass.

Propagating error backward through the network



backprop.computational-graph

The Chain Rule: Intuition

The Core Mathematical Tool

If A affects B and B affects C:

$$\frac{\partial C}{\partial A} = \frac{\partial C}{\partial B} \cdot \frac{\partial B}{\partial A}$$

Example:

Temperature → Ice cream sales → Profit

How does temperature affect profit?

$$\frac{\partial \text{Profit}}{\partial \text{Temp}} = \frac{\partial \text{Profit}}{\partial \text{Sales}} \cdot \frac{\partial \text{Sales}}{\partial \text{Temp}}$$

Chain Rule: Foundation of Backpropagation

Chain Rule: The Key to Backprop



Chain Rule:

$$\frac{d}{dx} f(g(x)) = \frac{df}{dg} \frac{dg}{dx}$$

Intuition: Rate of change multiplies through the chain

Example:
 $f(u) = u^2, g(x) = 3x + 1$
 $\frac{df}{du} = 2u, \frac{dg}{dx} = 3$
 $\frac{d}{dx} f(g(x)) = 2(3x+1) \cdot 3 = 6(3x+1)$

Multi-Variable Chain Rule



Total Derivative (sum over all paths):

$$\frac{\partial v}{\partial x} = \frac{\partial v}{\partial z} \frac{\partial z}{\partial x} + \frac{\partial v}{\partial w} \frac{\partial w}{\partial x}$$

In neural networks: gradients flow back through ALL paths



chain_rule_visualization

"If A affects B and B affects C, how does A affect C?"

Chain of Effects

Fed Rate → Mortgages → Housing → Banks → Portfolio

How does Fed rate affect your portfolio?

$$\frac{\partial \text{Portfolio}}{\partial \text{Fed}} = \frac{\partial P}{\partial B} \cdot \frac{\partial B}{\partial H} \cdot \frac{\partial H}{\partial M} \cdot \frac{\partial M}{\partial F}$$

Each Link:

- Fed → Mortgages: rate sensitivity
- Mortgages → Housing: demand elasticity
- Housing → Banks: credit exposure
- Banks → Portfolio: position size

Effects propagate through chains of influence

Neural Network Parallel

Finance	Neural Net
Fed rate	Input x
Mortgages	Hidden layer 1
Housing	Hidden layer 2
Banks	Hidden layer 3
Portfolio	Output

Backprop does this automatically:

Chains together all the local sensitivities to get the total effect of each input/weight on the loss.

At the Output

For output weight $w^{(L)}$:

$$\frac{\partial \mathcal{L}}{\partial w^{(L)}} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z^{(L)}} \cdot \frac{\partial z^{(L)}}{\partial w^{(L)}}$$

Each Term:

- $\frac{\partial \mathcal{L}}{\partial \hat{y}}$: How loss changes with output
- $\frac{\partial \hat{y}}{\partial z^{(L)}}$: Activation derivative
- $\frac{\partial z^{(L)}}{\partial w^{(L)}}$: Input from previous layer

For MSE + Sigmoid:

$$\frac{\partial \mathcal{L}}{\partial w^{(L)}} = (\hat{y} - y) \cdot \hat{y}(1 - \hat{y}) \cdot a^{(L-1)}$$

Output Error ($\delta^{(L)}$)

Define the “error signal”:

$$\delta^{(L)} = \frac{\partial \mathcal{L}}{\partial z^{(L)}}$$

For MSE loss + sigmoid:

$$\delta^{(L)} = (\hat{y} - y) \cdot \sigma'(z^{(L)})$$

Then:

$$\frac{\partial \mathcal{L}}{\partial w^{(L)}} = \delta^{(L)} \cdot a^{(L-1)}$$

This is just error \times input!

At the output, error attribution is straightforward

The Key Insight

Hidden layer error comes from downstream:

$$\delta^{(l)} = ((W^{(l+1)})^T \delta^{(l+1)}) \odot \sigma'(z^{(l)})$$

In Words:

1. Take error from next layer ($\delta^{(l+1)}$)
2. Multiply by weights connecting to next layer
3. Scale by local activation derivative

Error Flows Backward:

Output \rightarrow Last hidden $\rightarrow \dots \rightarrow$ First hidden

Why This Works

Chain rule connects layers:

$$\frac{\partial \mathcal{L}}{\partial z^{(l)}} = \sum_j \frac{\partial \mathcal{L}}{\partial z_j^{(l+1)}} \cdot \frac{\partial z_j^{(l+1)}}{\partial z^{(l)}}$$

Gradient for Hidden Weight:

$$\frac{\partial \mathcal{L}}{\partial w^{(l)}} = \delta^{(l)} \cdot a^{(l-1)}$$

Same formula as output layer!

Hidden layer gradients require the chain rule

The Complete Training Loop

One Training Step

1. Forward Pass

- Compute all activations
- Get prediction \hat{y}

2. Compute Loss

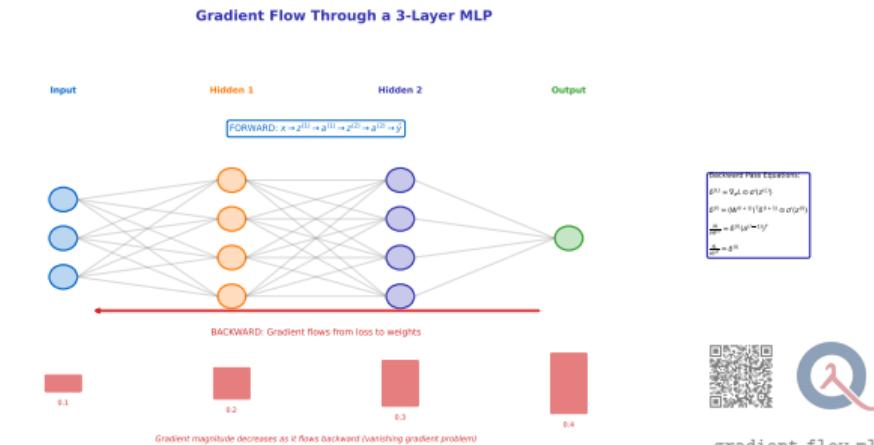
- $\mathcal{L} = \ell(\hat{y}, y)$

3. Backward Pass

- Compute $\delta^{(L)}$ at output
- Propagate backward to get all $\delta^{(l)}$
- Compute all weight gradients

4. Update Weights

- $\mathbf{W} \leftarrow \mathbf{W} - \eta \nabla \mathcal{L}$



Forward pass, compute loss, backward pass, update weights

Why “Backpropagation”?

The Name

“Back-propagation of errors”

Information Flow:

Forward:

- Data flows input → output
- Activations computed layer by layer

Backward:

- Errors flow output → input
- Gradients computed layer by layer

Symmetry:

Each layer: one forward operation, one backward operation.

Historical Note

The Algorithm:

- Werbos (1974): first derivation
- Rumelhart et al. (1986): popularized
- Now standard in all deep learning

Modern Perspective:

Backprop is just automatic differentiation applied to neural networks.

Frameworks (PyTorch, TensorFlow):

Compute gradients automatically – you just specify the forward pass!

Error information flows from output to input

"Why do deeper networks make training harder? What happens to gradients as they flow backward through many layers?"

Consider:

Vanishing Gradients:

- Sigmoid: max derivative 0.25
- Through 10 layers: 0.25^{10}
- Early layers get tiny gradients
- Learn extremely slowly

Exploding Gradients:

- If derivatives > 1
- Gradients grow exponentially
- Weights become huge
- Training diverges

Solutions: ReLU, batch normalization, residual connections, careful initialization.

Think-Pair-Share: 3 minutes

Batch Gradient Descent

Definition

Use **all** training data to compute gradient:

$$\nabla \mathcal{L} = \frac{1}{m} \sum_{i=1}^m \nabla \ell(\hat{y}^{(i)}, y^{(i)})$$

Then update weights once.

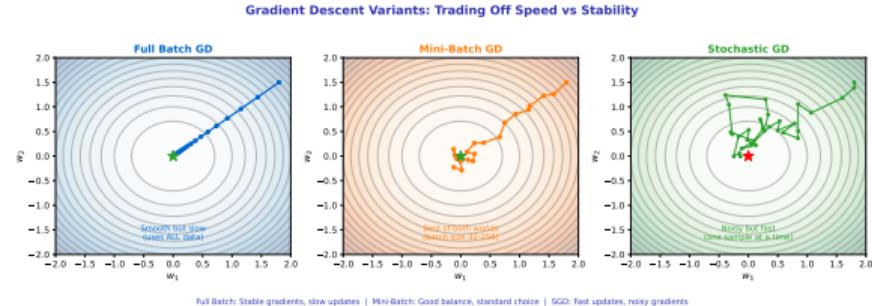
Advantages:

- + Stable gradient estimate
- + Deterministic updates
- + Guaranteed descent direction

Disadvantages:

- Slow for large datasets
- Must load all data in memory
- One update per full pass

Compute gradient using the entire dataset



batch_vs_stochastic

Definition

Update after **each** single example:

$$\nabla \mathcal{L} \approx \nabla \ell(\hat{y}^{(i)}, y^{(i)})$$

One sample = one update.

Advantages:

- + Very fast updates
- + Can handle huge datasets
- + Noise helps escape local minima
- + Online learning possible

Disadvantages:

- Noisy gradient estimate
- Erratic convergence
- May not settle at minimum

Update after each single example

Why “Stochastic”?

Random sampling of training examples introduces randomness into gradient.

Expected Value:

$$\mathbb{E}[\nabla \ell^{(i)}] = \nabla \mathcal{L}$$

On average, SGD points in the right direction.

Variance:

Individual updates are noisy, but noise can help exploration.

Mini-Batch: The Sweet Spot

Definition

Use small batches of B examples:

$$\nabla \mathcal{L} \approx \frac{1}{B} \sum_{i=1}^B \nabla \ell(\hat{y}^{(i)}, y^{(i)})$$

Typical B : 32, 64, 128, 256

Advantages:

- + Reduced variance vs SGD
- + GPU parallelization
- + Reasonable memory usage
- + Frequent updates

The Modern Default

Batch Size Trade-offs

Size	Noise	Speed
1 (SGD)	High	Fast updates
32-256	Medium	Best practice
Full batch	Low	Slow updates

Large Batch Issues:

- May converge to sharp minima
- Worse generalization
- Need learning rate scaling

Balance between efficiency and noise

Epochs: Full Passes Through Data

Definition

Epoch = one complete pass through all training data.

With Mini-Batches:

- 10,000 samples
- Batch size 100
- 100 updates per epoch

Typical Training:

- 10-1000 epochs
- Monitor loss curve
- Stop when converged

Training Timeline

Stage	Behavior
Early epochs	Loss drops quickly
Middle epochs	Progress slows
Late epochs	Diminishing returns

When to Stop?

- Loss stops improving
- Validation loss increases (overfitting!)
- Resource constraints

Training typically requires multiple epochs

Training Curves

What to Plot

- Training loss vs. epoch
- Validation loss vs. epoch
- Learning rate schedule
- Gradient norms (debugging)

Healthy Training:

- Both losses decrease
- Validation tracks training
- Smooth convergence

Warning Signs:

- Training drops, validation rises
- Loss oscillates wildly
- Loss becomes NaN



Monitoring progress during training

Worked Example: One Training Step

Simple 2-2-1 Network

Given:

- Input: $\mathbf{x} = (0.5, 0.8)^T$
- Target: $y = 1$
- Current weights (simplified)

Forward Pass:

$$z^{(1)} = W^{(1)}\mathbf{x} + b^{(1)}$$

$$a^{(1)} = \sigma(z^{(1)})$$

$$z^{(2)} = W^{(2)}a^{(1)} + b^{(2)}$$

$$\hat{y} = \sigma(z^{(2)}) = 0.62$$

Loss and Backward

Loss:

$$\mathcal{L} = \frac{1}{2}(y - \hat{y})^2 = \frac{1}{2}(1 - 0.62)^2 = 0.072$$

Backward Pass:

$$\begin{aligned}\delta^{(2)} &= (0.62 - 1) \cdot 0.62(1 - 0.62) \\ &= -0.089\end{aligned}$$

Weight Gradient:

$$\frac{\partial \mathcal{L}}{\partial W^{(2)}} = \delta^{(2)} \cdot a^{(1)}$$

Update:

$$W^{(2)} \leftarrow W^{(2)} - 0.1 \cdot \nabla W^{(2)}$$

Following the numbers through one training step

The Vanishing Gradient Problem

The Problem

Gradients shrink as they flow backward:

$$\delta^{(l)} \propto \prod_{k=l}^{L-1} \sigma'(z^{(k)})$$

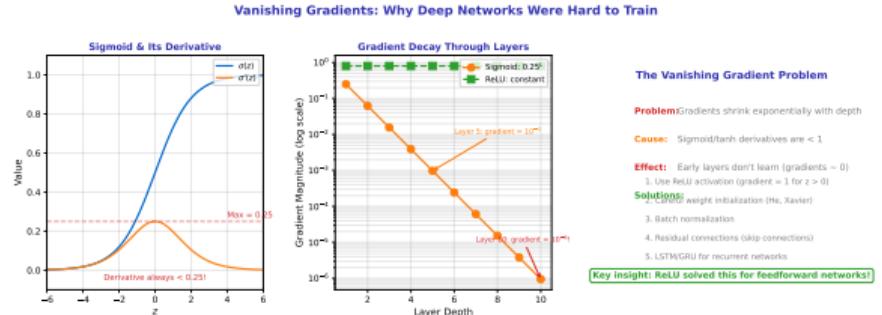
For sigmoid: $\sigma'(z) \leq 0.25$

Through 10 layers: gradient $\times 10^{-6}$

Symptoms:

- Early layers don't learn
- Deep networks fail to train
- Loss plateaus quickly

Deep networks: gradients can become vanishingly small



vanishing_gradient дем

This Module: Intuition

We covered:

- Why backprop works (chain rule)
- How errors flow backward
- Update rule intuition
- Training dynamics

What We Skipped:

- Full mathematical derivation
- Matrix calculus details
- Vectorized implementations
- Automatic differentiation theory

Appendix B Contains:

1. Chain rule setup
2. Output layer error derivation
3. Hidden layer recursion formula
4. Complete gradient equations
5. Weight and bias gradients
6. Algorithm pseudocode

For the mathematically curious:

The appendix provides the rigorous derivation with all matrix calculus steps.

See Appendix B for complete backpropagation derivation

What Is Overfitting?

Definition

Overfitting: When a model learns the training data too well, including its noise, and fails to generalize.

Analogy:

A student who memorizes exam answers but doesn't understand the material.

Symptoms:

- Training loss: very low
- Test loss: high
- Model is “too confident”

Why It Happens

Model Complexity:

- Too many parameters
- Can fit any training data perfectly
- Including noise

Limited Data:

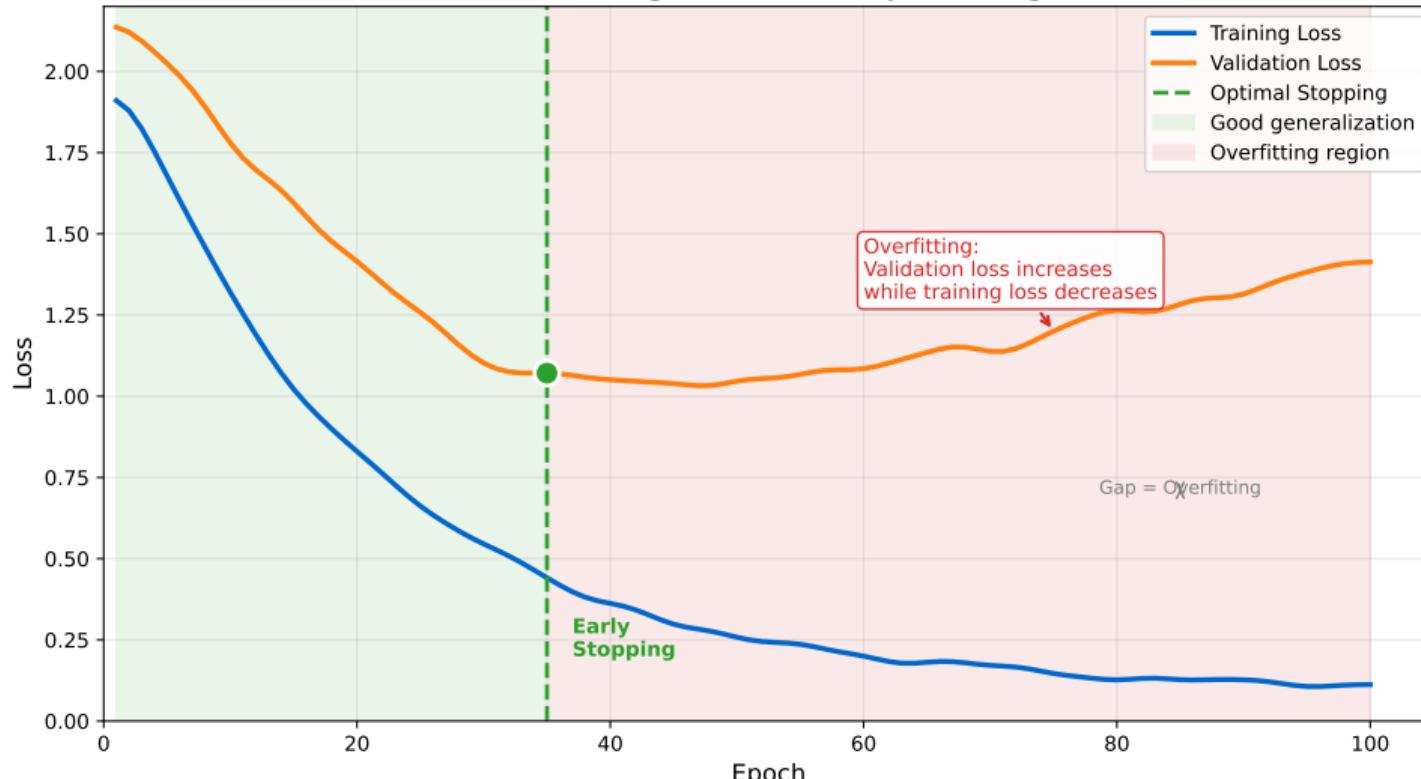
- Not enough examples
- Training set not representative
- Noise gets learned as signal

Training Too Long:

- Model eventually memorizes
- Needs early stopping

When your model memorizes instead of learns

Overfitting: When to Stop Training



Finance Analogy: Like backtesting a strategy that works perfectly on historical data but fails in live trading



The Trap

Every trading strategy looks good on historical data – that's how you found it!

The Process:

1. Try many strategies
2. Keep the one that worked best
3. By construction, it fits the past
4. Future performance? Unknown.

Multiple Testing:

- Try 1000 random strategies
- Best one has Sharpe 2.0
- Is it skill or luck?

"Every strategy looks good on historical data"

Why Finance Overfits Easily

1. Limited Data

- 20 years = 5000 trading days
- Few independent observations

2. Low Signal-to-Noise

- Markets are noisy
- Easy to fit noise

3. Non-Stationarity

- Regimes change
- Past may not predict future

4. Look-Ahead Bias

- Using future information
- Subtle but deadly

Data Limitations

Domain	Samples
ImageNet	1,200,000
MNIST	60,000
Stock returns (daily, 10y)	2,520
Stock returns (monthly, 50y)	600
Market crashes	~10

The Problem:

Neural networks have thousands of parameters but only thousands of data points.

Signal vs Noise

Image Classification:

- A cat is always a cat
- Signal is strong and consistent
- R^2 can reach 99%+

Stock Prediction:

- Returns are mostly random
- Signal is weak and changing
- R^2 of 1% is excellent!

Implication:

Standard ML practices don't directly transfer to finance.

Limited data, high noise, non-stationary markets

Train/Validation/Test Split

1. Training Set (60-80%)

- Used to fit weights

2. Validation Set (10-20%)

- Used to tune hyperparameters
- Monitor for overfitting

3. Test Set (10-20%)

- Final evaluation only
- Touch only once!

Key Rule:

Never use test data for decisions.

Always monitor out-of-sample performance

Warning Signs

Overfitting Indicators:

- Training loss \ll validation loss
- Validation loss starts increasing
- Model predictions are “too confident”
- Performance degrades out-of-sample

For Finance:

- Backtest Sharpe \gg live Sharpe
- Strategy “stops working”
- Drawdowns worse than expected

"How would you know if your stock prediction model is overfitting? What specific symptoms would you look for?"

Consider:

In Training:

- Training/validation gap
- Validation loss trend
- Prediction confidence

In Production:

- Live vs. backtest performance
- Regime sensitivity
- Transaction cost impact

Best Practice: Always maintain a truly out-of-sample test set that you evaluate only once.

Think-Pair-Share: 3 minutes

Solutions (Module 4)

1. L1/L2 Regularization

- Penalize large weights
- Simpler models

2. Dropout

- Randomly disable neurons
- Ensemble effect

3. Early Stopping

- Stop before overfitting
- Use validation loss

4. Data Augmentation

- Create more training data
- Finance: bootstrap?

Finance-Specific

1. Walk-Forward Validation

- Respect time ordering
- Rolling windows

2. Cross-Validation Variants

- Purged CV
- Combinatorial CV

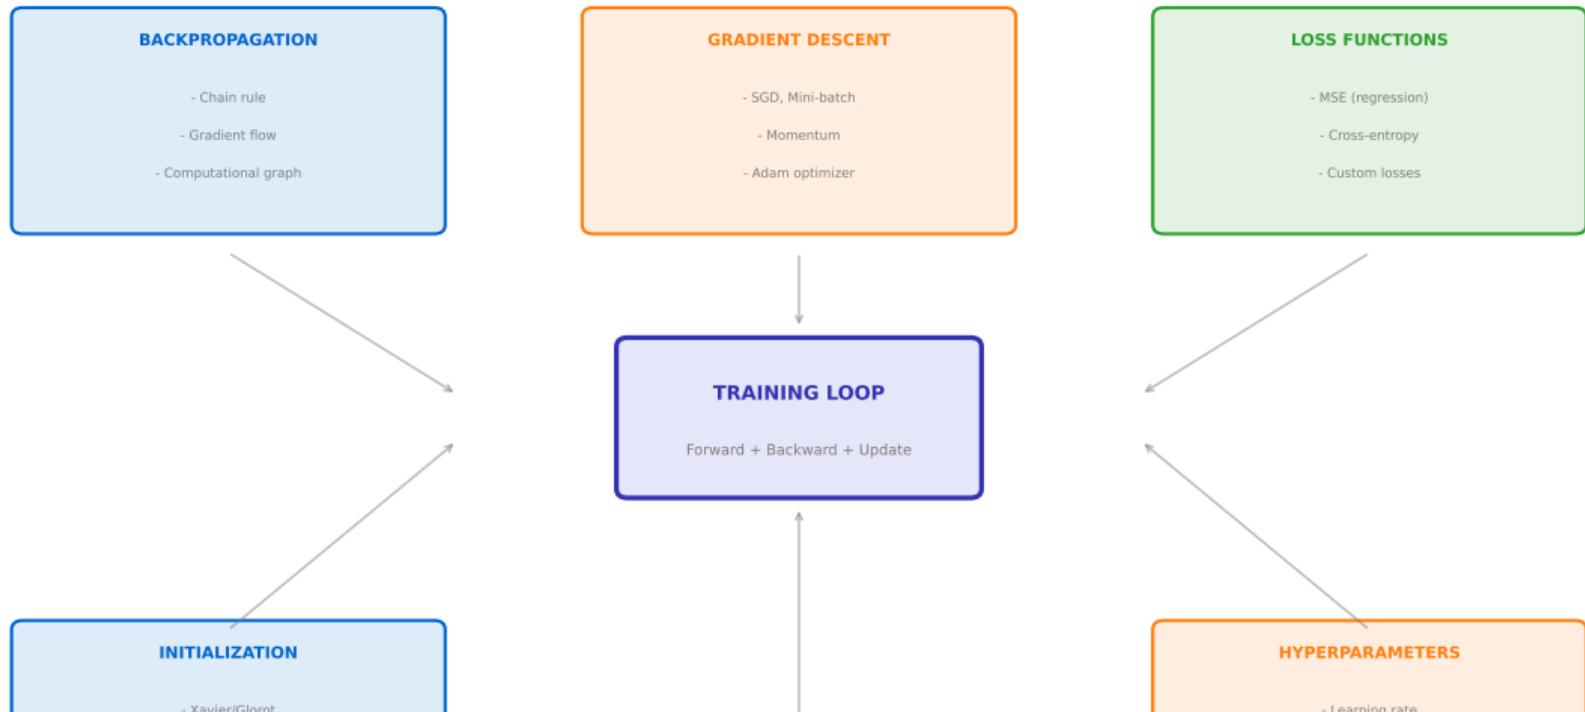
3. Ensemble Methods

- Average multiple models
- Reduce variance

Module 4 will cover these in detail.

Module 4 will cover solutions: regularization, dropout, early stopping

Module 3 Summary: Training Neural Networks



What We Learned

1. Loss Functions

- Measure prediction error
- MSE, cross-entropy
- Define what “good” means

2. Gradient Descent

- Follow the slope downhill
- Learning rate matters
- Batch vs stochastic

3. Backpropagation

- Chain rule for credit assignment
- Error flows backward
- Enables efficient gradient computation

4. Training Dynamics

- Epochs and batches
- Monitoring with curves
- Vanishing gradients

5. Overfitting

- Memorizing vs learning
- Train/val/test split
- Finance-specific challenges

The Big Picture:

We can now train neural networks. But making them work well requires more...

From measuring error to updating weights

Modules 1-3 Foundation

1. Module 1: Architecture

- Perceptron basics
- Linear decision boundaries
- Limitations (XOR)

2. Module 2: MLPs

- Hidden layers
- Non-linear activation
- Universal approximation

3. Module 3: Training

- Gradient descent
- Backpropagation
- Overfitting awareness

You Can Now:

- Explain how neural networks compute
- Understand the training process
- Recognize overfitting
- Follow the math (or know where to look)

What's Missing:

- Practical regularization
- Real-world applications
- Finance case studies
- Modern developments

Modules 1-3: The complete neural network foundation

Key Questions for Reflection

Think about these as you move to Module 4:

1. Loss vs. Profit:

Why might minimizing MSE not maximize trading profit? What loss function would better align with trading goals?

2. Overfitting in Finance:

With only 20 years of daily data, how many parameters can we safely learn? What's the ratio of samples to parameters you'd be comfortable with?

3. Non-Stationarity:

If market regimes change, what does that mean for our training strategy? Should we weight recent data more heavily?

4. The Efficient Market Hypothesis:

If markets are efficient, can neural networks find persistent patterns? What would success look like?

Reflect on the learning process

"Theory meets practice. How do we actually use neural networks in finance?"

Coming Up:

- Regularization techniques
 - L1/L2, dropout, early stopping
- Financial data challenges
 - Non-stationarity, noise
- Complete case study
 - Stock prediction end-to-end

Also:

- Modern architectures overview
 - CNN, RNN, Transformers
- Limitations and ethics
 - Black-box decisions
 - Regulatory concerns
- Future directions
 - Where the field is heading

Mathematical details: See Appendix B-D for derivations

Next: Regularization, case studies, and modern developments